

# EECS151/251A

## Introduction to Digital Design and ICs

### Lecture 10: RISC-V Pipelining Sophia Shao



#### Tesla's Dojo

To say Tesla is merely interested in machine learning is an understatement. The electric car maker built an in-house supercomputer named Dojo, optimized for training its machine learning models. Unlike many other supercomputers, Dojo isn't using off-the-shelf CPUs and GPUs, such as from AMD, Intel, or Nvidia. Instead, Tesla designed their own microarchitecture tailored to their needs, letting them make tradeoffs that more general architectures can't make. In this article, we're going to take a look at that architecture, based on Tesla's presentations at Hot Chips. The architecture doesn't have a separate name, so for simplicity, whenever we mention Dojo further down we're talking about the architecture.

<https://chipsandcheese.com/2022/09/01/hot-chips-34-teslas-dojo-microarchitecture/>



# Review


- We have covered the implementation of the base ISA for RV32I!!!
  - Get yourself familiar with the ISA Spec.
- Instruction type:
  - R-type
  - I-type
  - S-type
  - B-type
  - J-type
  - U-type
- Implementation suggested is straightforward, yet there are modalities in how to implement it well at gate level.
- Single-cycle datapath is slow – need to pipeline it



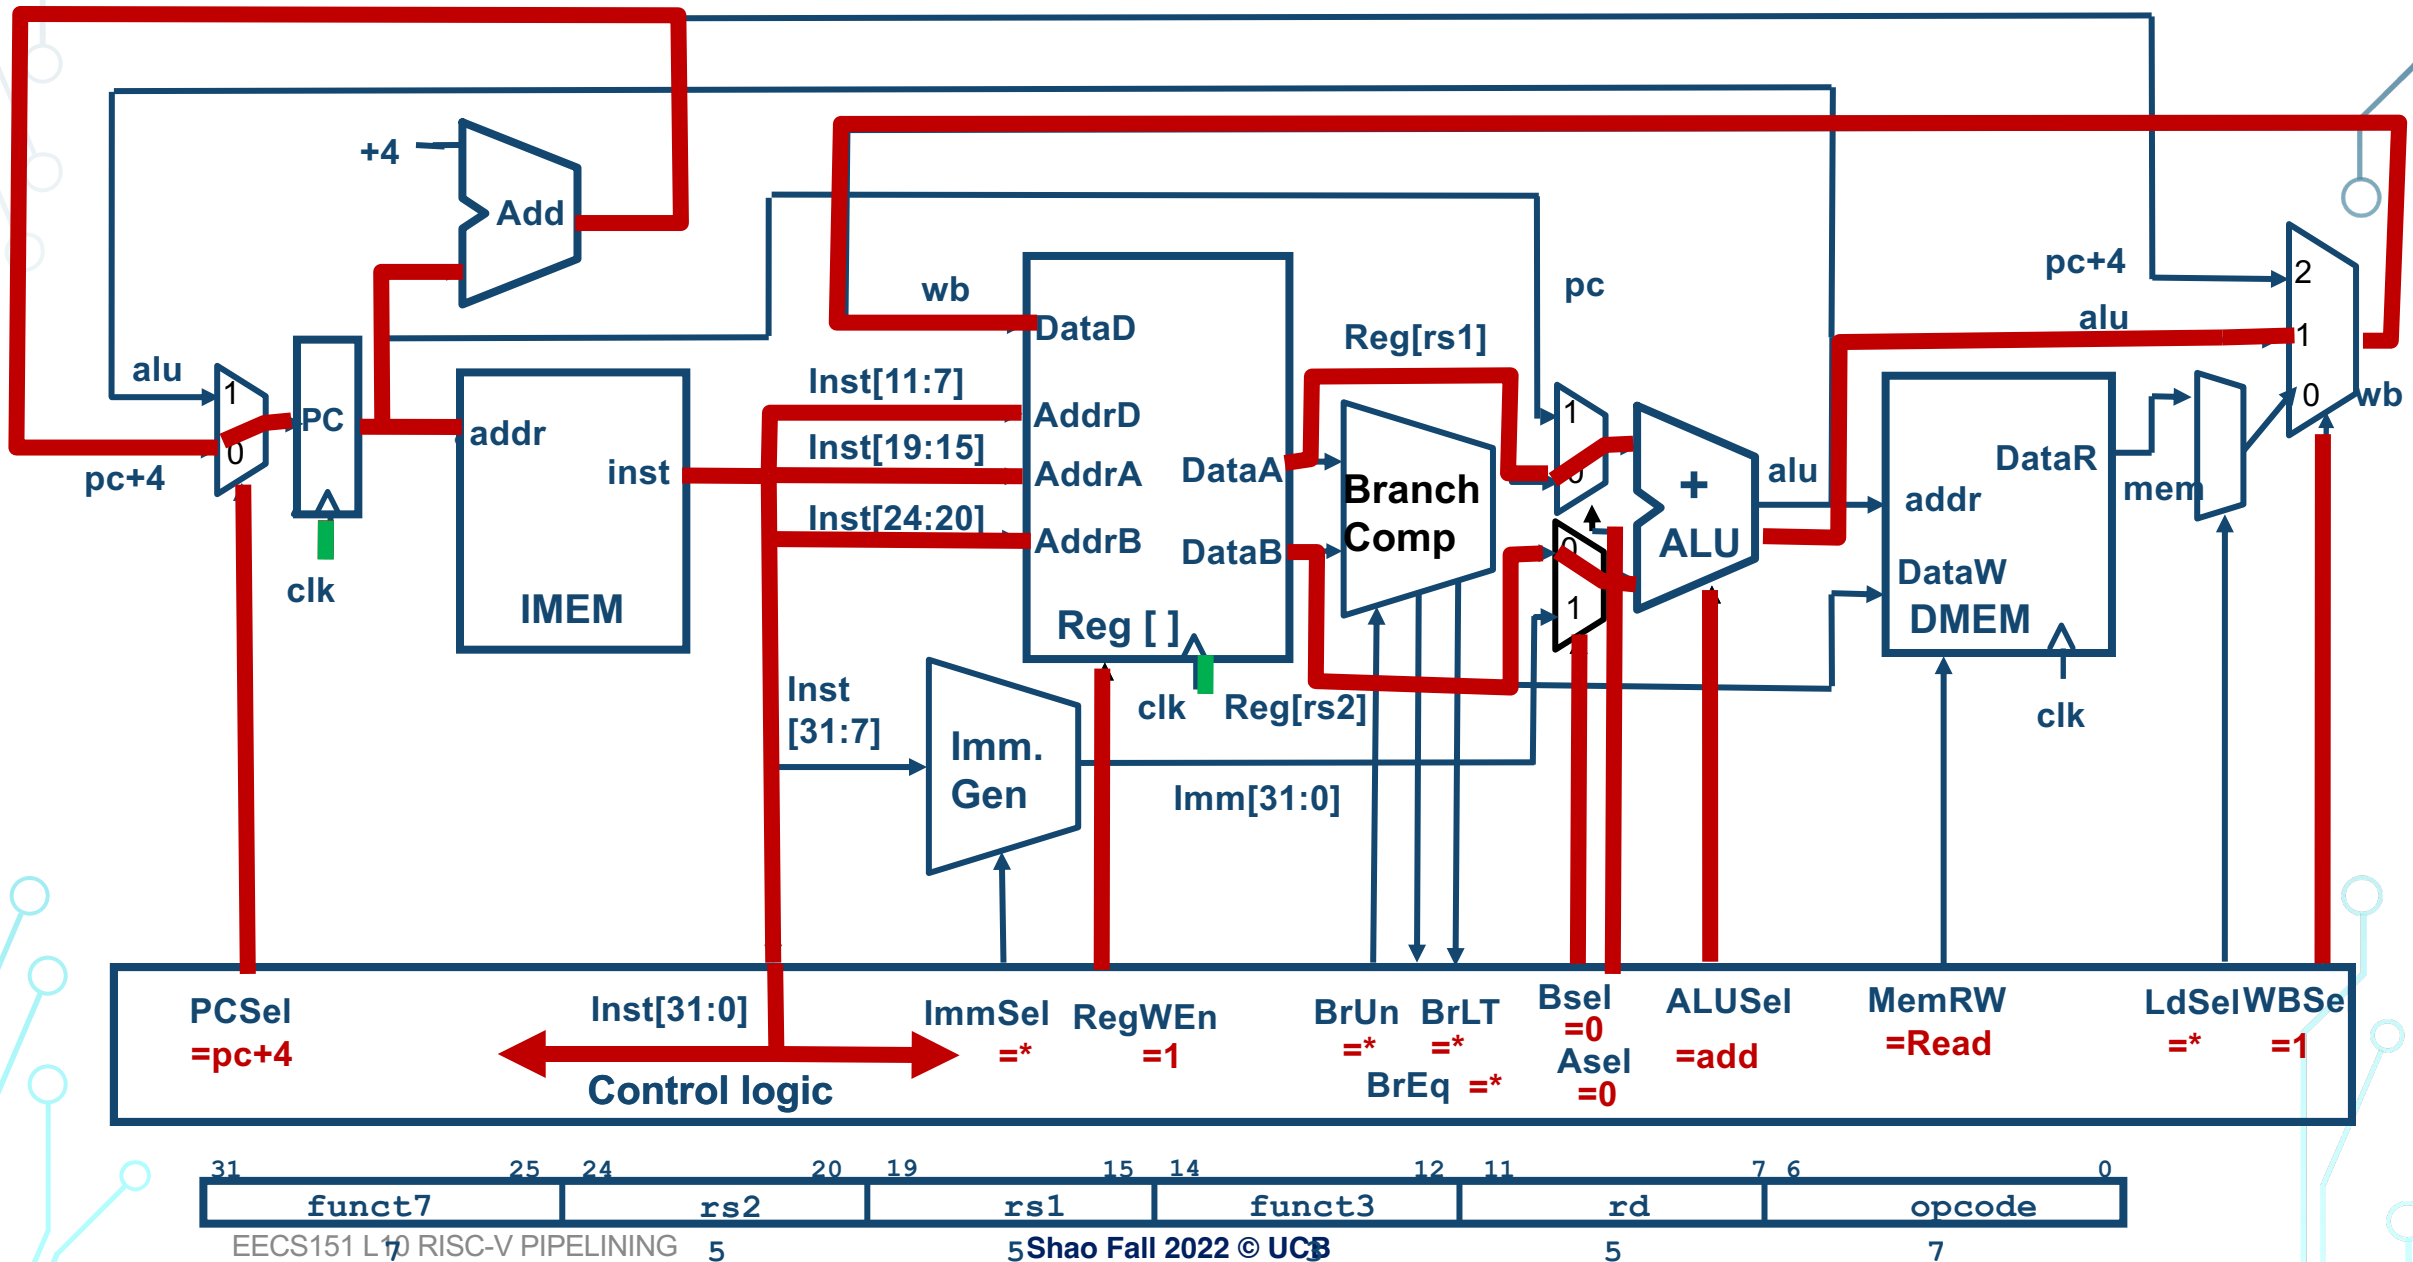
## • RISC-V Datapath & Control

- R-type
- I-type
- S-type
- B-type
- J-type
- U-type
- **Control Logic**

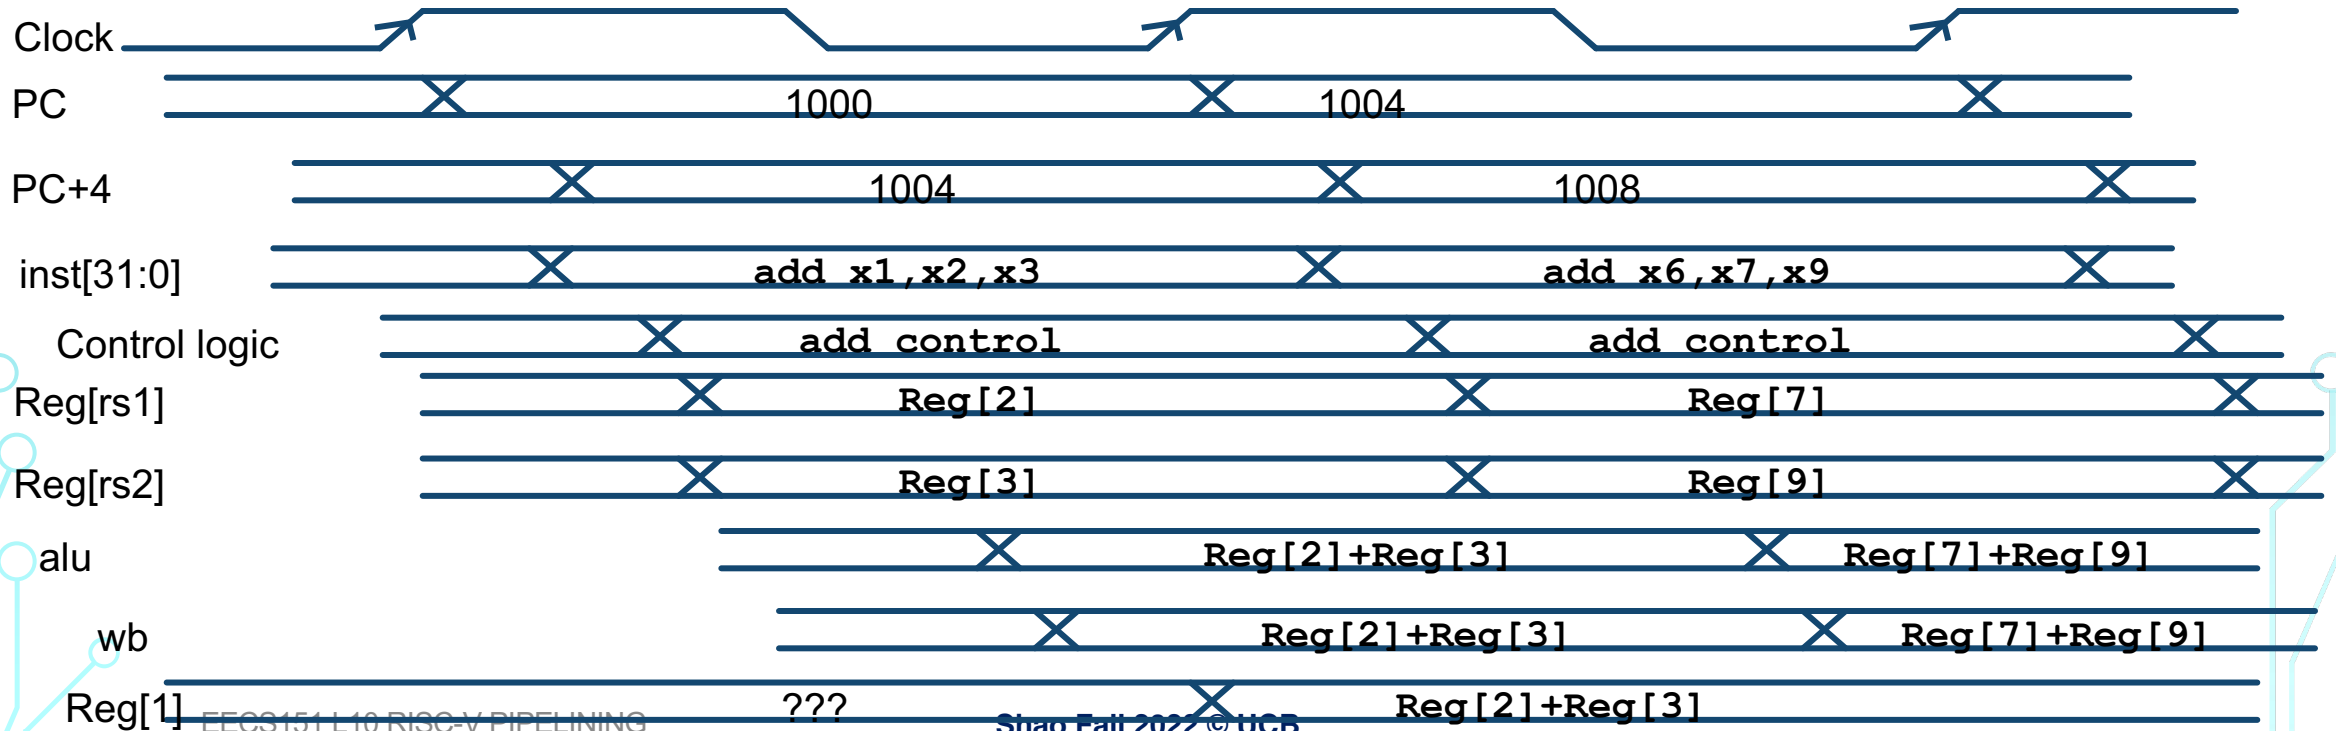
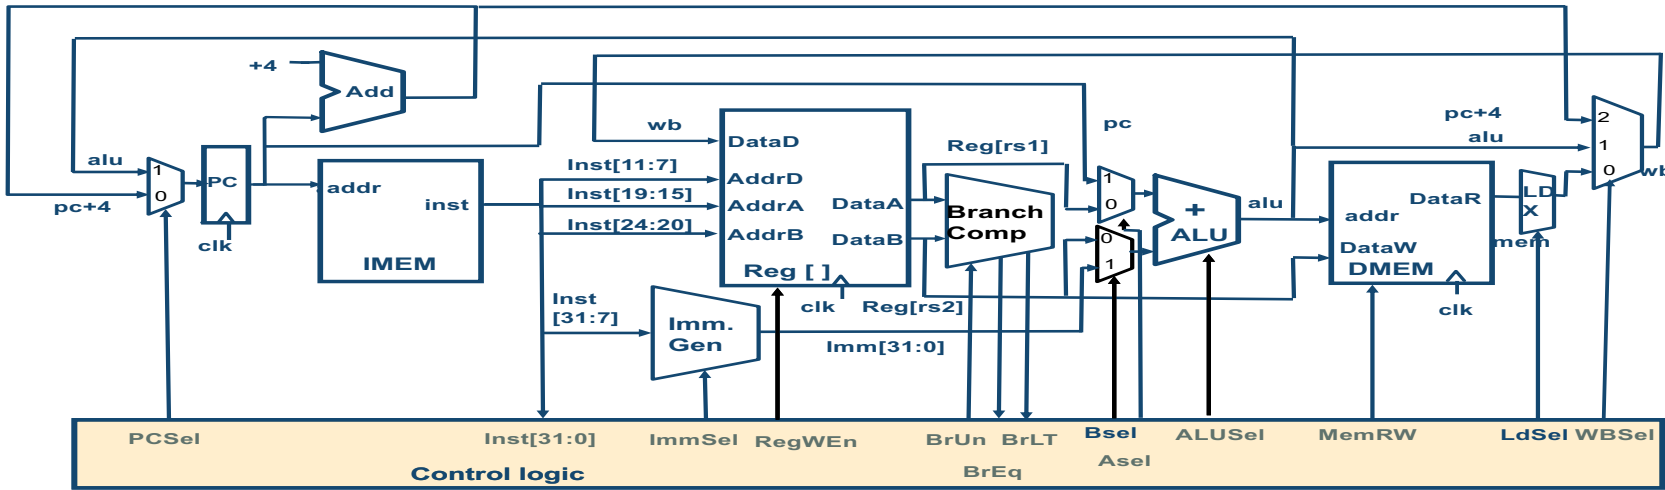
# Complete RV32I Datapath with Control



# Example: add



# add Execution



# Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
<i>(R-R Op)</i>	*	*	+4	*	*	Reg	Reg	<i>(Op)</i>	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auipc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU



# RV32I, a nine-bit ISA!

imm[31:12]				rd	011011	LUI
imm[31:12]				rd	001011	AUIPC
imm[20:10:11:19:12]				rd	110111	JAL
imm[11:0]				rd	110011	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	110001	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	110001	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	110001	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	110001	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	110001	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	110001	BGEU
imm[11:0]				rd	000001	LB
imm[11:0]				rd	000001	LH
imm[11:0]				rd	000001	LW
imm[11:0]				rd	000001	LBU
imm[11:0]				rd	000001	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	010001	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	010001	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	010001	SW
imm[11:0]				rd	001001	ADDI
imm[11:0]				rd	001001	SLTI
imm[11:0]				rd	001001	SLTIU
imm[11:0]				rd	001001	XORI
imm[11:0]				rd	001001	ORI
imm[11:0]				rd	001001	ANDI
0000000	shamt	rs1	001	rd	001001	SLLI
0000000	shamt	rs1	101	rd	001001	SRLI
0100000	shamt	rs1	101	rd	001001	SRAI
0000000	rs2	rs1	000	rd	011001	ADD
0100000	rs2	rs1	000	rd	011001	SUB
0000000	rs2	rs1	001	rd	011001	SLL
0000000	rs2	rs1	010	rd	011001	SLT
0000000	rs2	rs1	011	rd	011001	SLTU
0000000	rs2	rs1	100	rd	011001	XOR
0000000	rs2	rs1	101	rd	011001	SRL
0100000	rs2	rs1	101	rd	011001	SRA
0000000	rs2	rs1	110	rd	011001	OR
0000000	rs2	rs1	111	rd	011001	AND

inst[30]

inst[14:12]

inst[6:2]

**Instruction type encoded using only 9 bits inst[30],inst[14:12], inst[6:2]**



# Control Realization Options

- ROM
  - “Read-Only Memory”
  - Regular structure
  - Can be easily reprogrammed
    - fix errors
    - add instructions
- Combinatorial Logic
  - Decoder is typically hierarchical
    - First decode opcode, and figure out instruction type
    - E.g. branches are  $\text{Inst}[6:2] = 11000$
    - Then determine the actual instruction
      - $\text{Inst}[30] + \text{Inst}[14:12]$
  - Modularity helps simplify and speed up logic
    - Narrow problem space for logic synthesis

# Combinational Logic Control

- Simplest example: BrUn

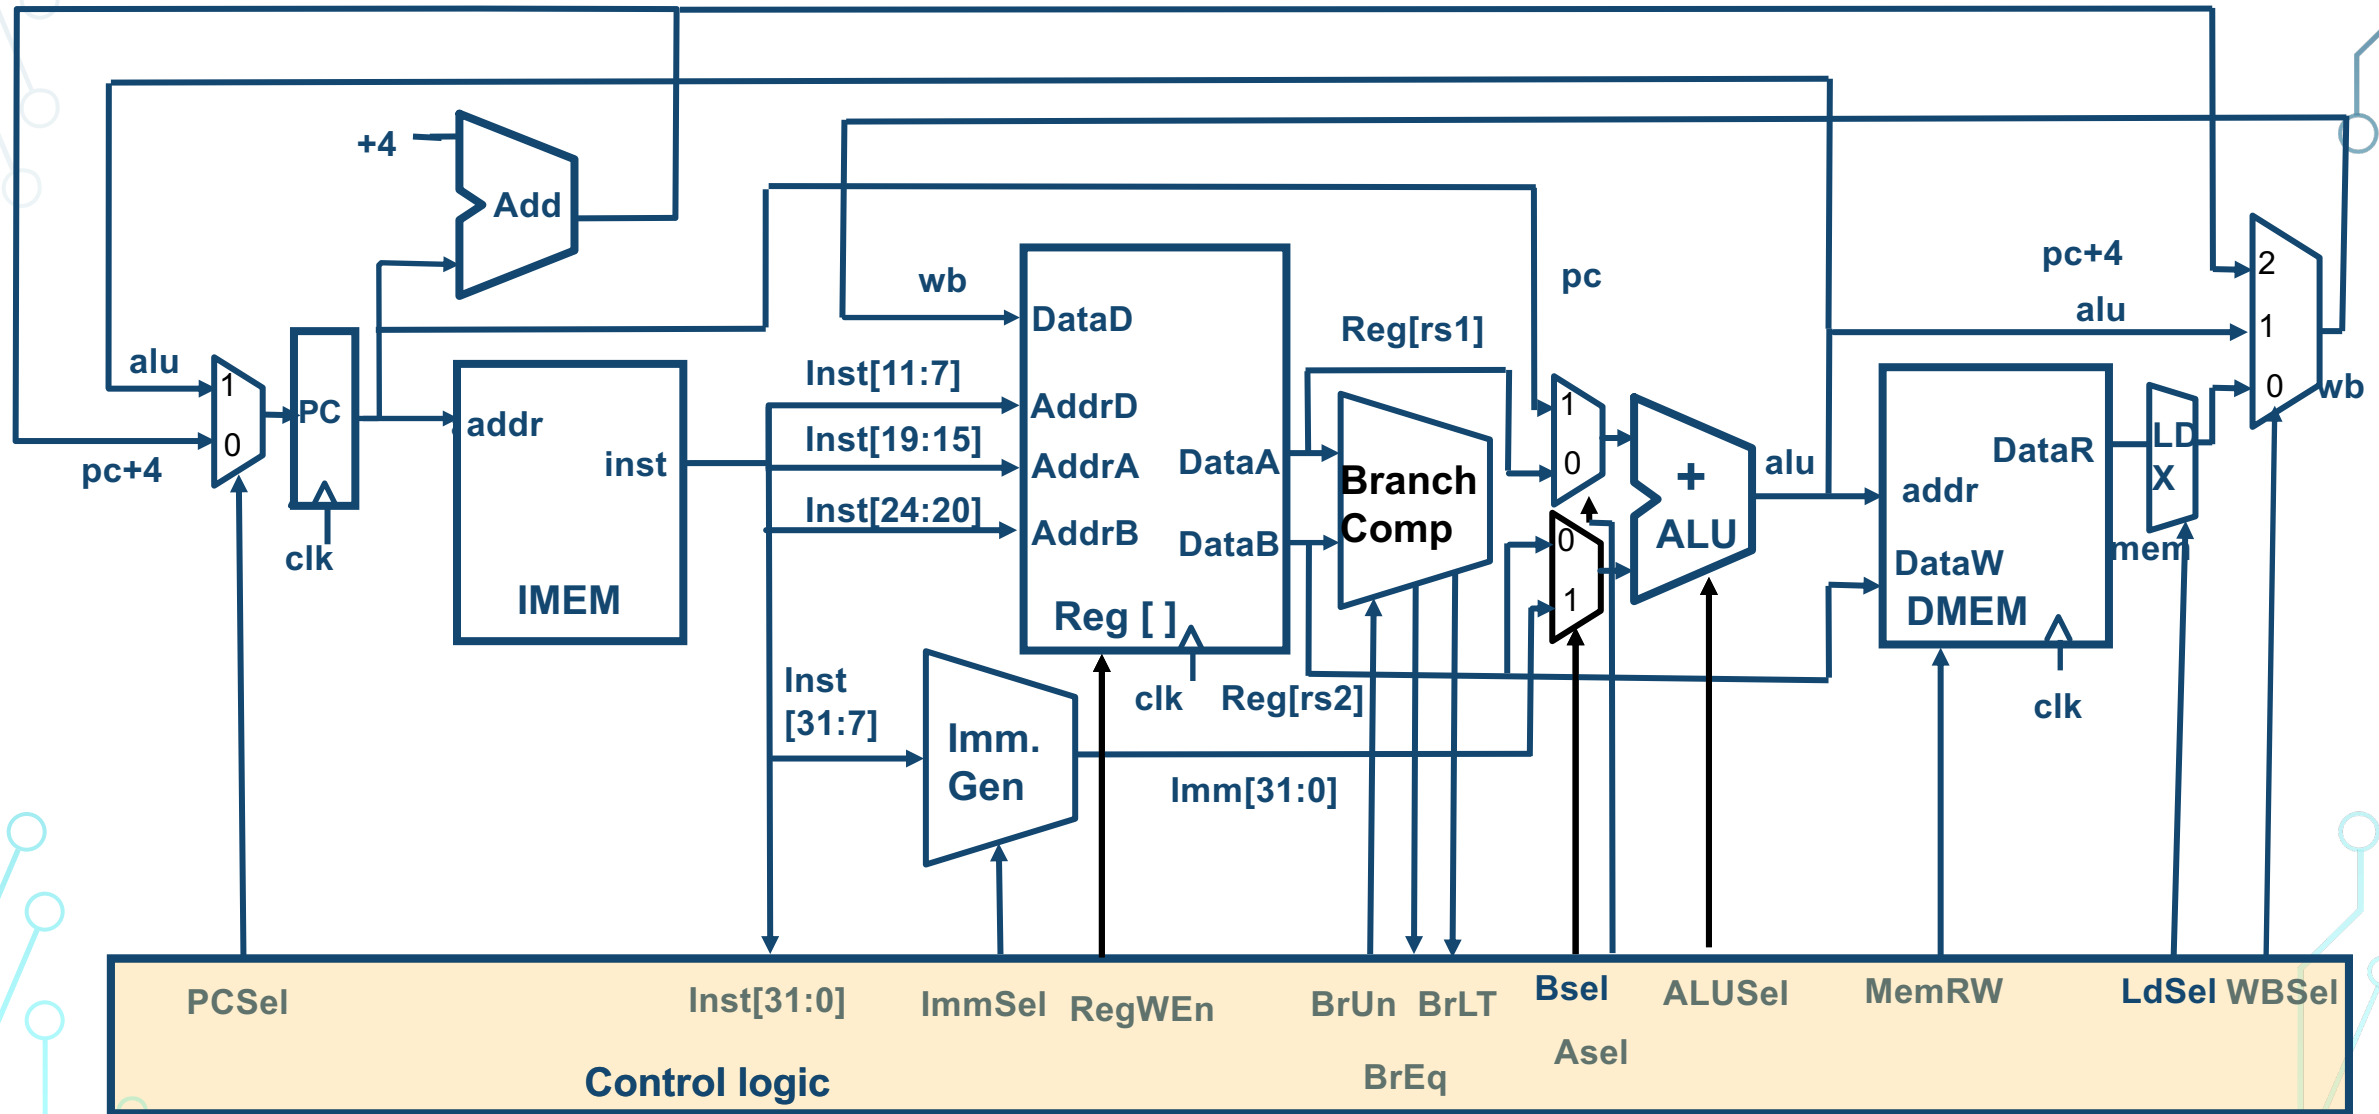
inst[14:12]				inst[6:2]		
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	BGEU

		inst[14:13]			
inst[12]		00	01	11	10
0					
1					

## • How to decode whether BrUn is 1?

- **Branch = Inst[6] • Inst[5] • !Inst[4] • !Inst[3] • !Inst[2]**
- **BrUn = Inst [13] • Branch**

# Complete RV32I Datapath with Control

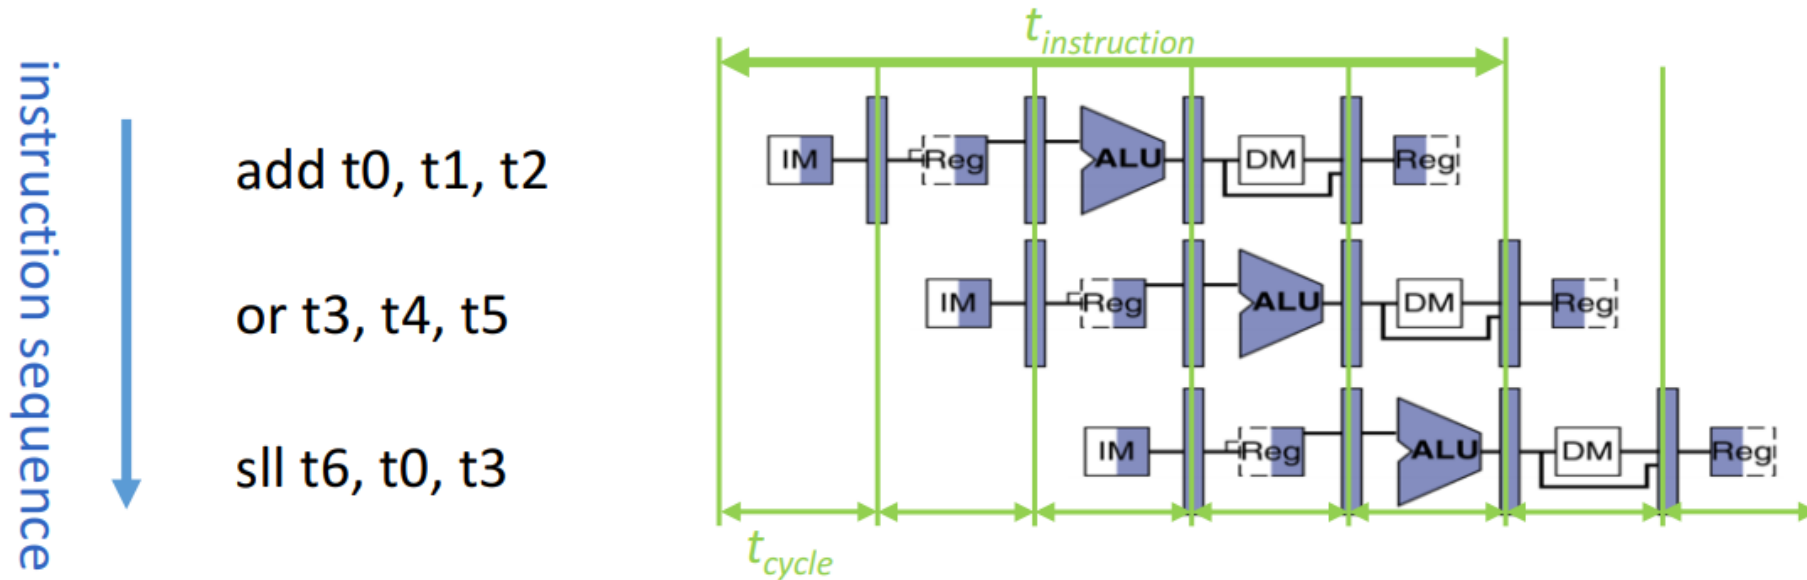




- **RISC-V Pipelining**

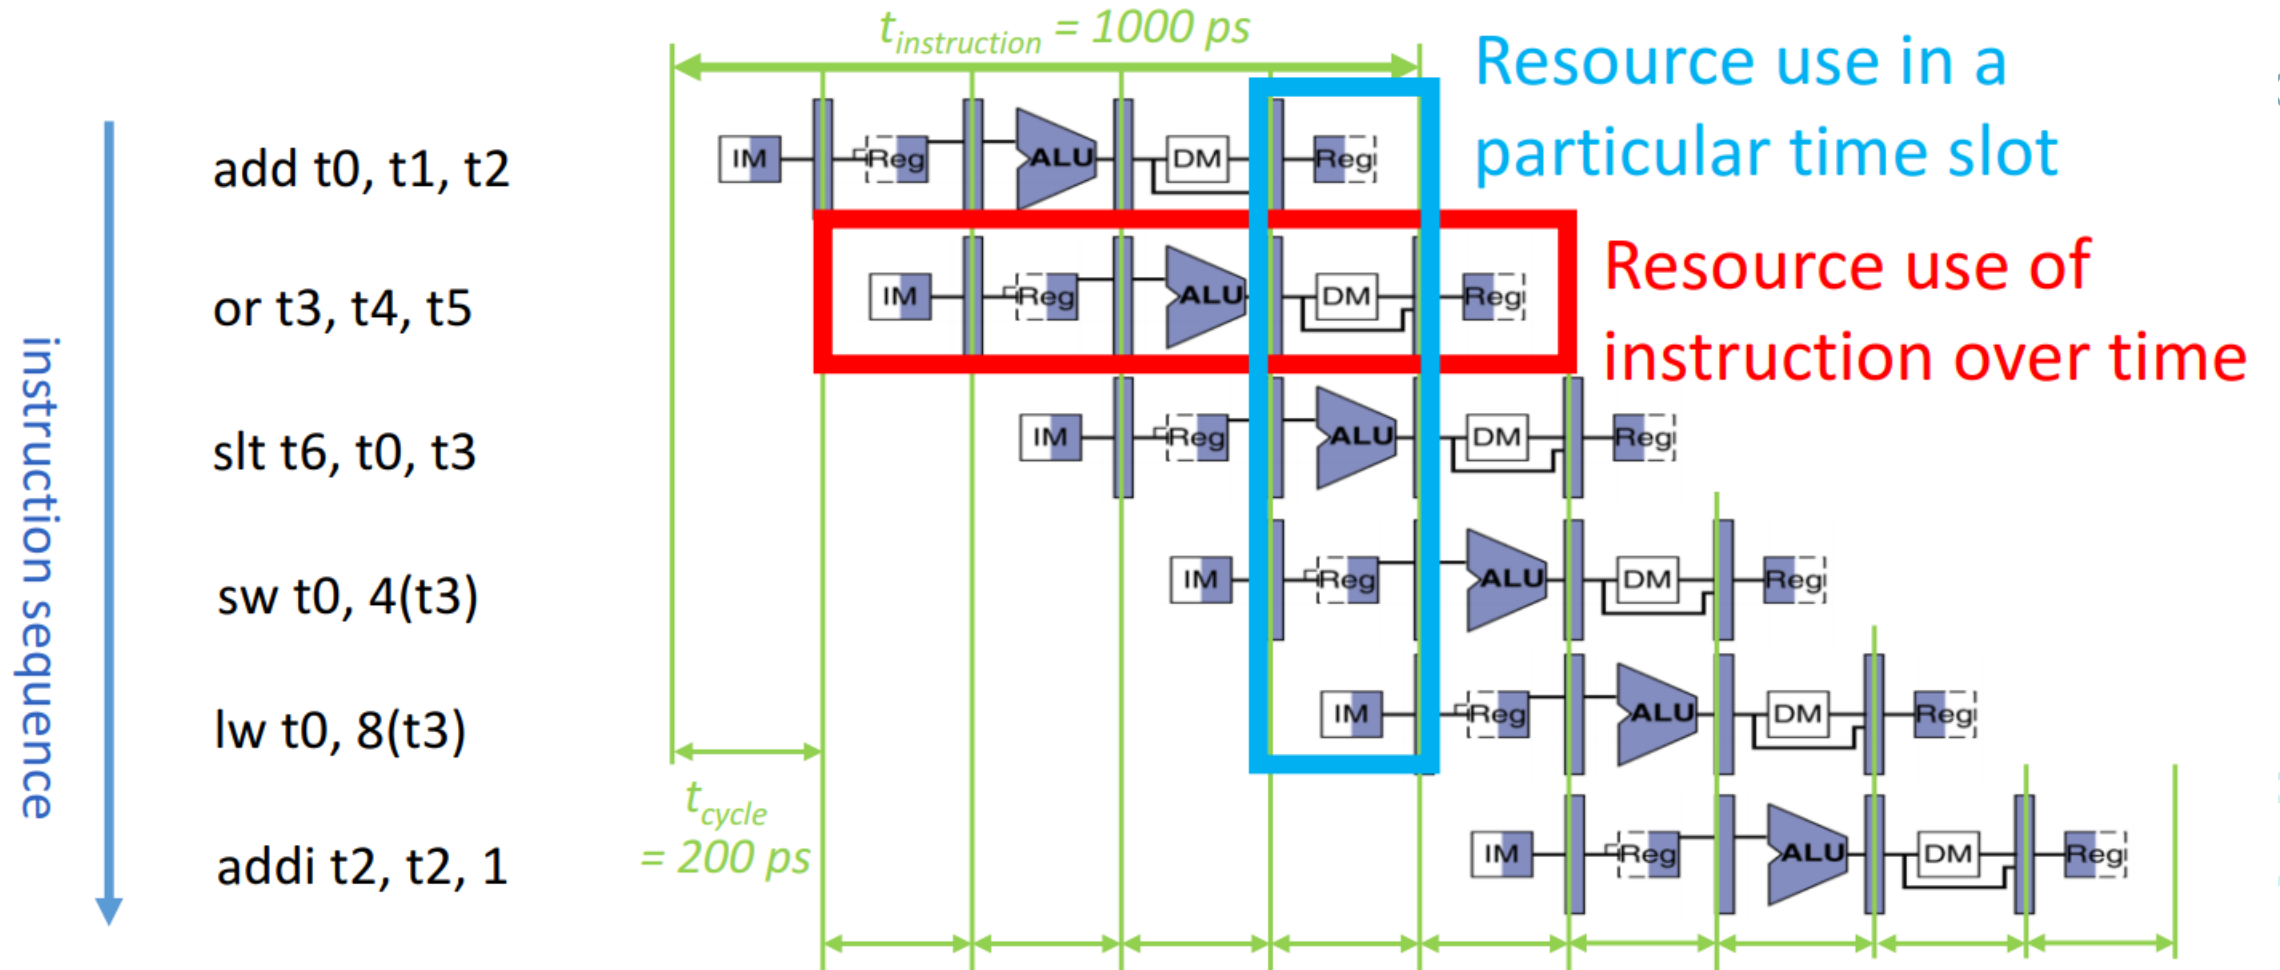
- **5-Stage Pipeline**
- **Pipeline Hazards**
  - **Structural**
  - **Data**
  - **Control**

# Pipelining with RISC-V



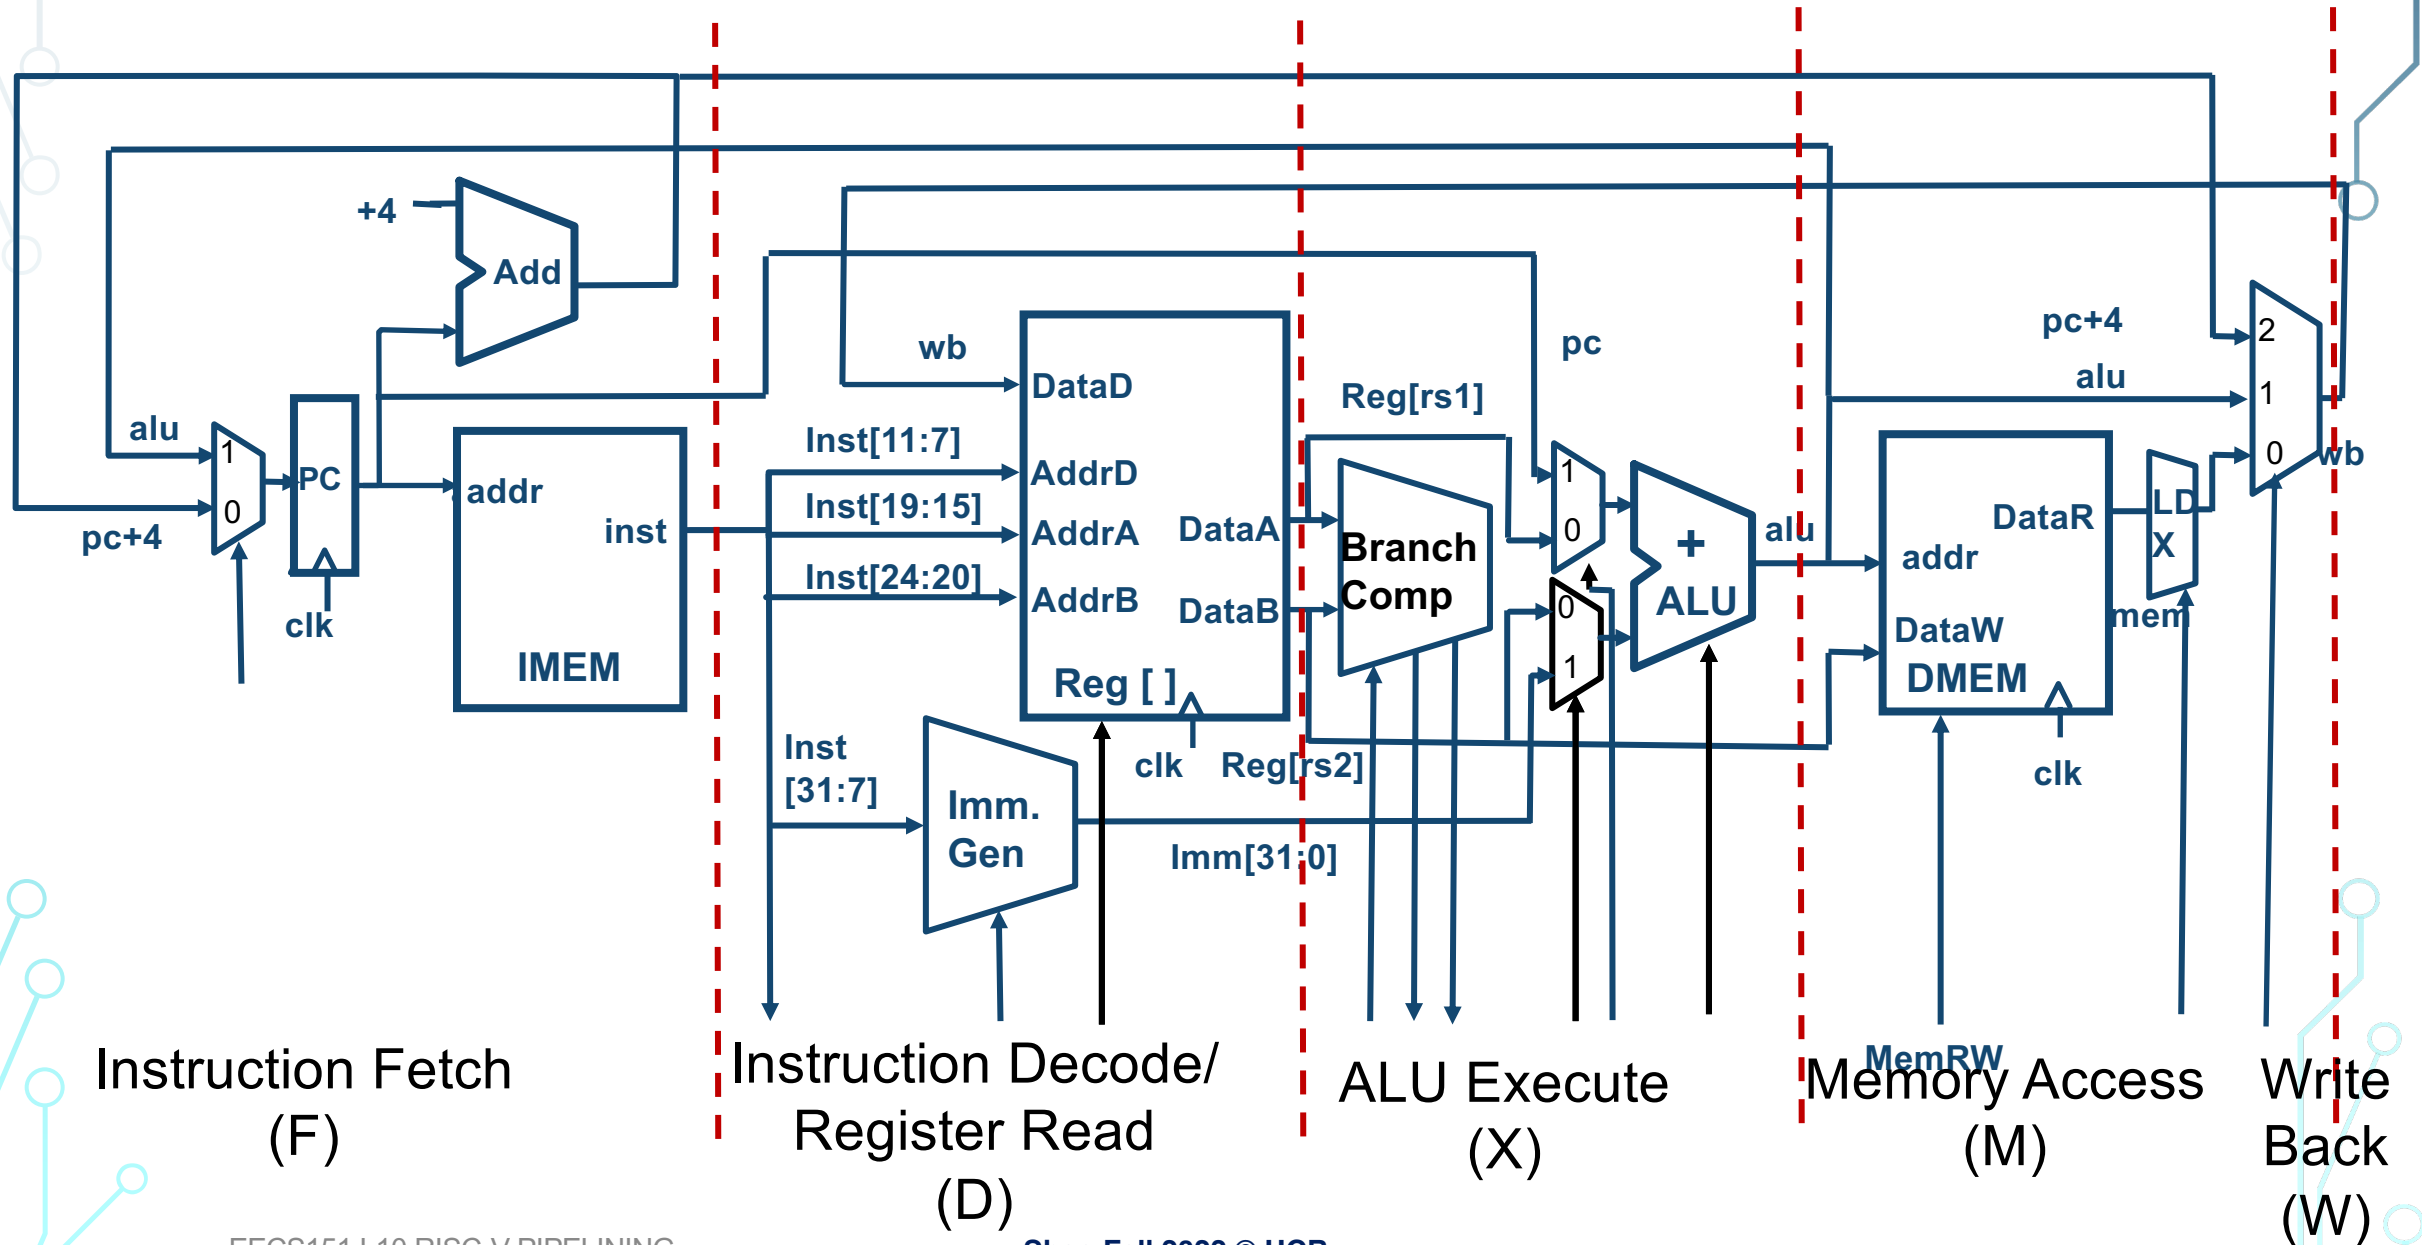
	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, $f_s$	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x

# Pipelining with RISC-V

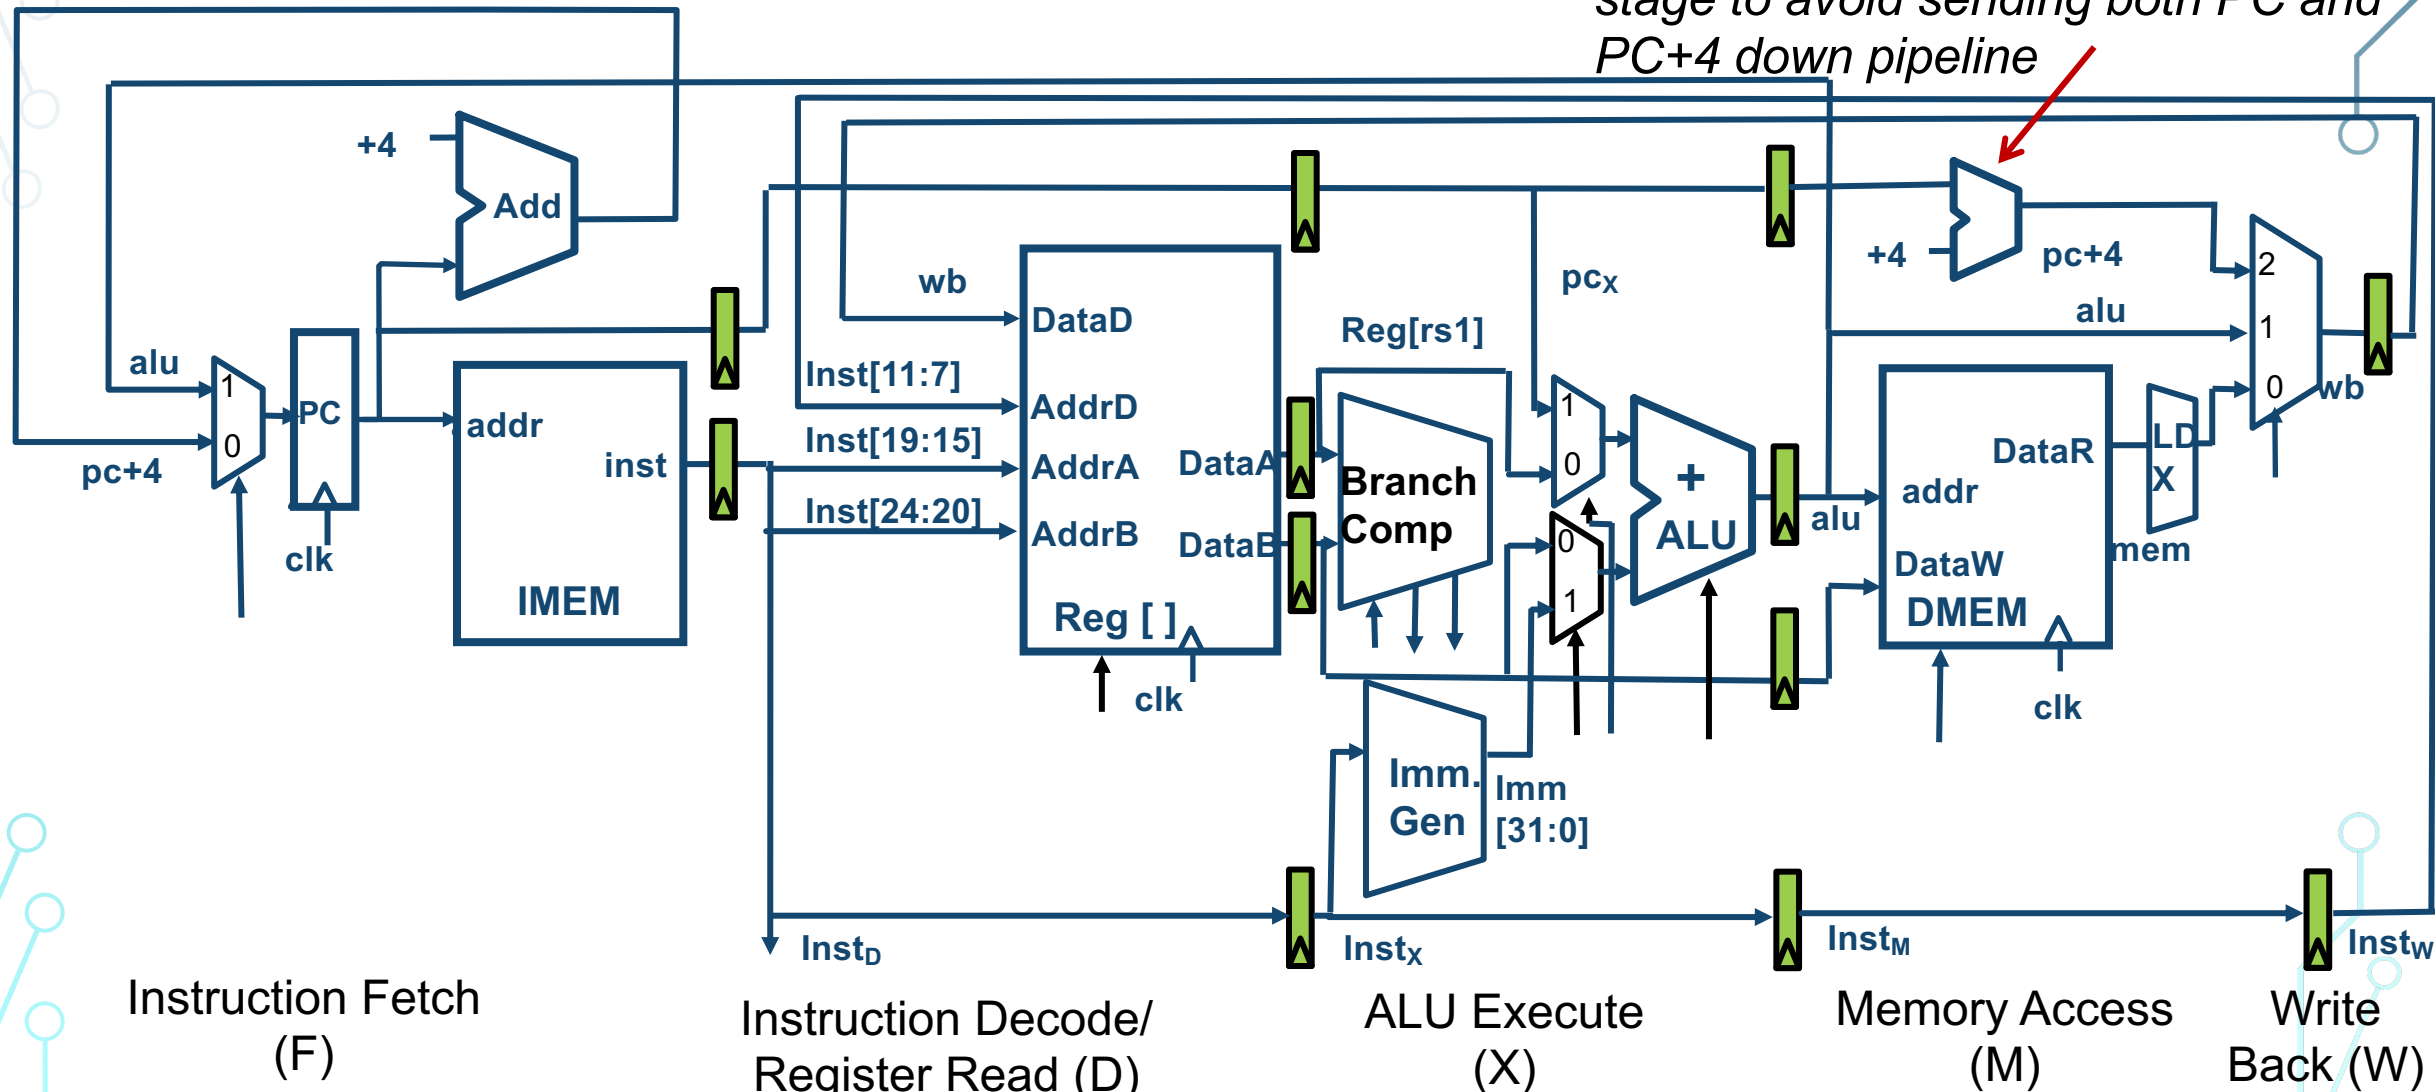




# Complete RV32I Datapath with Control

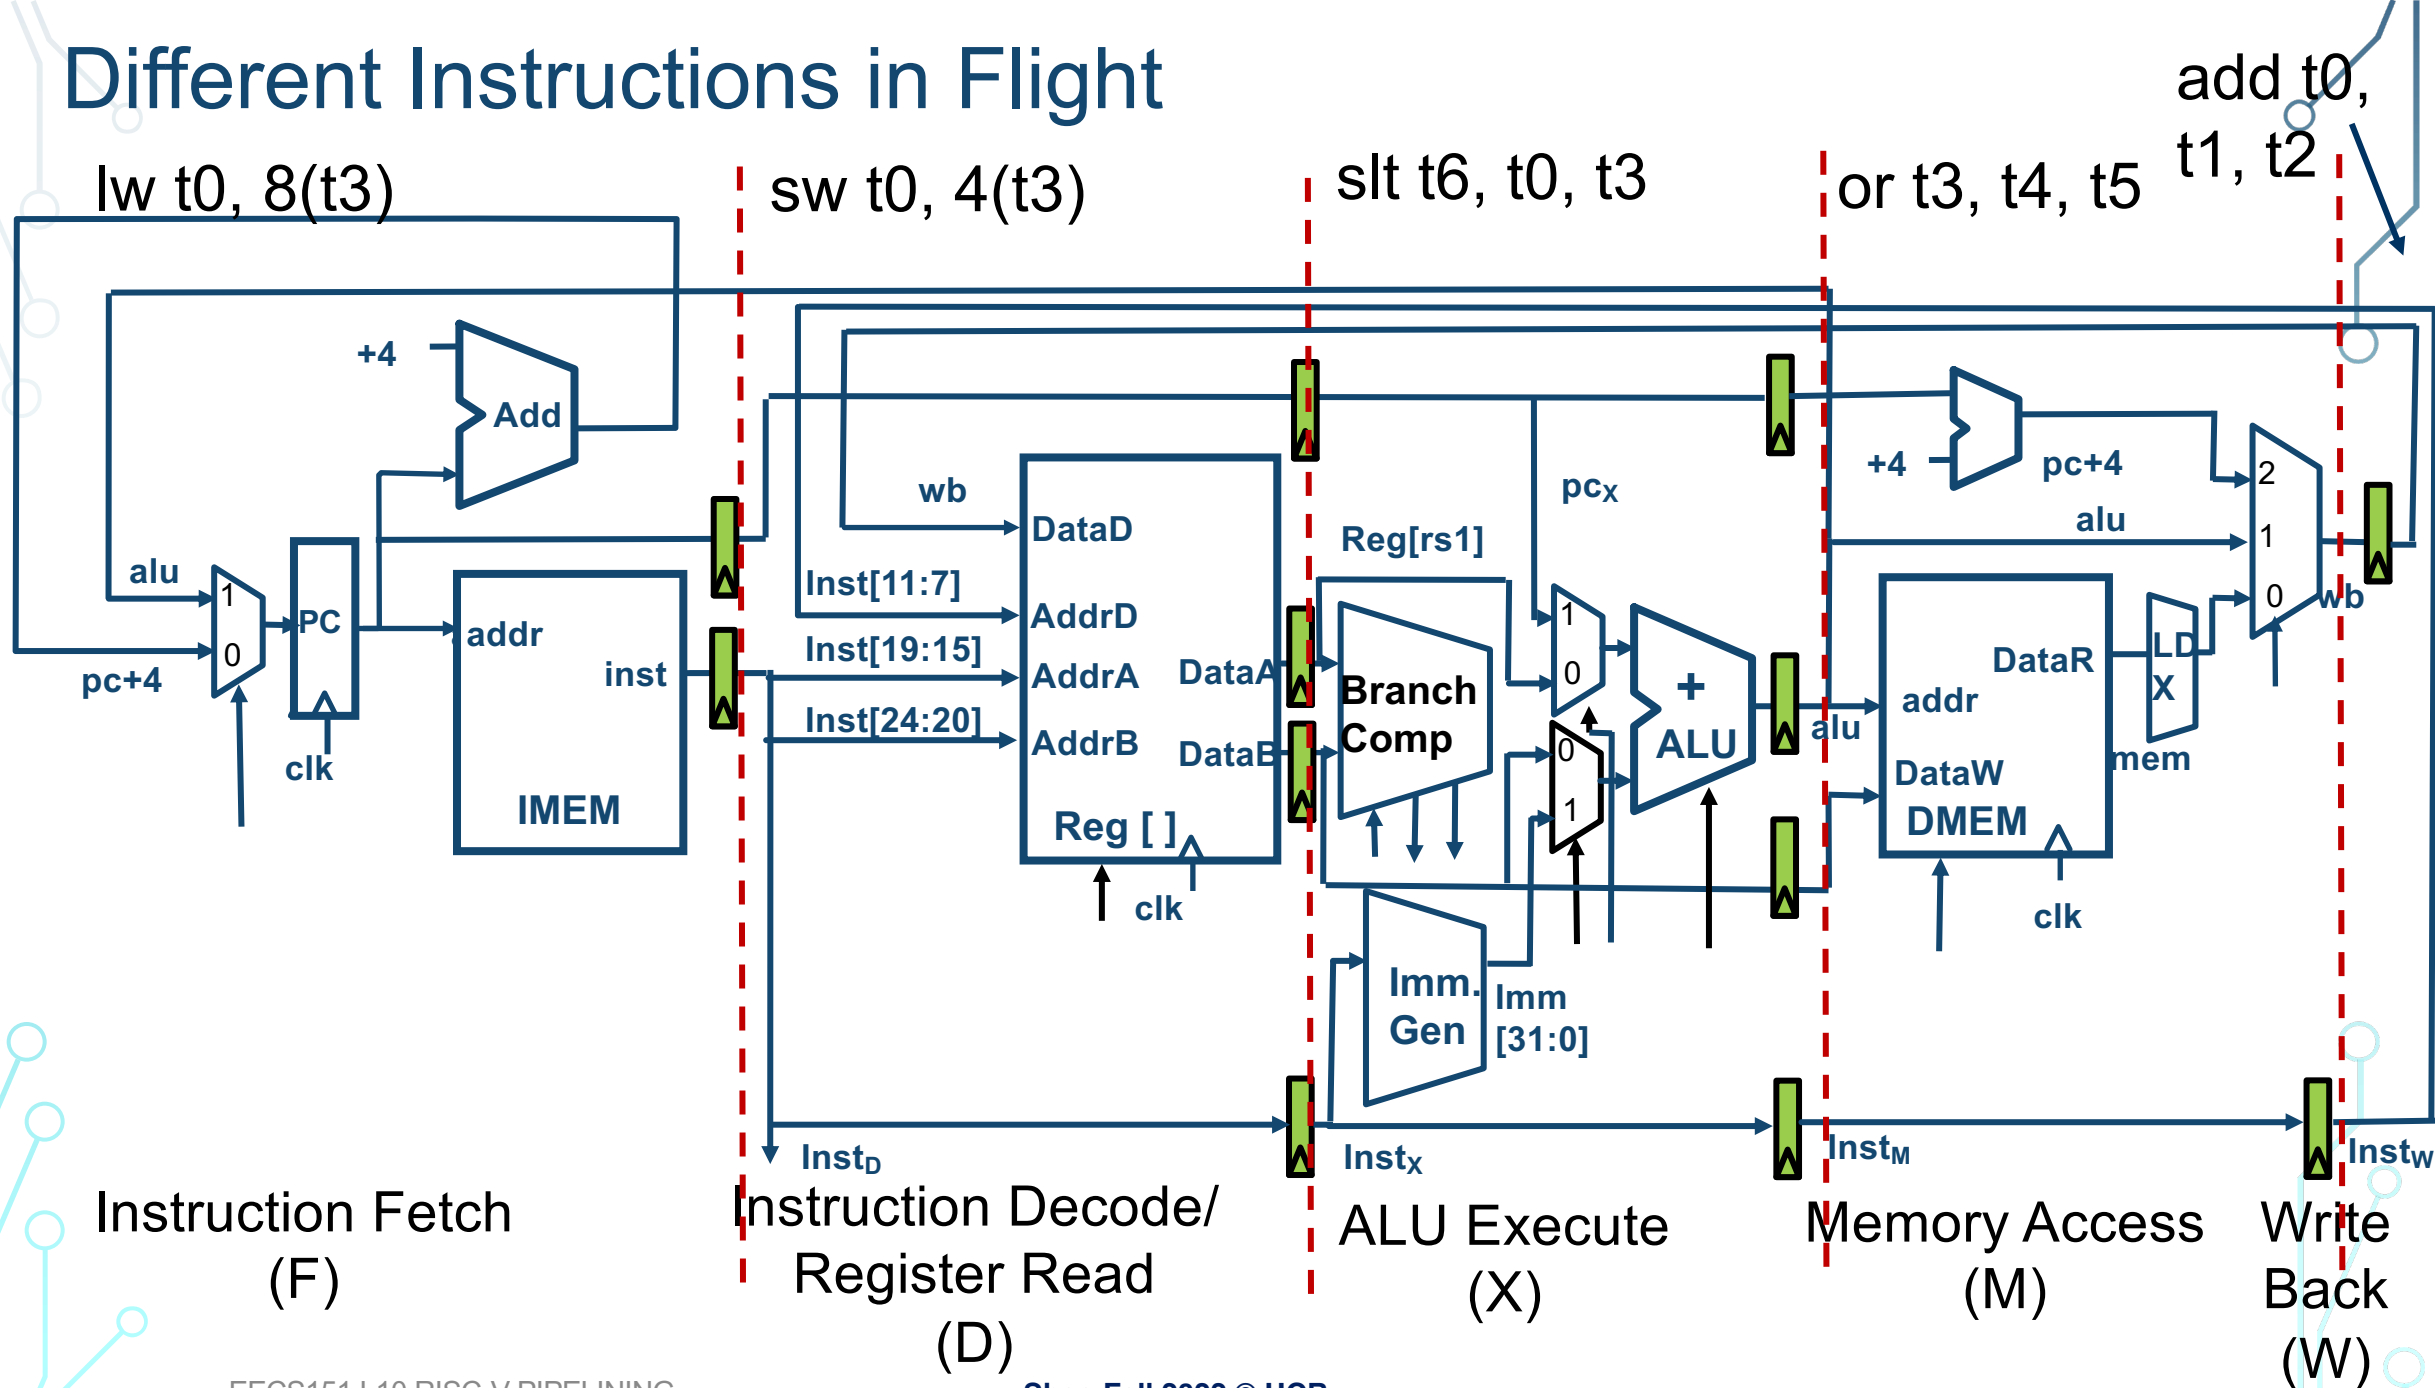


# Pipelining RV32I Datapath




*Must pipeline instruction along with data, so control operates correctly in each stage*

# Different Instructions in Flight



# Pipelined Control

- Control signals derived from instruction
  - As in single-cycle implementation
  - Information is stored in pipeline registers for use by later stages



# Administrivia

- FPGA Guest Lecture next week
- No new lab this week.
- New homework.



- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
  - **Structural**
  - **Data**
  - **Control**



# Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

## 1) *Structural hazard*

- A required resource is busy (e.g. needed in multiple stages)

## 2) *Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

## 3) *Control hazard*

- Flow of execution depends on previous instruction

# Structural Hazard

- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take it in turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
- *Can always solve a structural hazard by adding more hardware*

# Structural Hazards I: Regfile

- Each instruction:
  - can read up to two operands in decode stage
  - can write one value in writeback stage
- Avoid structural hazard by having separate “ports”
  - two independent read ports and one independent write port
- Three accesses per cycle can happen simultaneously
  - We will see a circuit design for multi-ported register files later in the class
- Processors with multiple execution paths have larger number of ports

# Structural Hazard II : Memory Access

instruction sequence


add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



- Instruction and data memory used simultaneously  
✓ Use two separate memories

# Structural Hazards – Summary

- Conflict for use of a resource
- In RISC-V pipeline with a single memory
  - Load/store requires data access
  - Without separate memories, instruction fetch would have to *stall* for that cycle
    - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
  - Or separate instruction/data caches
- RISC ISAs (including RISC-V) designed to avoid structural hazards
  - e.g. at most one memory access/instruction



- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
  - **Structural**
  - **Data**
  - **Control**



# Data Hazard: Register Access

- Separate ports, but what if write to same value as read?
- Does **sw** in the example fetch the old or new value?


add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



# Register Access Policy

instruction sequence

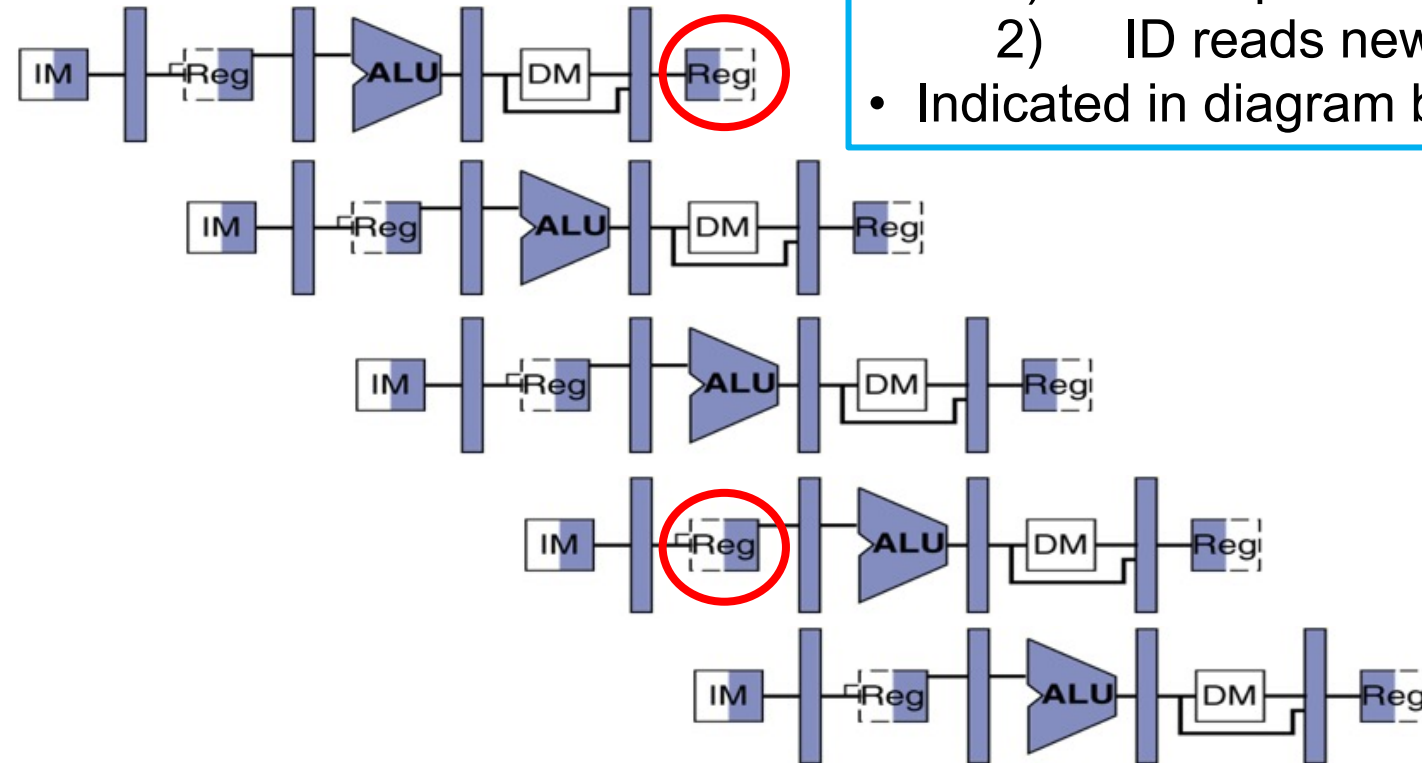
add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



- Exploit high speed of register file (100 ps)
  - 1) WB updates value
  - 2) ID reads new value
- Indicated in diagram by shading

***Might not always be possible to write then read in same cycle, especially in high-frequency designs. Check assumptions in any question.***

# Data Hazard: ALU Result

Value of **s0**

5	5	5	5	5/9	9	9	9	9
---	---	---	---	-----	---	---	---	---


add **s0**, t0, t1

sub t2, **s0**, t0

or t6, **s0**, t3

xor t5, t1, **s0**

sw **s0**, 8(t3)

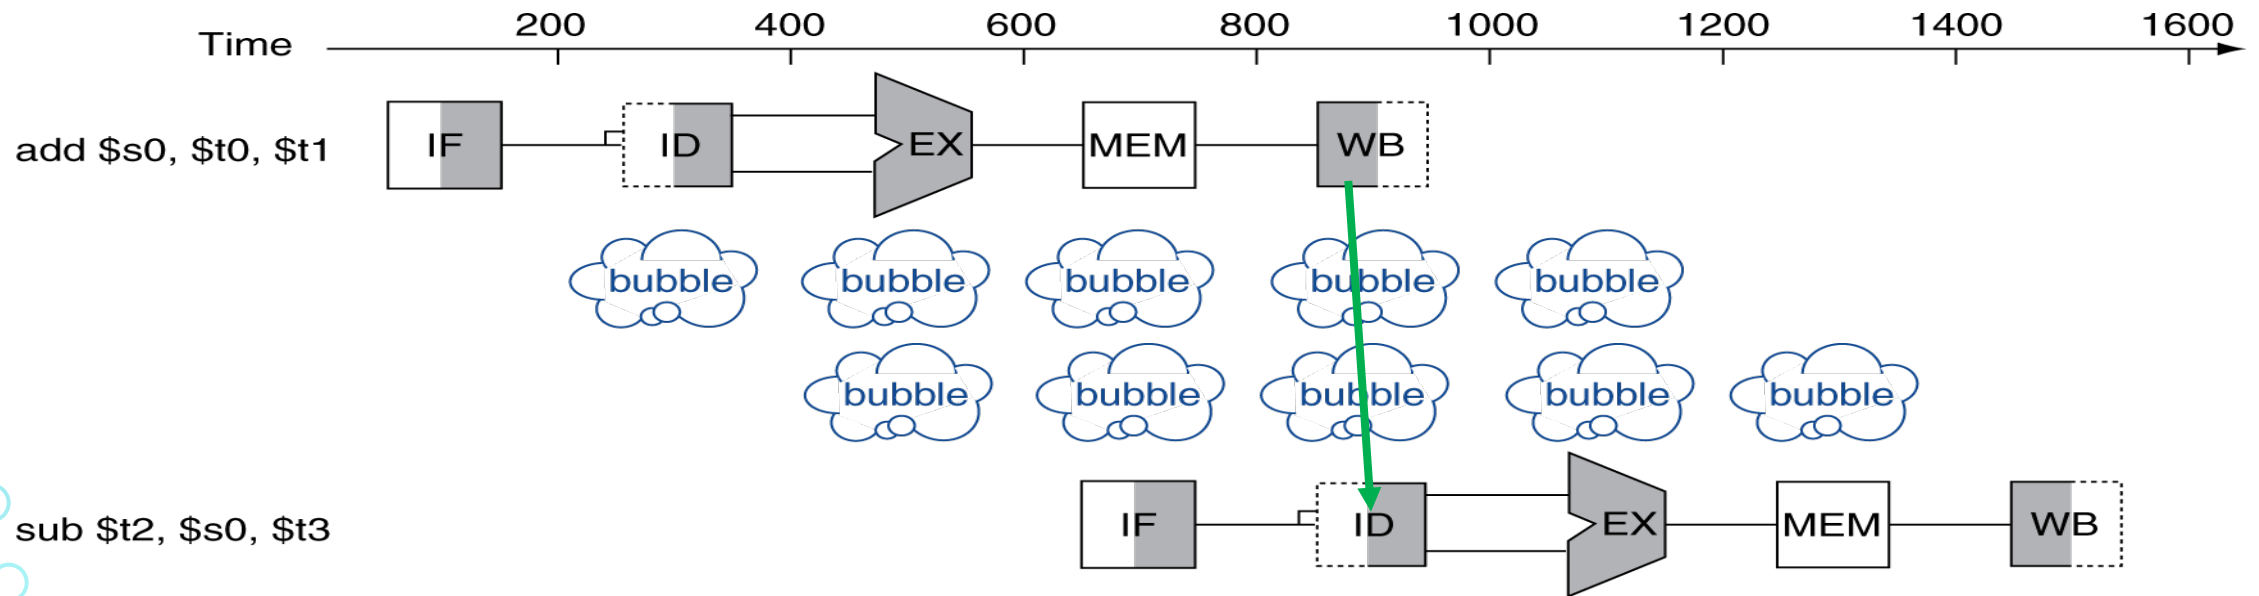


Without some fix, **sub** and **or** will calculate wrong result!

# Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

- add     s0, t0, t1
- sub     t2, s0, t3



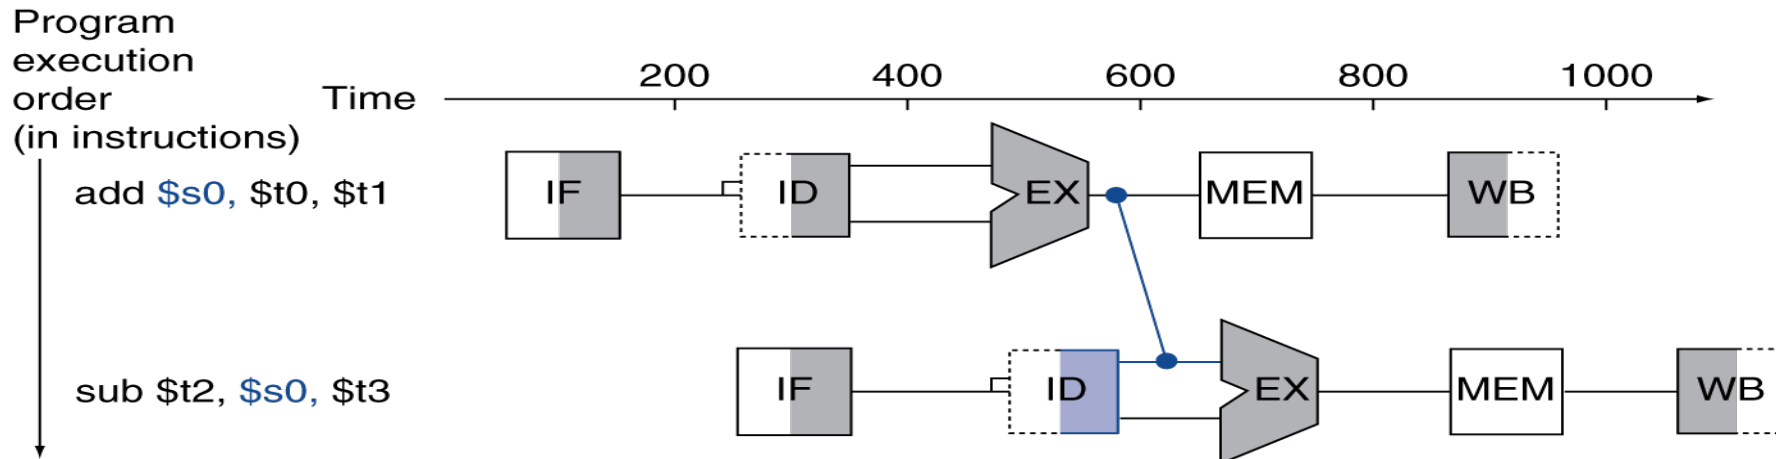
- Bubble:
  - effectively NOP: affected pipeline stages do “nothing”

# Stalls and Performance

- Stalls reduce performance
  - But stalls are required to get correct results
- Compiler can arrange code or insert NOPs (writes to register x0) to avoid hazards and stalls
  - But requires knowledge of the pipeline structure


# Solution 2: Forwarding

- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra connections in the datapath





# Solution 2: Forwarding



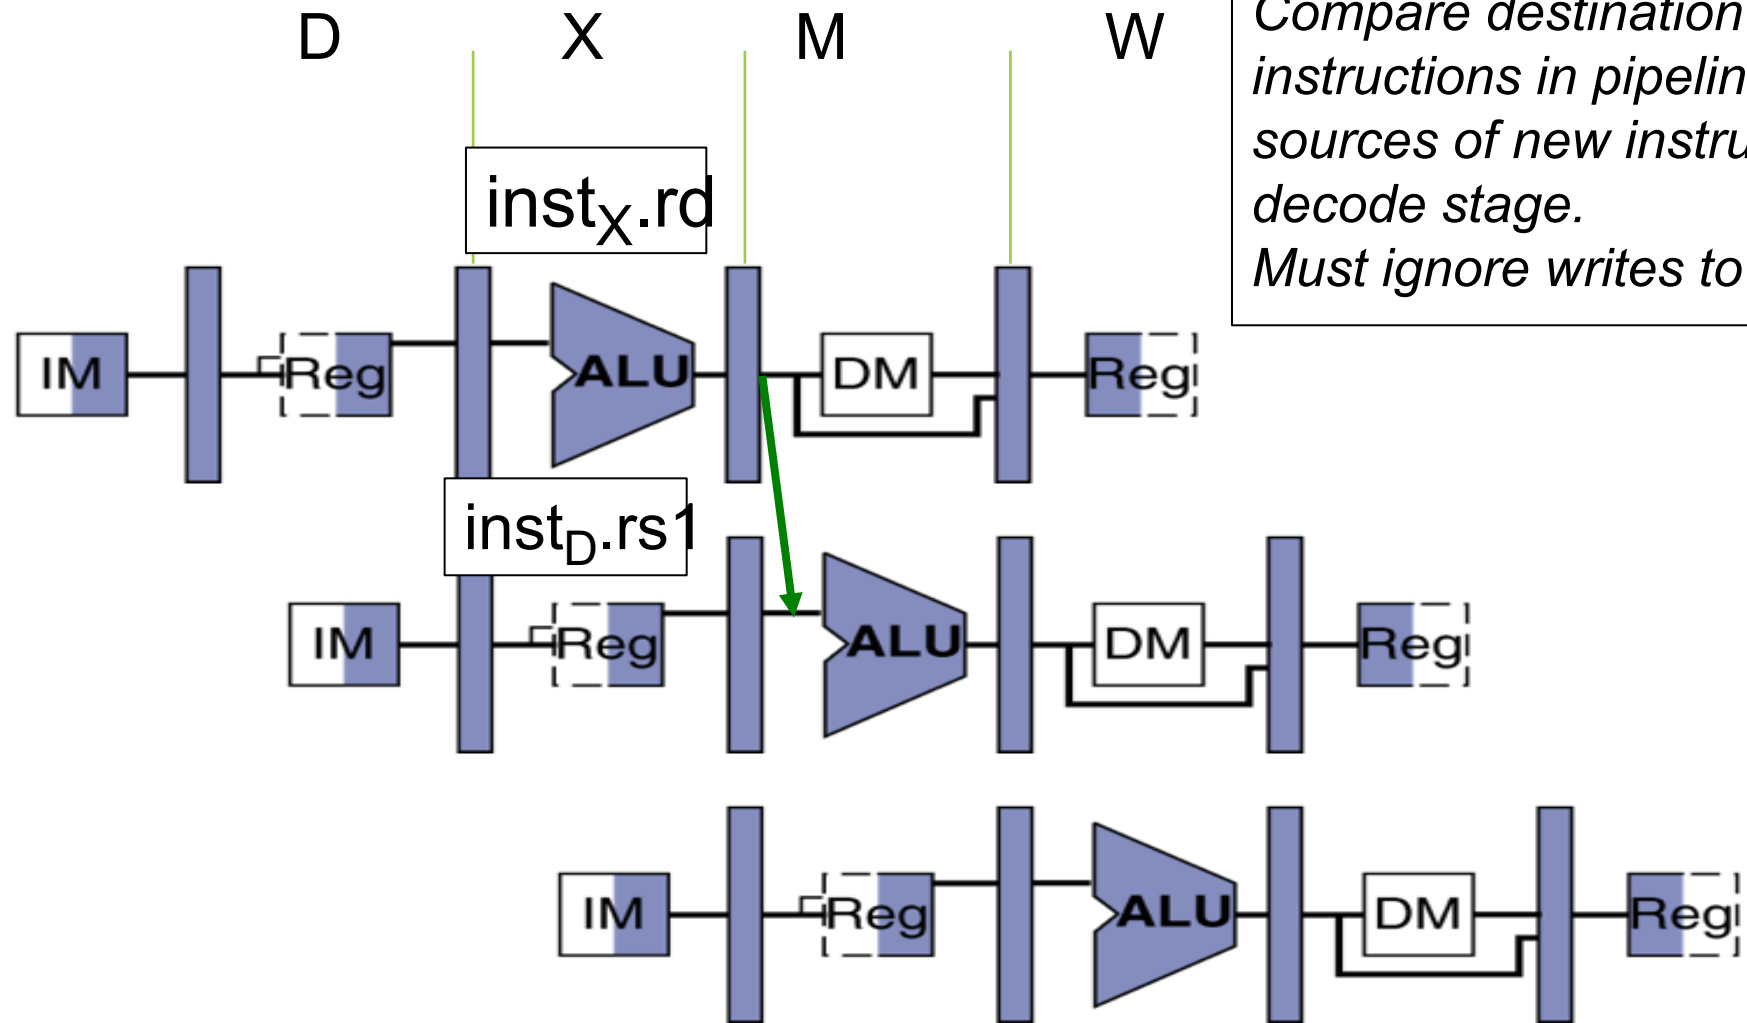
Forwarding: grab operand from pipeline stage,  
rather than register file

# Detect Need for Forwarding (example)

add **t0**, t1, t2

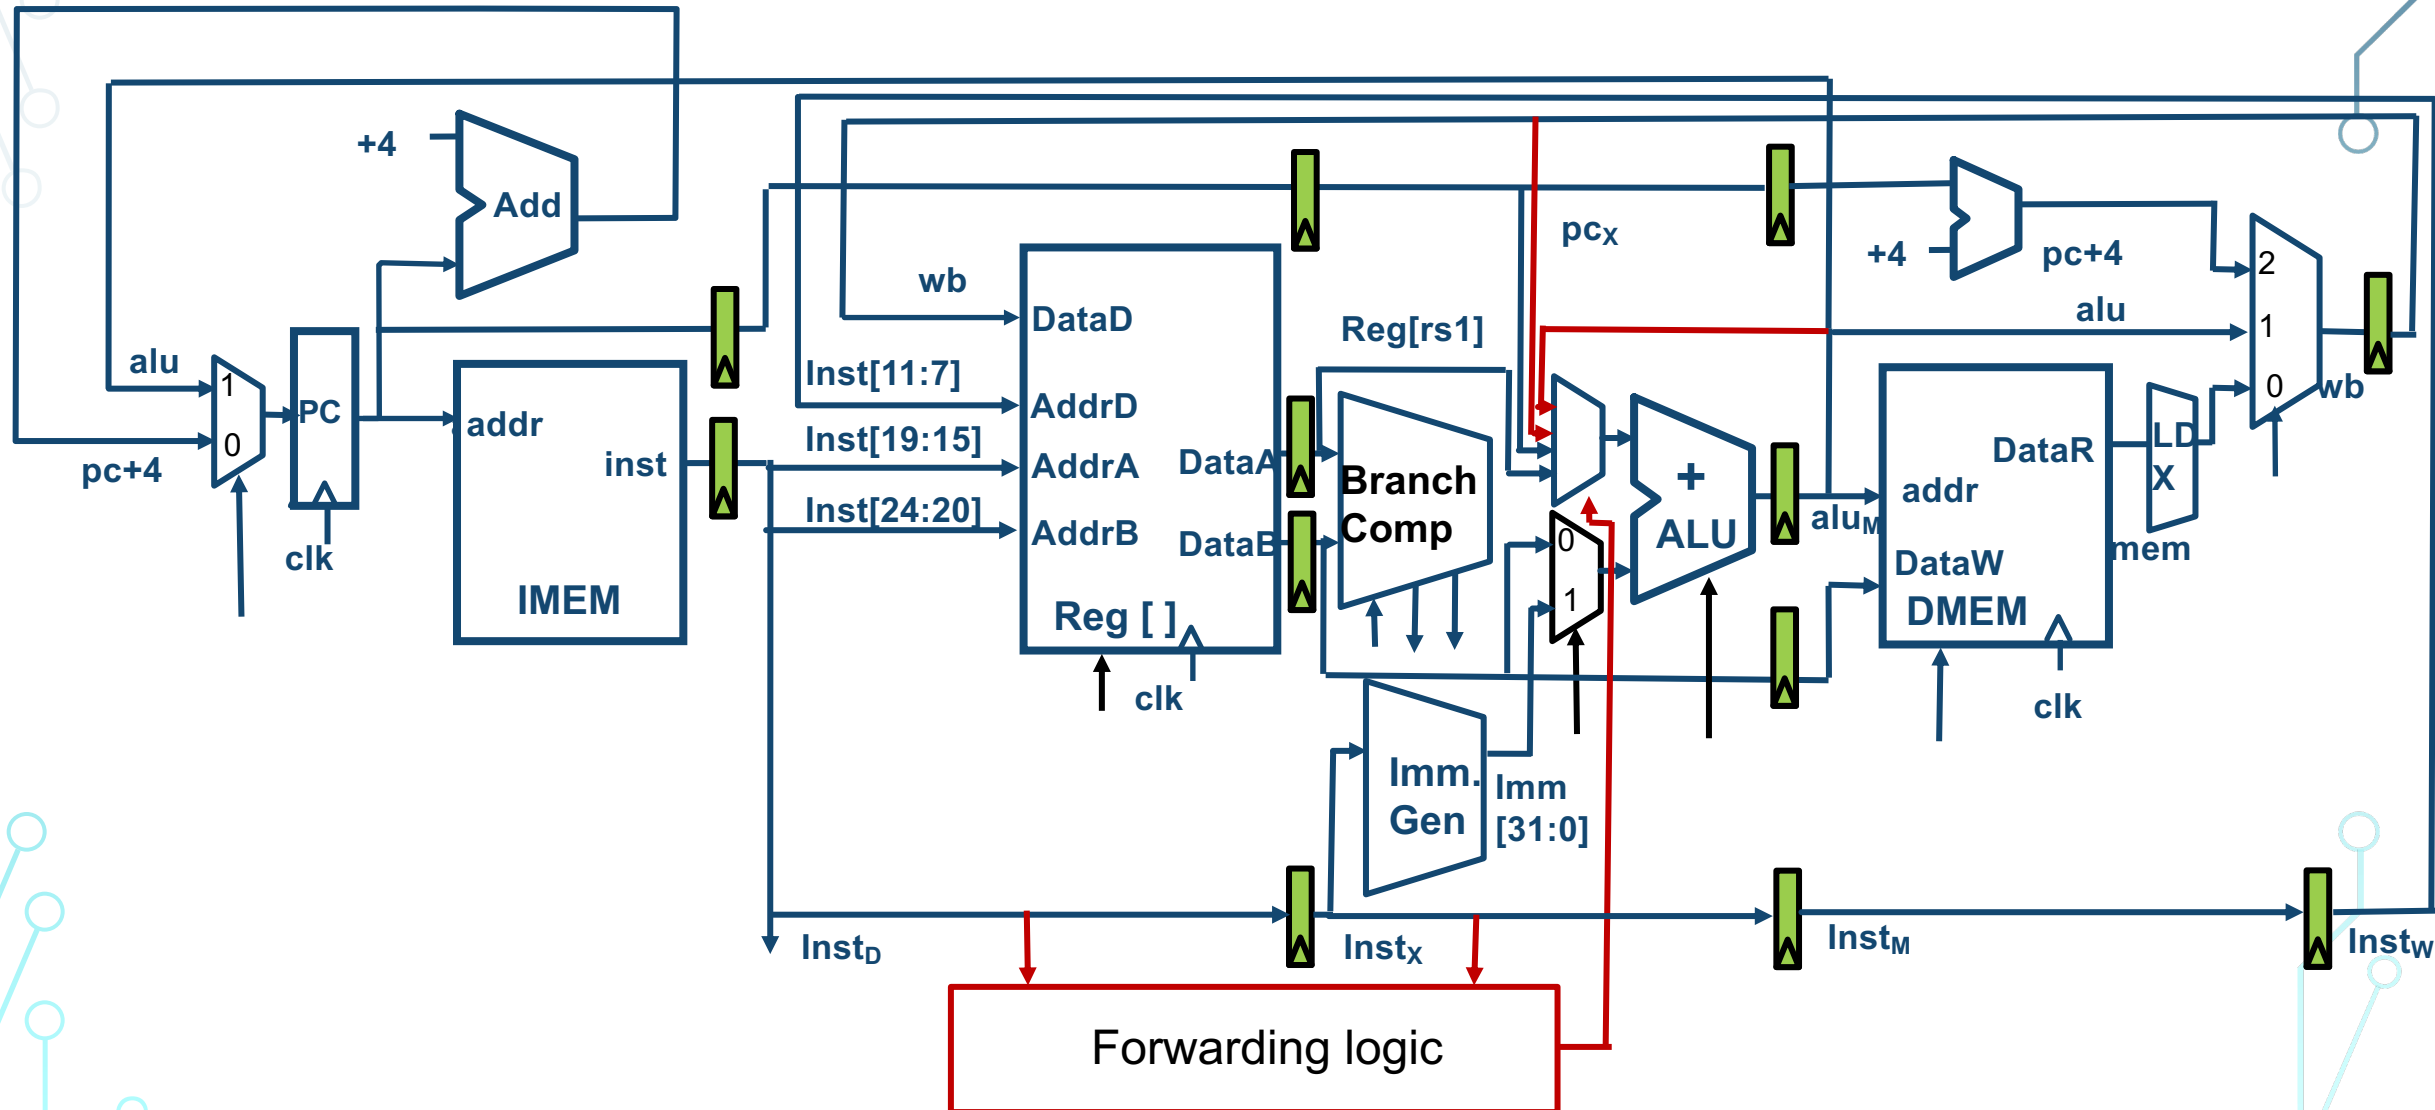
or t3, **t0**, t5

sub t6, **t0**, t3



*Compare destination of older instructions in pipeline with sources of new instruction in decode stage.  
Must ignore writes to x0!*

# ForwardingA Path

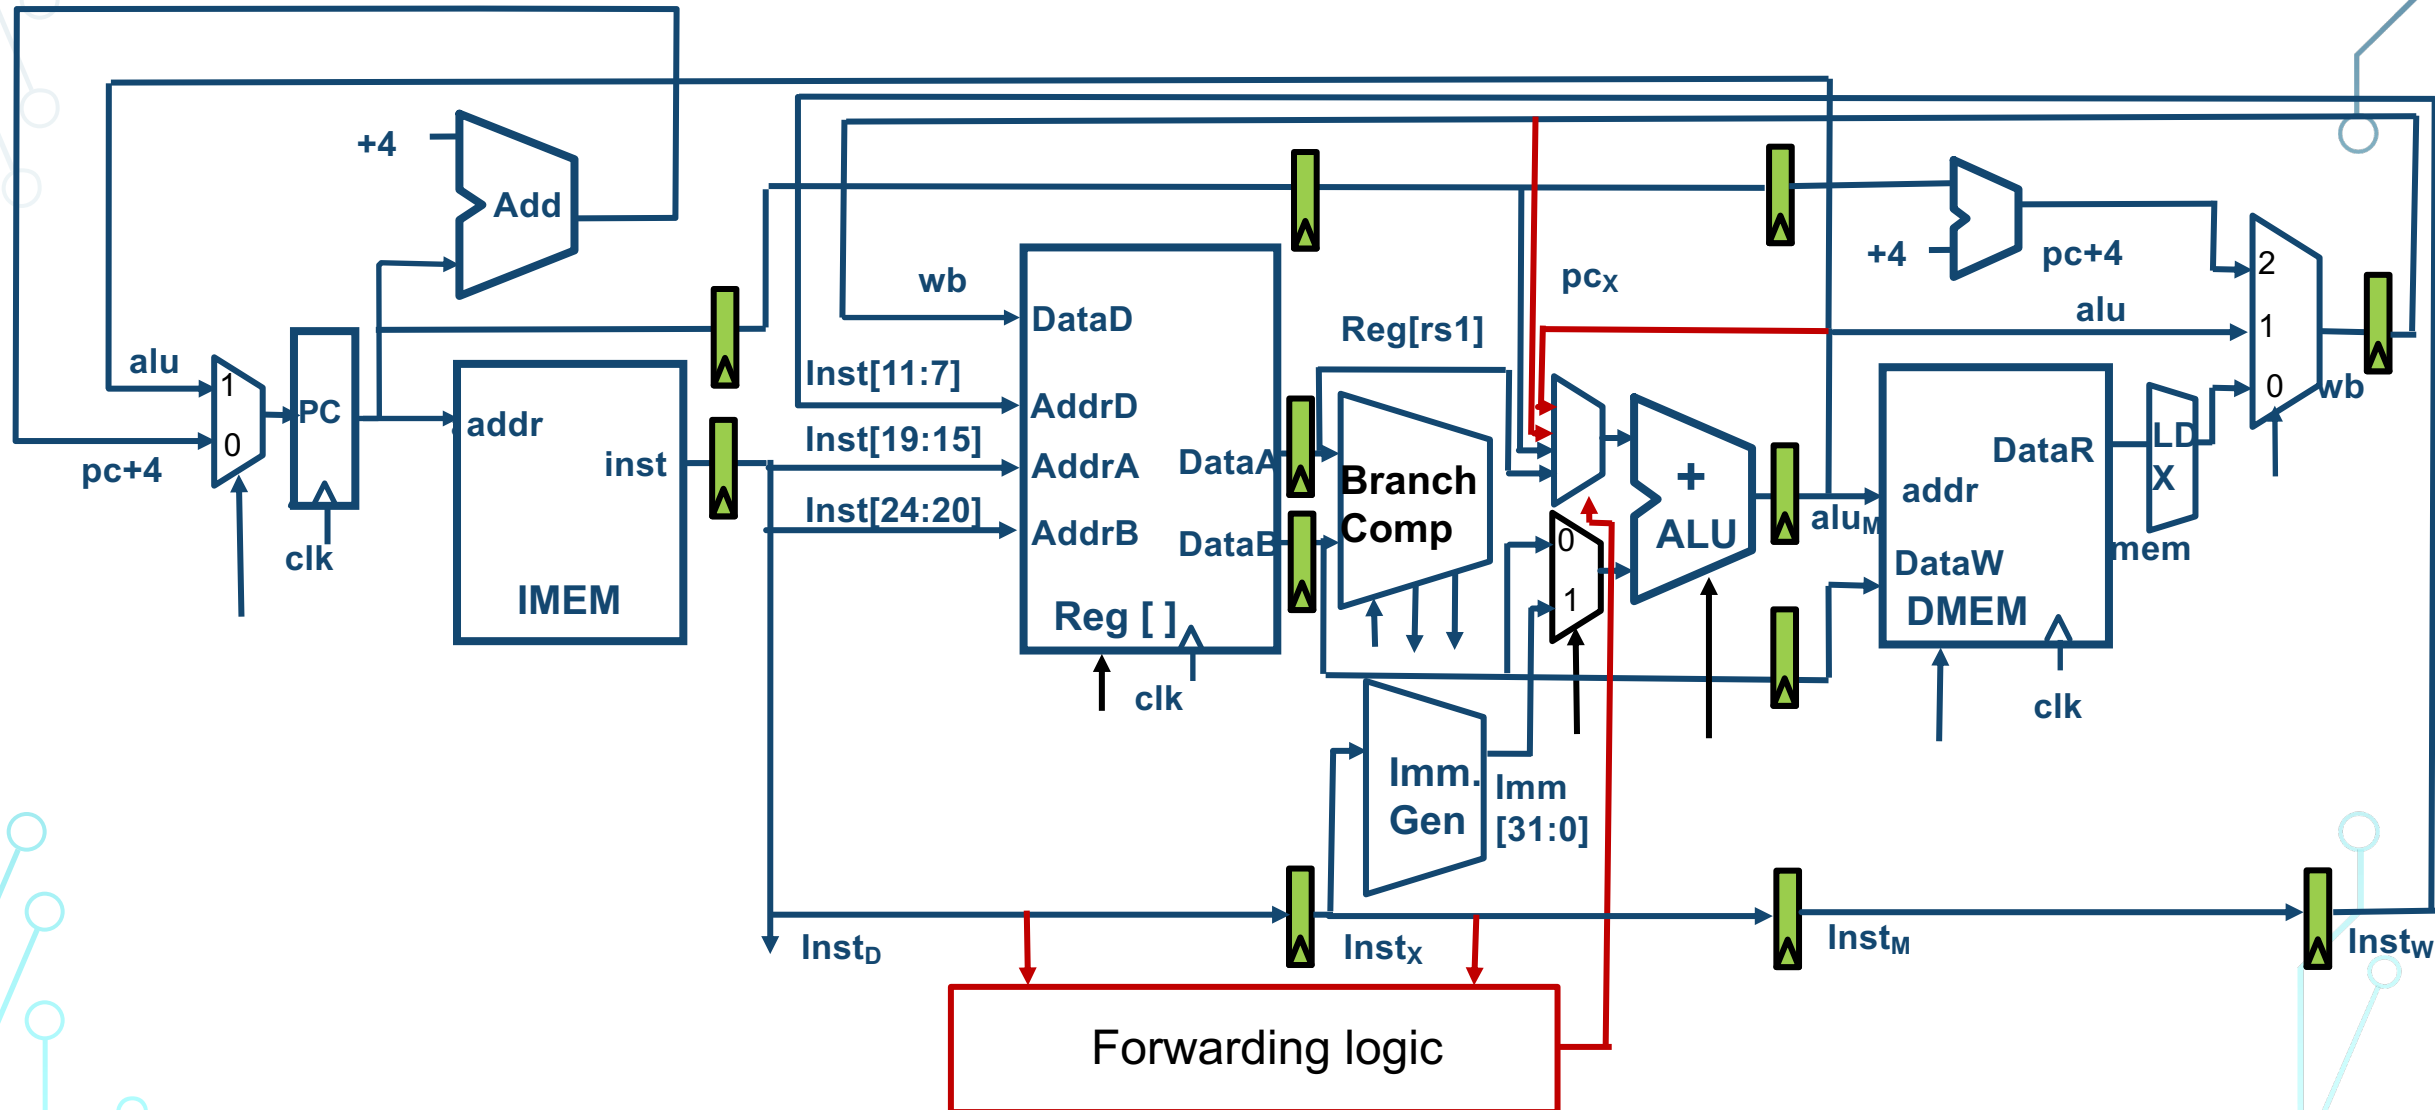





- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
  - **Structural**
  - **Data**
  - **Control**

# Pipelined Datapath: When is the target PC calculated?



# Control Hazards



# Observation

- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs

# Kill Instructions after Branch if Taken


beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

Add t2, t3, s0

label: xxxxxx



Taken branch

Convert to NOP

Convert to NOP

Convert to NOP

PC updated  
reflecting branch  
outcome



# Reducing Branch Penalties

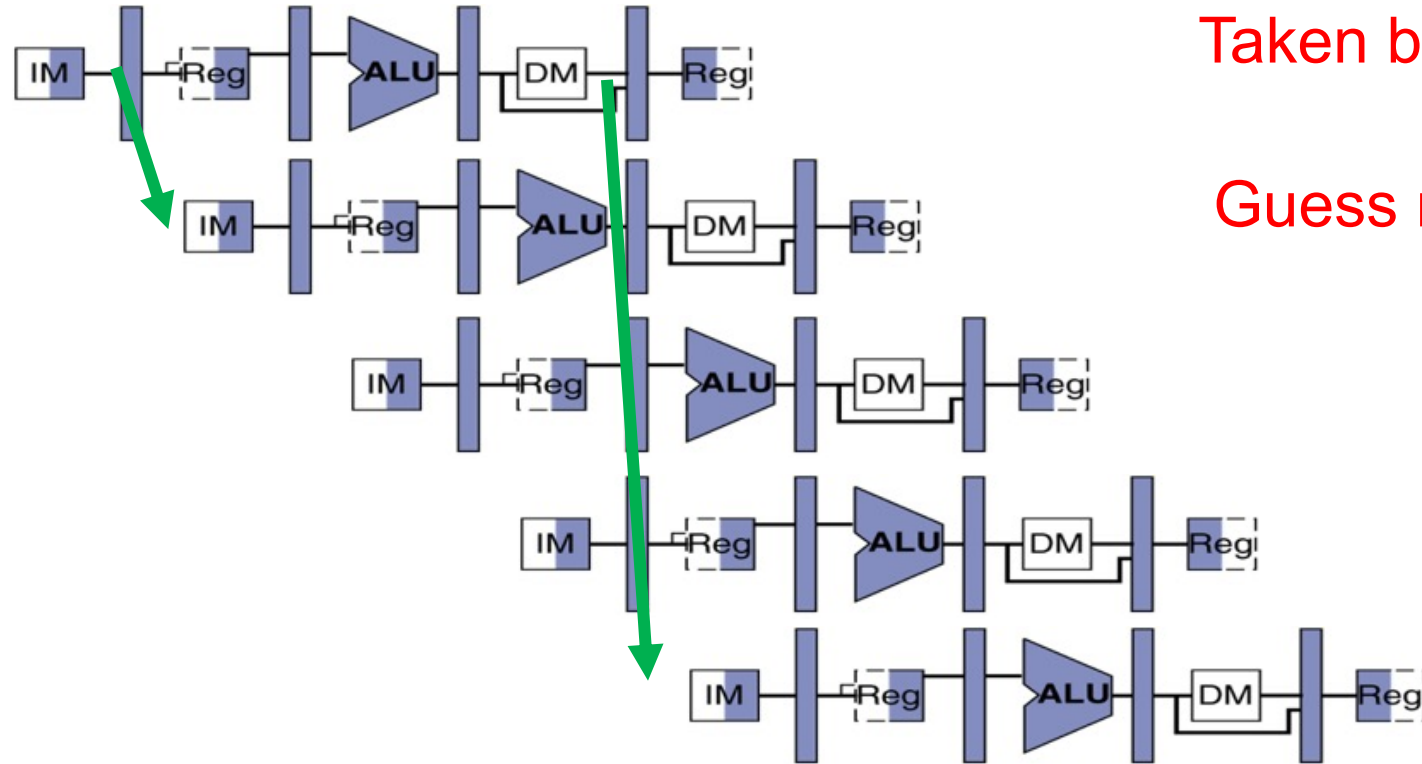
- In our datapath, every taken branch in simple pipeline costs 3 dead cycles
- To improve performance, use “branch prediction” to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect

# Branch Prediction

beq t0, t1, label

label: .....

.....



Taken branch

Guess next PC!

Check guess correct

# Summary

- We have covered single-cycle, multi-cycle and pipelined datapaths
- Pipeline Hazards:
  - Structural
  - Data
  - Control hazards
- Need to be understood and handled appropriately
  - More resources
  - Forwarding
  - Stall