

# EECS251B

# Advanced Digital Circuits and Systems

## Lecture 3 – Design Productivity

Vladimir Stojanović

Tuesdays and Thursdays 9:30-11am

Cory 521

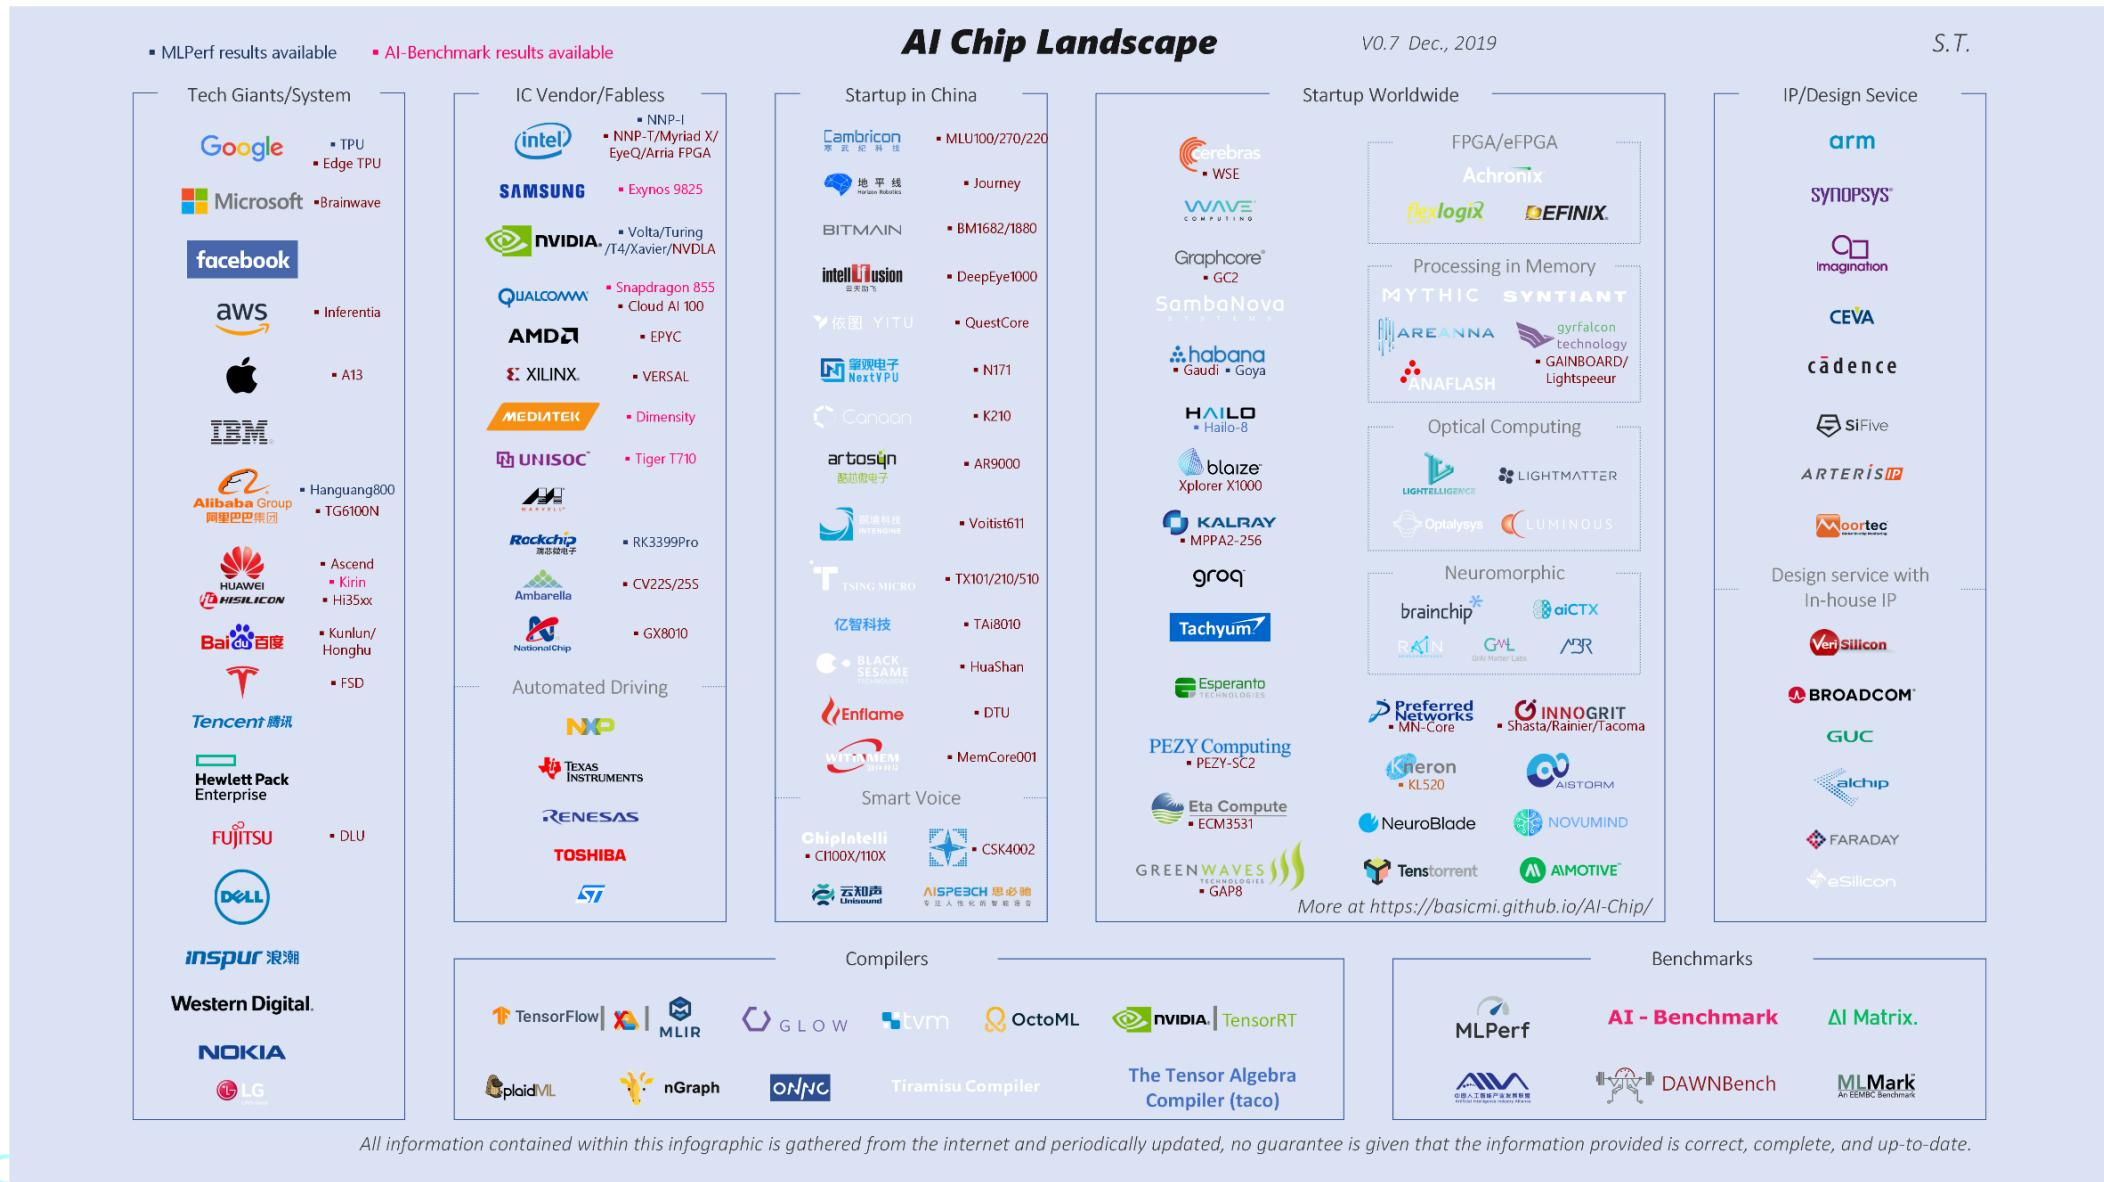


- **Design Productivity**
- High-Level Synthesis
- Chisel
- SystemVerilog (next lecture)

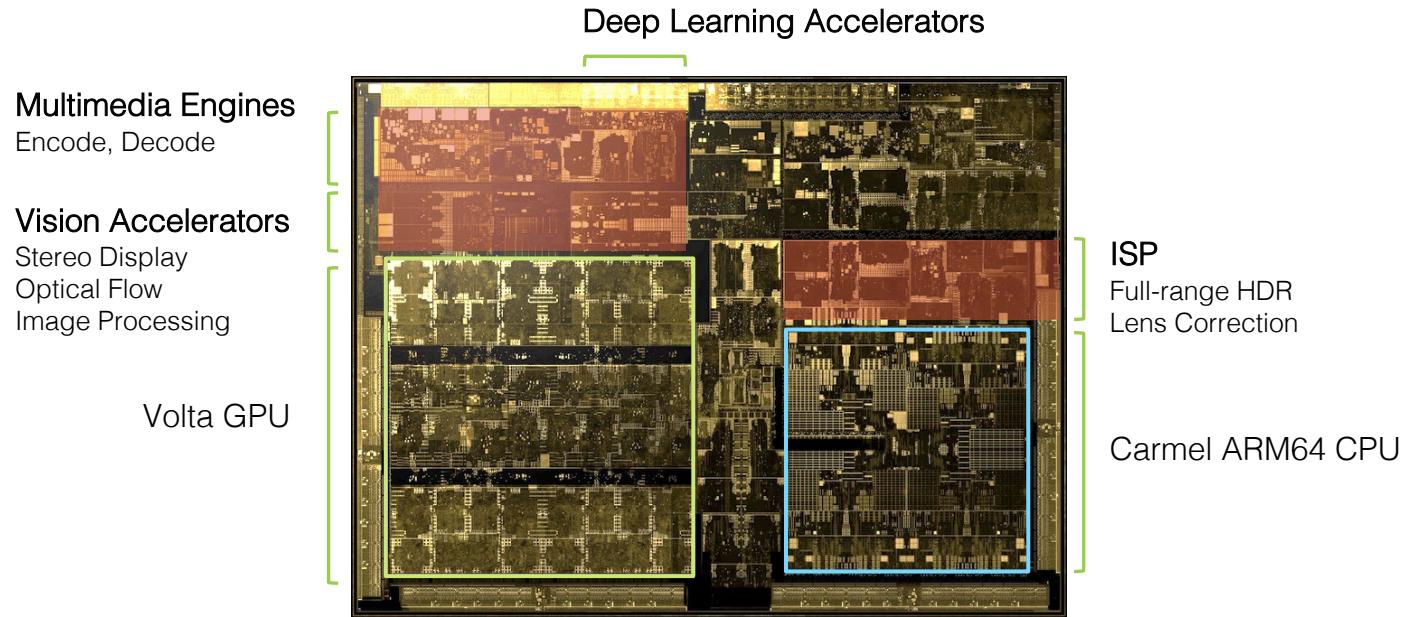
# Resources

- Chisel
  - Chisel bootcamp: <https://github.com/freechipsproject/chisel-bootcamp>
  - Professor Scott Beamer's lectures on Agile Hardware Design in Jupyter notebooks (<https://github.com/agile-hw/lectures>)
  - Resources: <https://www.chisel-lang.org/chisel3/docs/resources/resources.html>
  - Scala course: <https://www.coursera.org/specializations/scala>
- High-Level Synthesis
  - Professor Zhiru Zhang's course on High-Level Synthesis (<https://www.cs.l.cornell.edu/courses/ece5775/schedule.html>)
  - High-Level Synthesis Blue Book
  - Ryan Kastner, Janarbek Matai, and Stephen Neuendorffer, [Parallel Programming for FPGAs](#), arXiv, 2018.

# Golden Age of Computer Architecture



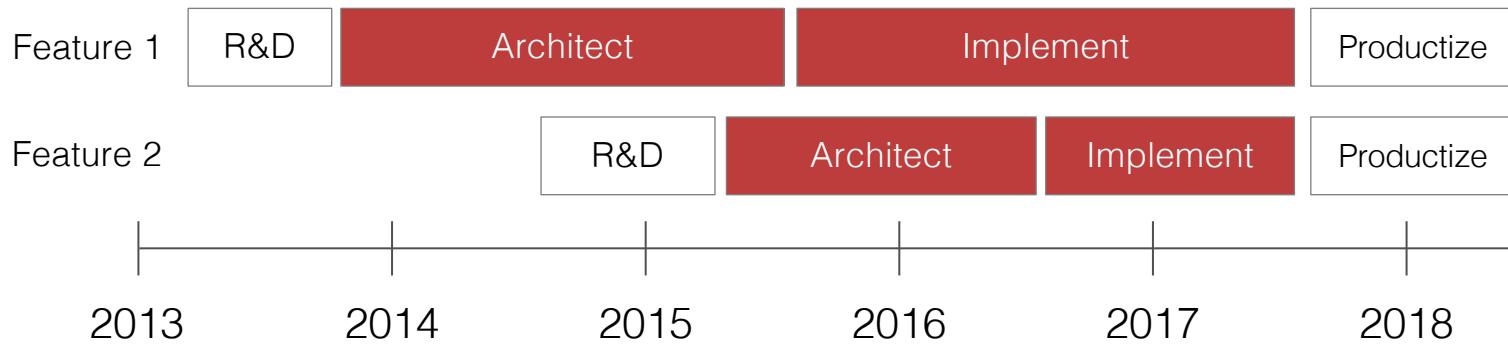
# Today's Systems-on-Chip are complex



"Most Complex SoC Ever Made  
~8,000 Engineering Years"

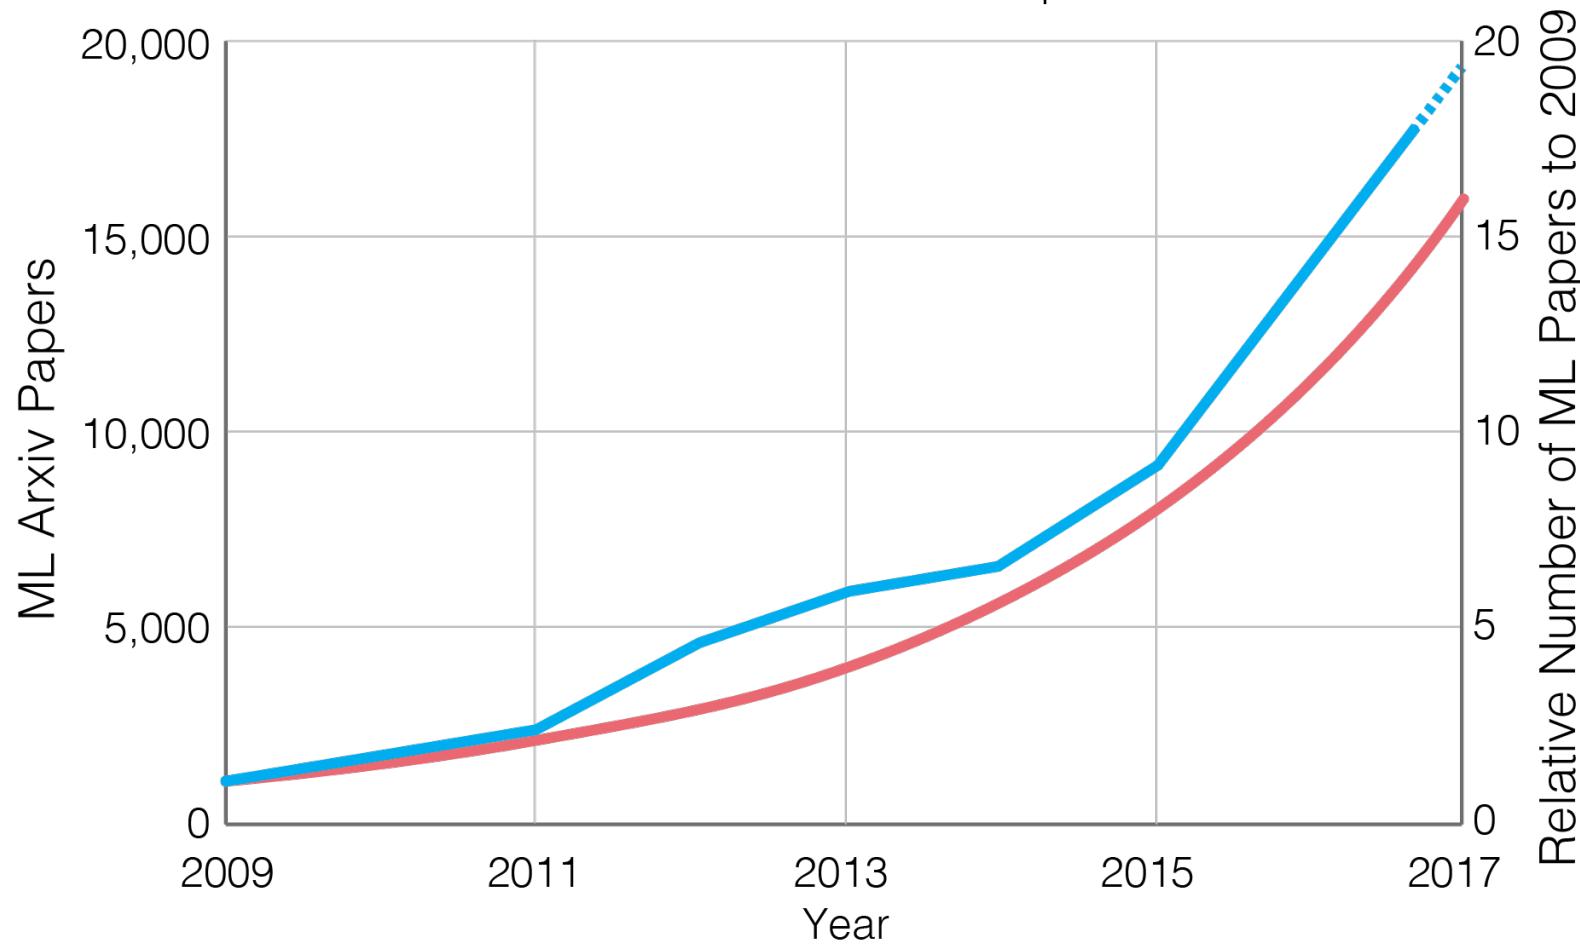
\* NVIDIA Xavier Announcement, CES'2018

# Hardware design efforts are relatively slow



- The number of unique features in hardware is limited by design effort

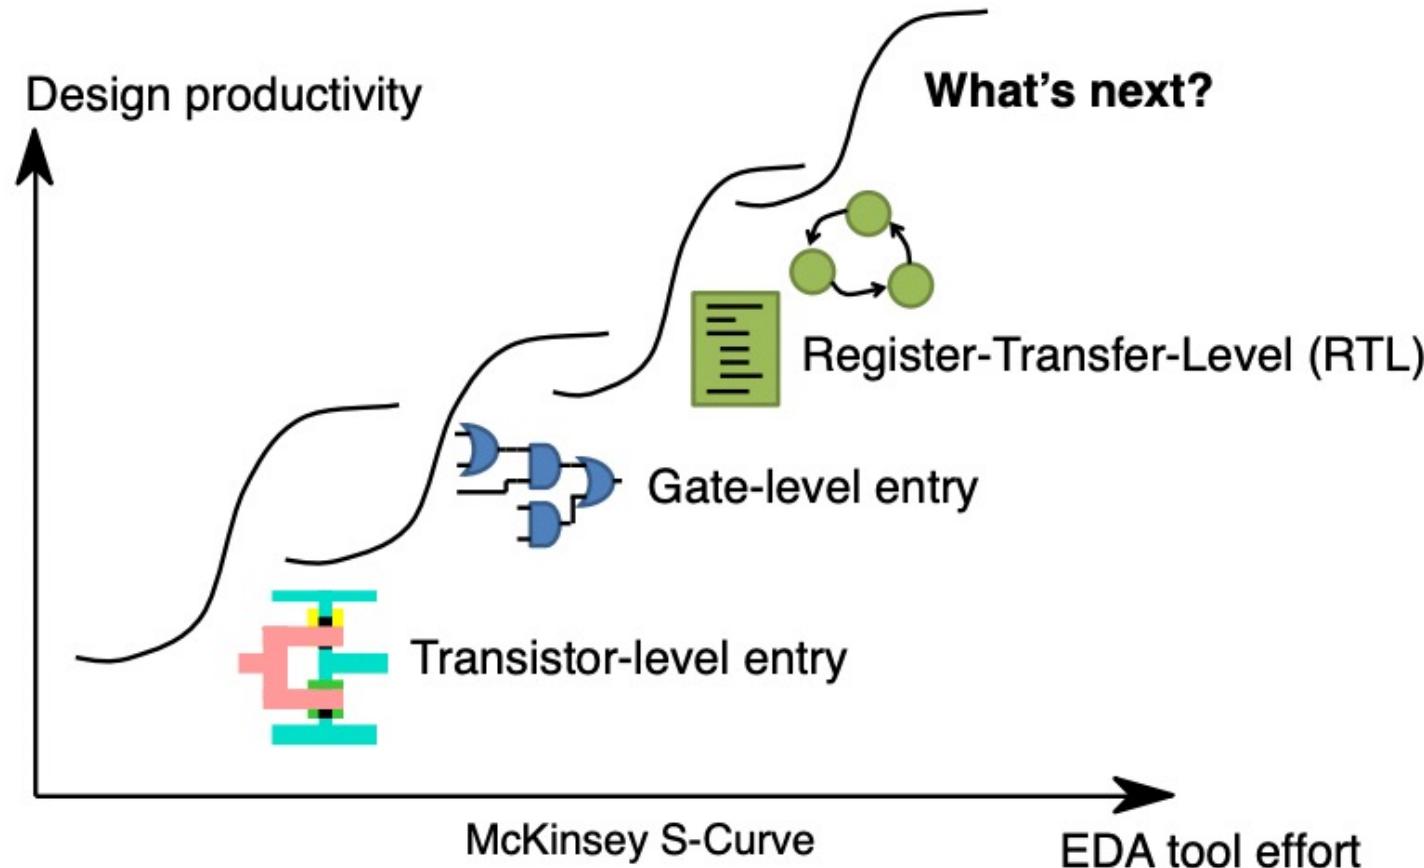
# While algorithms are changing relatively quickly





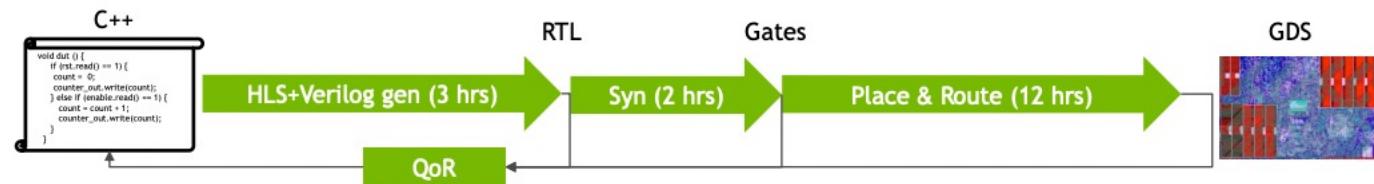
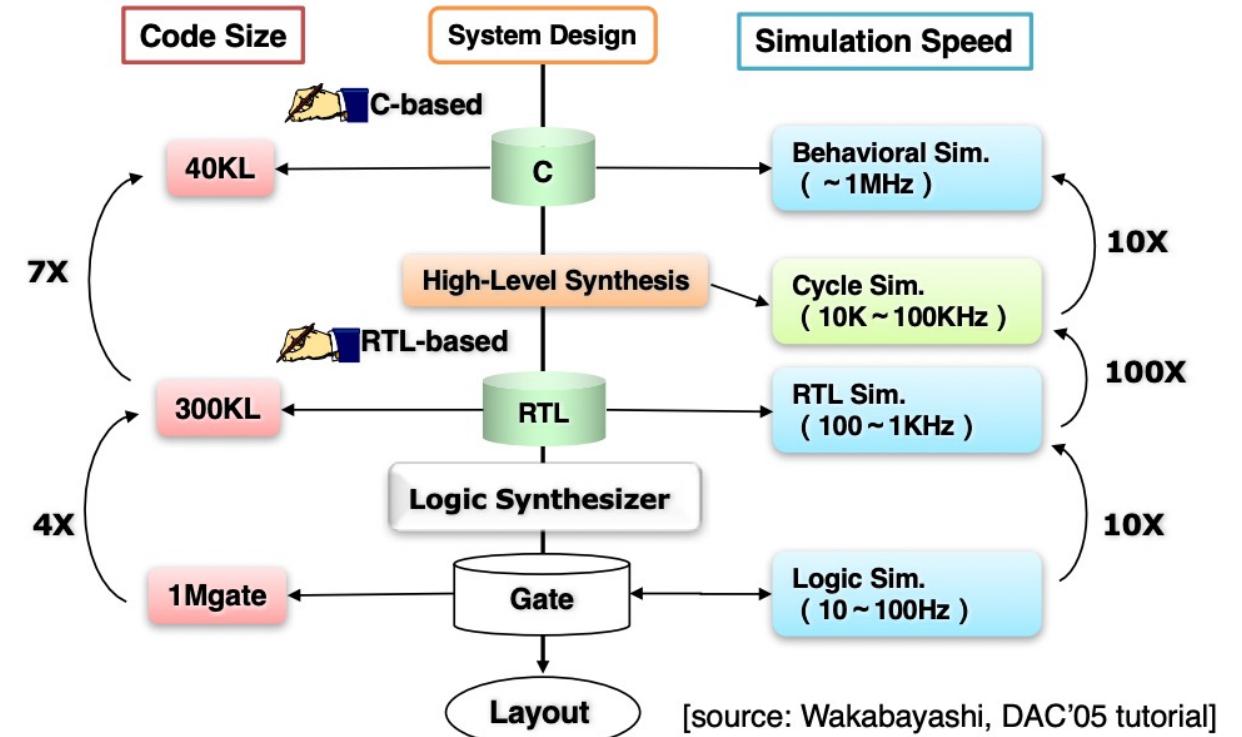
- Design Productivity
- **High-Level Synthesis**
- Chisel
- SystemVerilog (next lecture)

# Design Abstraction Evolution



# High-Level Synthesis (HLS)

- Automated design process that transforms a high-level functional specification to optimized register-transfer level (RTL) descriptions for efficient hardware implementation
- Why HLS?
  - Productivity
    - Lower design complexity and faster simulation speed
  - Portability
    - Single source -> multiple implementations
  - Permutability
    - Rapid design space exploration -> higher quality of result (QoR)



[Source: NVIDIA, 2019]

# High-Level Synthesis (HLS)

- Automated design process that transforms a high-level functional specification to optimized register-transfer level (RTL) descriptions for efficient hardware implementation
- Why HLS?
  - Productivity
    - Lower design complexity and faster simulation speed
  - Portability
    - Single source -> multiple implementations
  - Permutability
    - Rapid design space exploration -> higher quality of result (QoR)

```
module dut(rst, clk, q);
    input rst;
    input clk;
    output q;
    reg [7:0] c;

    always @ (posedge clk)
    begin
        if (rst == 1b'1) begin
            c <= 8'b00000000;
        end
        else begin
            c <= c + 1;
        end
    end

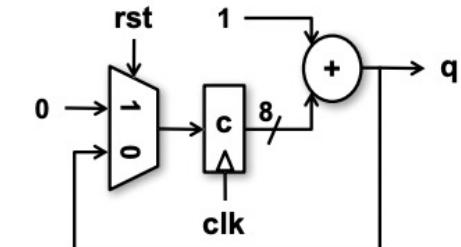
    assign q = c;
endmodule
```

RTL Verilog

```
uint8 dut() {
    static uint8 c;
    c+=1;
}
```

vs.

High-Level  
Synthesis



An 8-bit counter

# Case study: FIR Filter

$$y[n] = \sum_{i=0}^N b_i x[n - i]$$

$x[n]$  input signal

$y[n]$  output signal

$N$  filter order

$b_i$  *i*th filter coefficient

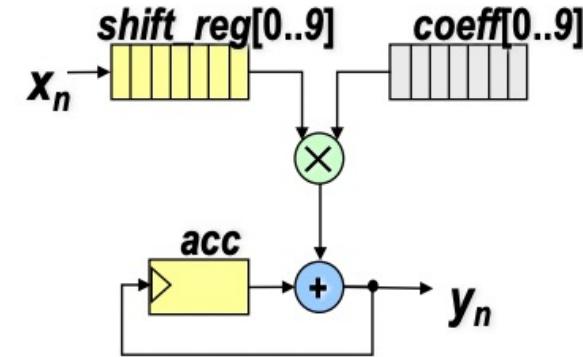
```
//original FIR
#define SIZE 128
#define N 10
void fir(int input[SIZE], int out[SIZE]) {
    // FIR coefficients
    int coeff[N] = {13,-2,9,...,74};
    // exact translation from the FIR equation
    for (int n = 0; n < SIZE; n++) {
        int acc = 0;
        for (int i = 0; i < N; i++) {
            if (n - i >= 0)
                acc += coeff[i] * input[n-i];
        }
        output[n] = acc;
    }
}
```

[Example adapted from Prof. Zhiru Zhang's HLS tutorial.]

# Case study: FIR Filter: Default Microarchitecture

```
void fir(int input[SIZE], int out[SIZE]) {  
    // FIR coefficients  
    int coeff[N] = {13, -2, 9, ..., 74};  
    // shift registers  
    int shift_reg[N] = {0};  
    for (int n = 0; n < SIZE; n++) {  
        int acc = 0;  
        for (int j = N - 1; j>0; j--) {  
            shift_reg[j] = shift_reg[j-1];  
        }  
        shift_reg[0] = input[n];  
        for (int i = 0; i < N; i++) {  
            acc += coeff[i] * shift_reg[i];  
        }  
        output[n] = acc;  
    }  
}
```

$$y[n] = \sum_{i=0}^N b_i x[n-i]$$



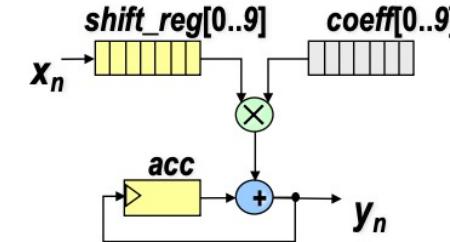
## Possible optimizations

- Loop unrolling
- Array partitioning
- Pipelining

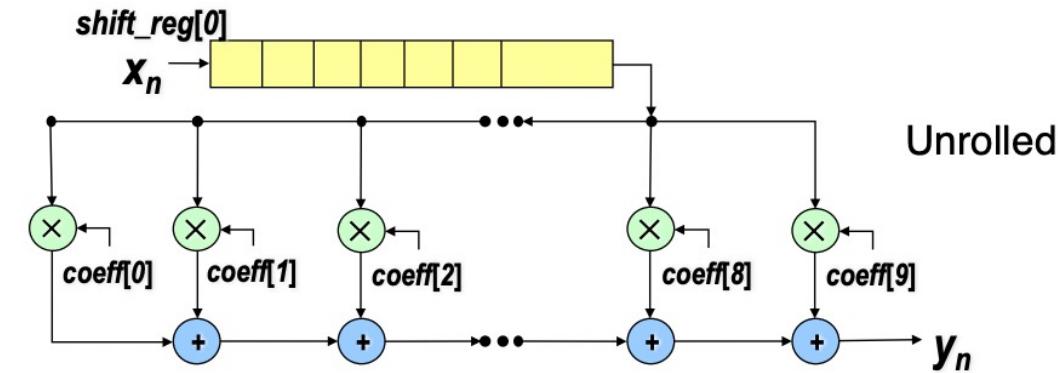
[Example adapted from Prof. Zhiru Zhang's HLS tutorial.]

# Case study: FIR Filter: Unrolling

```
void fir(int input[SIZE], int out[SIZE]) {  
    // FIR coefficients  
    int coeff[N] = {13,-2,9,...,74};  
    // shift registers  
    int shift_reg[N] = {0};  
    for (int n = 0; n < SIZE; n++) {  
        int acc = 0;  
        for (int j = N - 1; j>0; j--) {  
            #pragma HLS unroll  
            shift_reg[j] = shift_reg[j-1];  
        }  
        shift_reg[0] = input[n];  
        for (int i = 0; i < N; i++) {  
            #pragma HLS unroll  
            acc += coeff[i] * shift_reg[i];  
        }  
        output[n] = acc;  
    }  
}
```



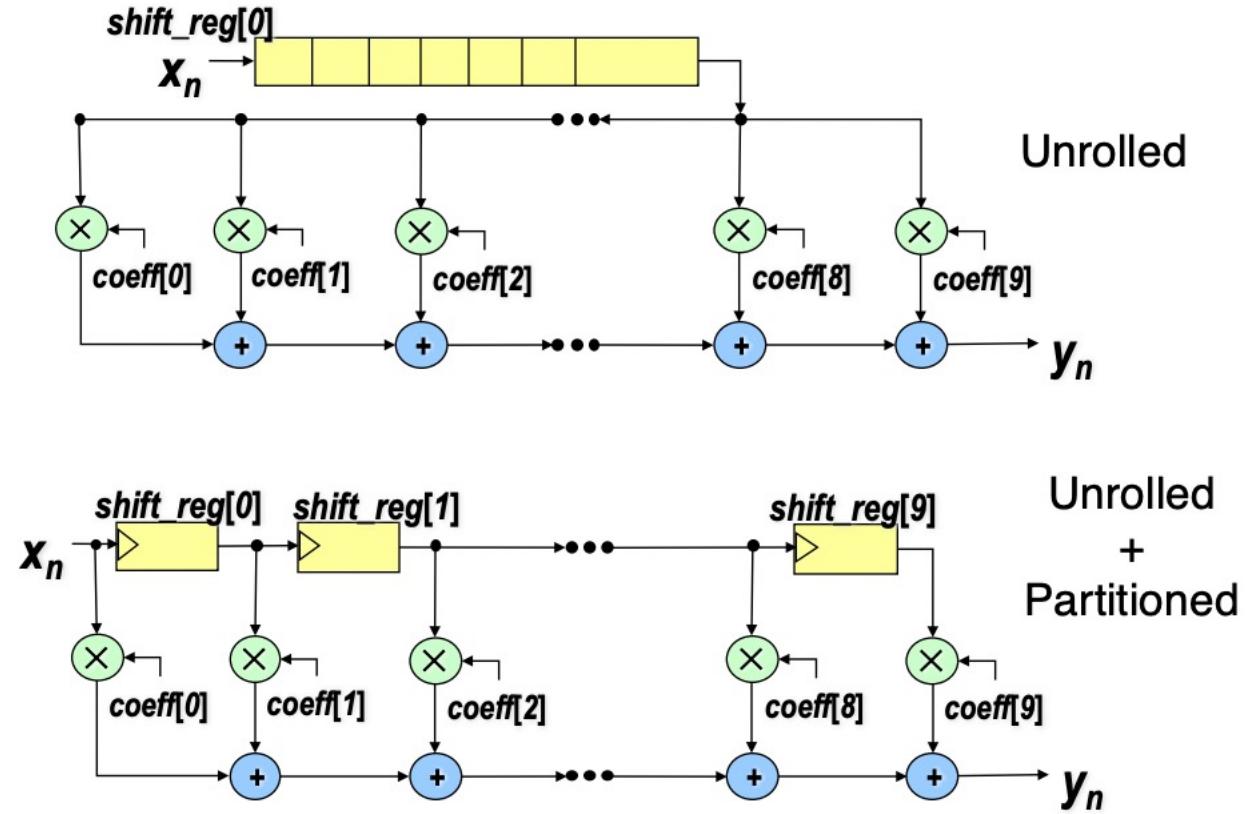
Default



Unrolled

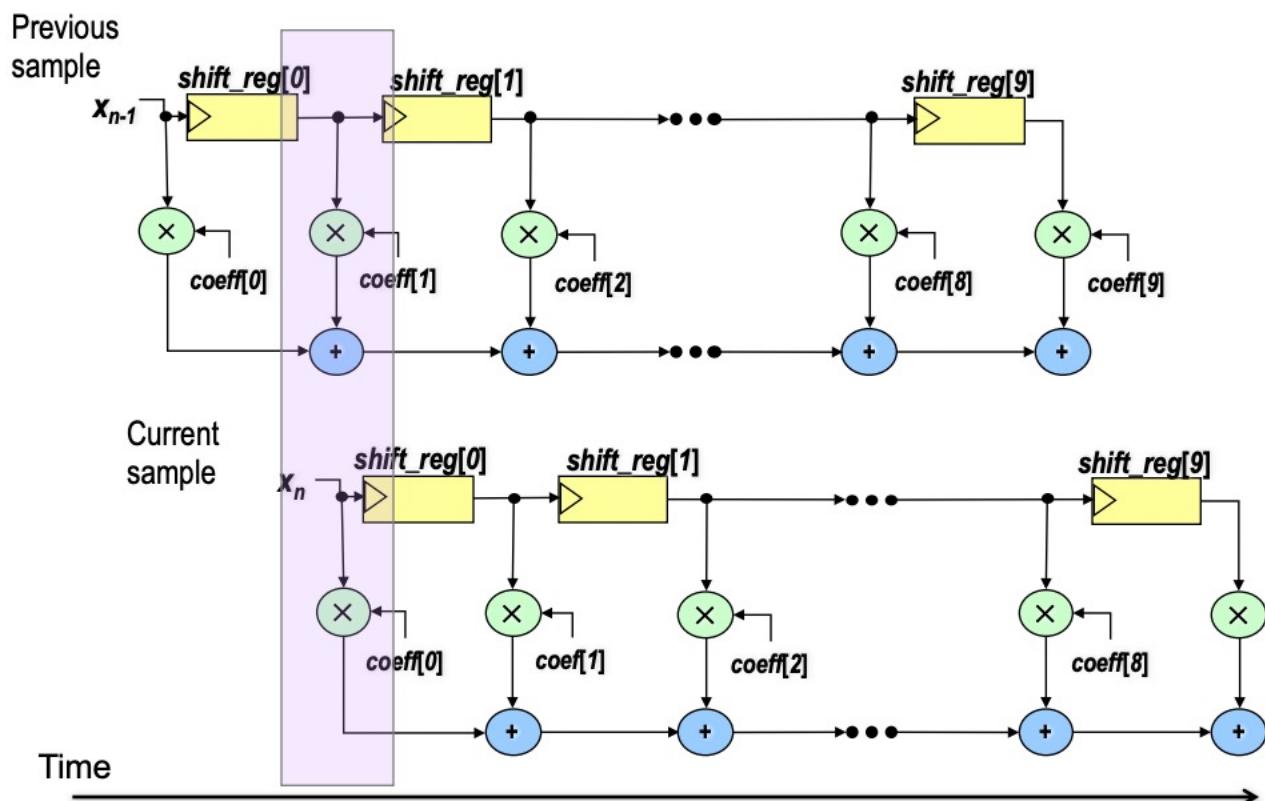
# Case study: FIR Filter: Partitioning

```
void fir(int input[SIZE], int out[SIZE])  
{  
    // FIR coefficients  
    int coeff[N] = {13, -2, 9, ..., 74};  
    // shift registers  
    int shift_reg[N] = {0};  
    #pragma HLS ARRAY_PARTITION  
    variable=shift_reg complete dim=0  
    for (int n = 0; n < SIZE; n++) {  
        for (int j = N - 1; j>0; j--) {  
            ...  
            ...  
            for (int i = 0; i < N; i++) {  
                ...  
            }  
        }  
    }  
}
```



# Case study: FIR Filter: Pipelining

```
void fir(int input[SIZE], int  
out[SIZE]) {  
    // FIR coefficients  
    int coeff[N] = {13, -2, 9, ..., 74};  
    // shift registers  
    int shift_reg[N] = {0};  
    for (int n = 0; n < SIZE; n++) {  
        #pragma HLS pipeline II=1  
        int acc = 0;  
        ...  
        output[n] = acc;  
    }  
}
```



# High-Level Synthesis

- Exponential growth in HW design complexity calls for higher level of design abstraction.
- To raise hardware design level of abstraction
  - Use **high-level languages**,
    - e.g., C++ instead of Verilog
  - Use **automation**,
    - e.g., HLS tools to generate Verilog
  - Use **libraries/generators**,
    - e.g., MatchLib from NVIDIA



- Design Productivity
- High-Level Synthesis
- Chisel
- SystemVerilog (next lecture)

Adapted from slides by Hasan Genc and Brendan Sweeney

# Verilog/VHDL

- Old, extremely popular
- If you get a job writing RTL, it will probably be in one of these
- Cons
  - Poor parameterizability
    - Some companies write scripts to generate Verilog
  - Weak type-safety (esp. Verilog)
    - Weak notion of signed vs unsigned
    - Easy to accidentally truncate bits
  - Easy to generate non-synthesizable hardware by mistake
    - Originally a simulation language
      - Hardware primitives must be *inferred*
    - Huge market for Verilog linters!

*Old school hardware engineers will tell you that it's fine. It's fine that the language is so counter-intuitive that almost all people who initially approach Verilog write code that's not just wrong but nonsensical.*

Dan Luu

# SystemVerilog



- Superset of Verilog
- New features
  - Especially for verification
  - More expressive
    - Clearer datatypes, better abstractions, etc.
- More about SystemVerilog from Vighnesh (next lecture)
- Cons:
  - Also possible to generate non-synthesizable hardware by mistake
  - Some advanced features are not supported by all vendors



- Embeds an HDL in Scala
  - Verilog/VHDL/SystemVerilog are standalone languages
- Chisel describes the hardware...
  - ...and Scala parameterizes it
- Pros:
  - No confusing separation between synthesizable and simulation-only hardware
  - Easy to define new types and *typeclasses*
    - You don't have to always think in terms of "bits"
    - Operator overloading, implicit datatype conversion, etc.
  - Powerful metaprogramming and parameterizability
    - Don't just write hardware!
    - Write *programs* that generate hardware!
    - Just look at *Diplomacy* for proof

# But surely there must be a catch?

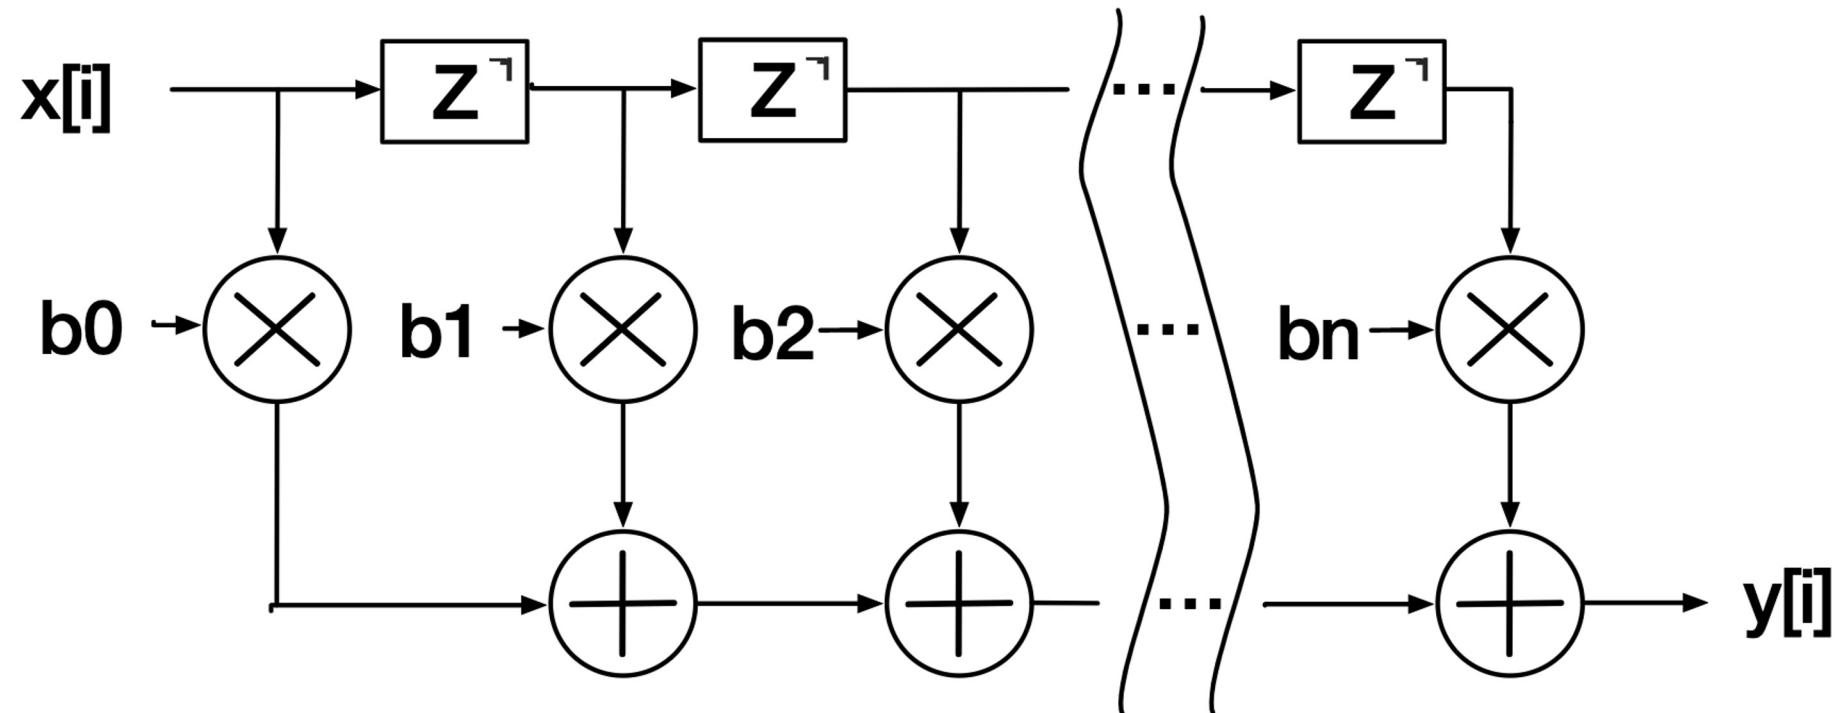
- Chisel cons:
  - Awkward support for latches and asynchronous hardware
  - Verification flow is not (yet) as advanced as with SystemVerilog
    - Some examples of Chisel working with UVM and SVA exist/being worked on (ChiselVerify, Dynamic Verification Library (verif), CHA)
  - Some people are scared of Scala
    - Some parts of the Scala ecosystem are painful to use
      - SBT, JVM, ...
- We will also talk about issues specific to the Chisel **compiler** later

# What about other options?

- There are *lots* of HDLs
- Embedded in lots of different languages
  - Including Python and Haskell
- Examples
  - BlueSpec
  - SpinalHDL
  - Magma
  - PyRTL
  - Clash
  - MyHDL

# Let's Make an FIR Filter

(example from the Chisel website)



# Prettier Verilog

This is how beginner typically use Chisel:

```
// 3-point moving sum implemented in the style of a FIR filter
class MovingSum3(bitWidth: Int) extends Module {
    val io = IO(new Bundle {
        val in = Input(UInt(bitWidth.W))
        val out = Output(UInt(bitWidth.W))
    })

    val z1 = RegNext(io.in)
    val z2 = RegNext(z1)

    io.out := (io.in * 1.U) + (z1 * 1.U) + (z2 * 1.U)
}
```

# Where Chisel Shines

- Write *generators*, not *circuits*
- Chisel makes functional programming *easy*
  - Maps, reduces, tabulates...
- But it doesn't *force* functional programming on you
  - For-loops and stateful statements are OK too!

```
// Generalized FIR filter parameterized by the convolution coefficients
class FirFilter(bitWidth: Int, coeffs: Seq[UInt]) extends Module {
    val io = IO(new Bundle {
        val in = Input(UInt(bitWidth.W))
        val out = Output(UInt(bitWidth.W))
    })
    // Create the serial-in, parallel-out shift register
    val zs = Reg(Vec(coeffs.length, UInt(bitWidth.W)))
    zs(0) := io.in
    for (i <- 1 until coeffs.length) {
        zs(i) := zs(i-1)
    }

    // Do the multiplies
    val products = VecInit.tabulate(coeffs.length)(i => zs(i) * coeffs(i))

    // Sum up the products
    io.out := products.reduce(_ + _)
}
```

# Where Chisel *Especially* Shines

- Types!
  - Powerful typeclasses
  - Great type-checking
    - Fewer bugs!

```
// Generalized FIR filter parameterized by the convolution coefficients
class FirFilter[T <: Data : Ring](t: T, coeffs: Seq[T]) extends Module {
    val io = IO(new Bundle {
        val in = Input(t)
        val out = Output(t)
    })
    // Create the serial-in, parallel-out shift register
    val zs = Reg(Vec(coeffs.length, t))
    zs(0) := io.in
    for (i <- 1 until coeffs.length) {
        zs(i) := zs(i-1)
    }
    // Do the multiplies
    val products = VecInit.tabulate(coeffs.length)(i => zs(i) * coeffs(i))
    // Sum up the products
    io.out := products.reduce(_ + _)
}
```

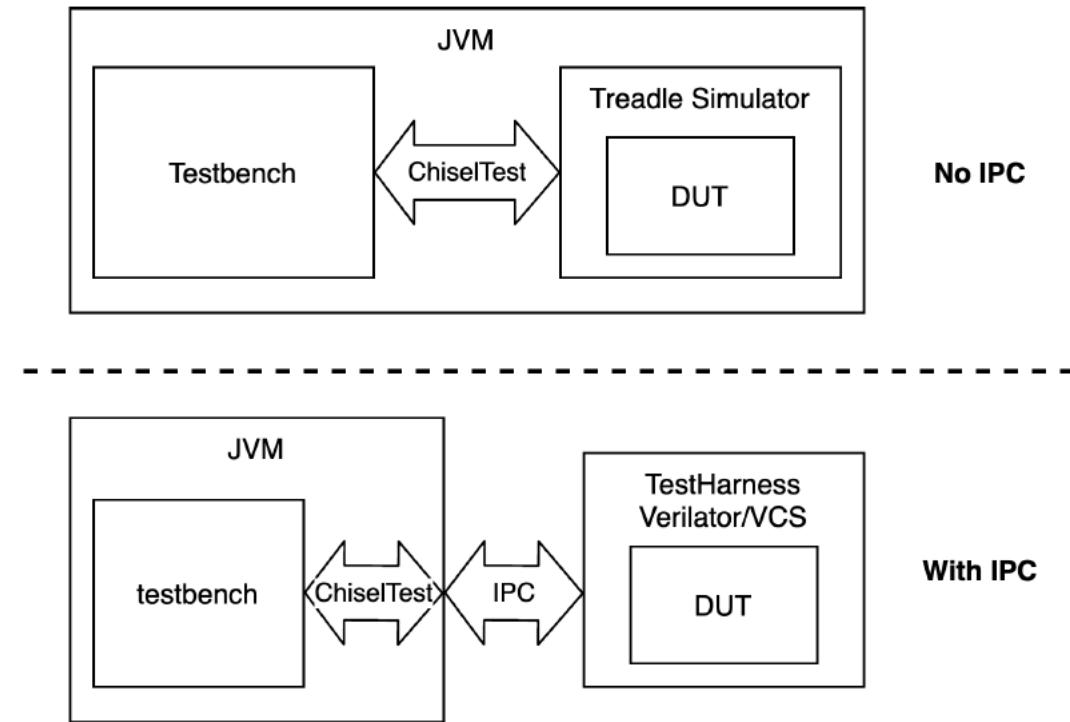
# Generators Make Hardware Reusable

Not unique to Chisel, but Chisel makes it *easier*

```
val movingSum3Filter = Module(new FirFilter(UInt(8.W) Seq(1.U, 1.U, 1.U)))
val delayFilter = Module(new FirFilter(SInt(16.W) Seq(0.S, 1.S)))
val triangleFilter = Module(new FirFilter(Float(8,24), Seq(1.Fl, 2.Fl, 3.Fl, 2.Fl, 1.Fl)))
```

# ChiselTest – A test harness for CHISEL RTL

```
import chisel3._  
import chiseltest._  
import org.scalatest.flatspec.AnyFlatSpec  
  
class SimpleTestExpect extends AnyFlatSpec  
  with ChiselScalatestTester {  
  "DUT" should "pass" in {  
    test(new DeviceUnderTest) { dut =>  
      dut.io.a.poke(0.U)  
      dut.io.b.poke(1.U)  
      dut.clock.step()  
      dut.io.out.expect(0.U)  
      dut.io.a.poke(3.U)  
      dut.io.b.poke(2.U)  
      dut.clock.step()  
      dut.io.out.expect(2.U)  
    }  
  }  
}
```



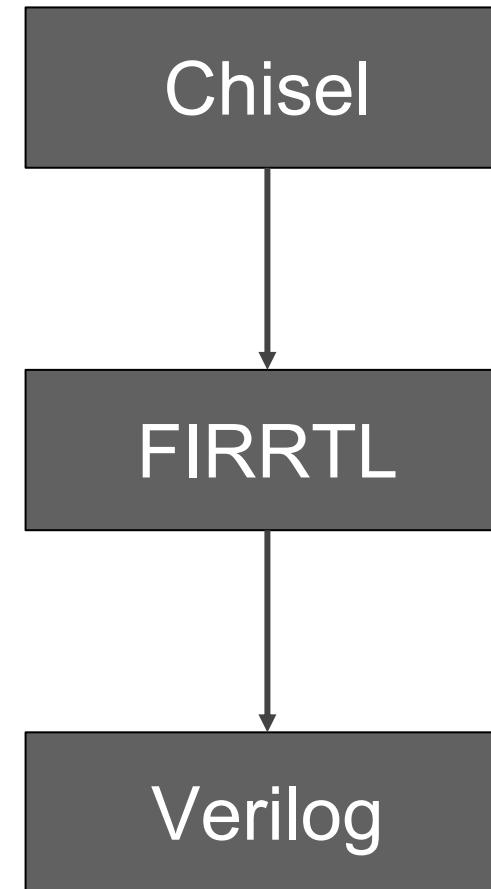
- Based on ScalaTest
  - Set input values with *poke*
  - Read output values with *peek*
  - Compare the values with *expect*

# Chisel high-points

- List manipulations are easy
  - Map, reduce, scan, etc.
- Parameterizing across *types* is easy
  - With an excellent type system to make sure you don't make mistakes
- Reusable hardware
  - Very convenient standard library
    - Queues, handshake-based interfaces, bit-manipulation functions, counters, etc.
  - RocketChip components are (kind-of) plug-and-play
    - TLBs, caches, entire CPUs, etc.
- Excellent IDE support
  - Write hardware in IntelliJ

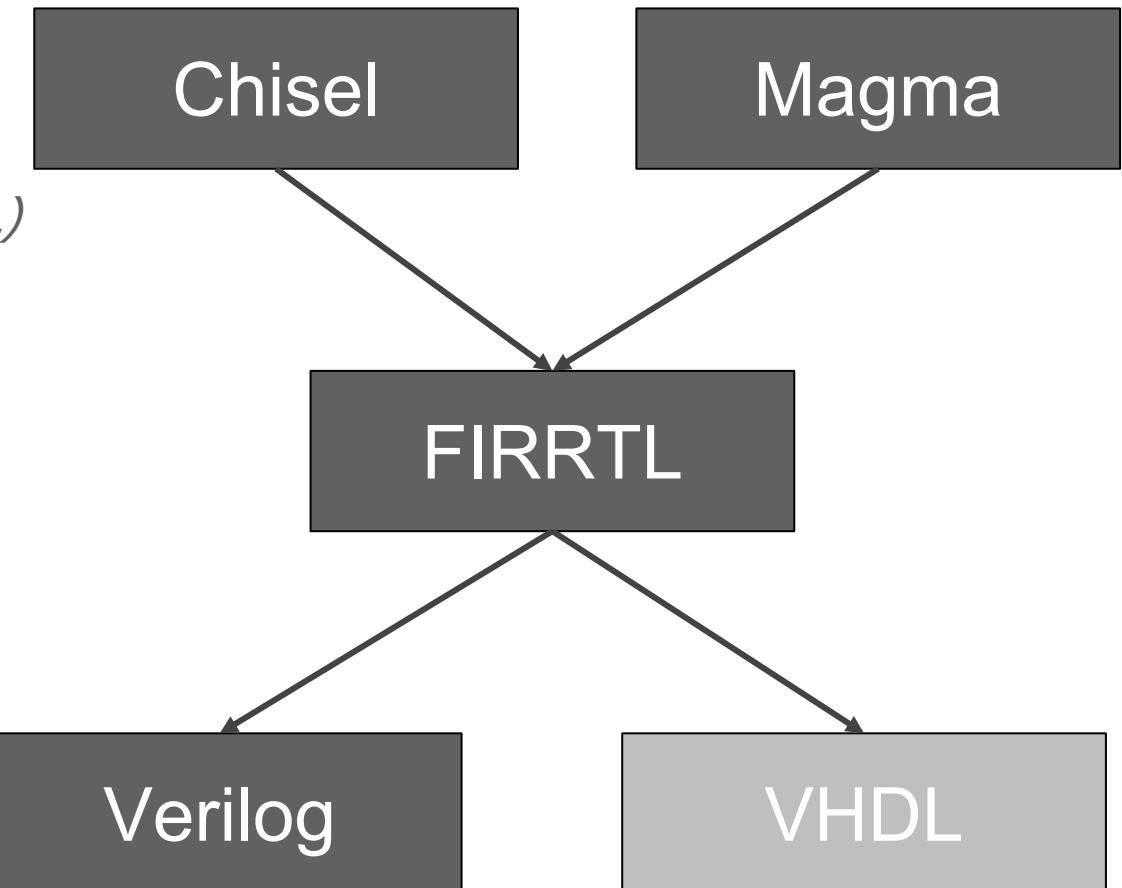
# Chisel Compiler

- Chisel has an intermediate representation (IR) called *FIRRTL (Flexible Intermediate Representation for RTL)*
  - *Chisel*: Clang
  - *FIRRTL*: LLVM
  - *Verilog*: Machine code
- FIRRTL supports *transformations* on FIRRTL data structures (abstract syntax tree – AST)
  - Early area estimations
  - Add hardware counters and instrumentation
  - Optimize RTL for ASIC vs FPGA



# Chisel Compiler

- Chisel has an intermediate representation (IR) called *FIRRTL (Flexible Intermediate Representation for RTL)*
  - *Chisel*: Clang
  - *FIRRTL*: LLVM
  - *Verilog*: Machine code
- FIRRTL supports *transformations* on FIRRTL data structures (abstract syntax tree – AST)
  - Early area estimations
  - Add hardware counters and instrumentation
  - Optimize RTL for ASIC vs FPGA



# CHISEL -> FIRRTL -> Verilog

## CHISEL

```
// 3-point moving sum implemented in the style of a FIR filter
class MovingSum3(bitWidth: Int) extends Module {
    val io = IO(new Bundle {
        val in = Input(UInt(bitWidth.W))
        val out = Output(UInt(bitWidth.W))
    })

    val z1 = RegNext(io.in)
    val z2 = RegNext(z1)

    io.out := (io.in * 1.U) + (z1 * 1.U) + (z2 * 1.U)
}
```

## FIRRTL

```
1 circuit FirFilter :
2     module FirFilter :
3         input clock : Clock
4         input reset : UInt<1>
5         output io : {flip in : UInt<8>, out : UInt<8>}
6
7         reg zs : UInt<8>[3], clock
8         zs[0] <= io.in
9         zs[1] <= zs[0]
10        zs[2] <= zs[1]
11        node _T = mul(zs[0], UInt<1>("h01"))
12        node _T_1 = mul(zs[1], UInt<1>("h01"))
13        node _T_2 = mul(zs[2], UInt<1>("h01"))
14        wire products : UInt<9>[3]
15        products[0] <= _T
16        products[1] <= _T_1
17        products[2] <= _T_2
18        node _T_3 = add(products[0], products[1])
19        node _T_4 = tail(_T_3, 1)
20        node _T_5 = add(_T_4, products[2])
21        node _T_6 = tail(_T_5, 1)
22        io.out <= _T_6
```

## Verilog

```
1 module FirFilter(
2     input          clock,
3     input          reset,
4     input [7:0]    io_in,
5     output [7:0]   io_out
6 );
7     reg [7:0] zs_0;
8     reg [7:0] zs_1;
9     reg [7:0] zs_2;
10    wire [8:0] products_0 = zs_0 * 1'h1;
11    wire [8:0] products_1 = zs_1 * 1'h1;
12    wire [8:0] products_2 = zs_2 * 1'h1;
13    wire [8:0] _T_4 = products_0 + products_1;
14    wire [8:0] _T_6 = _T_4 + products_2;
15    assign io_out = _T_6[7:0];
16    always @ (posedge clock) begin
17        zs_0 <= io_in;
18        zs_1 <= zs_0;
19        zs_2 <= zs_1;
20    end
21 endmodule
```

- FIR Filter Example

# Advantages of the Chisel Compiler

- Optimization passes can be shared across hardware designs
  - A similar concept exists for software compilers
- Attach arbitrary metadata for any Chisel expression
  - The FIRRTL compiler can use that metadata later
    - E.g. don't const-propagate, FPGA register preset
- FIRRTL makes it easy to build *tooling* for RTL designs
  - Automatically create:
    - Test harness generators
    - Macro compilers
      - Useful for SRAMs!
    - Simulators
  - Can benefit from other people's FIRRTL optimizers

# Problems With the Chisel Compiler

- Generated Verilog can be unreadable (e.g. fully unrolled arrays, ...)
  - This can make debugging and verification more difficult
    - Especially for some automated verification tools
  - Some HDLs do a much better job with this
- Slow
- Memory hungry
  - Makes it hard to run this on a low-end laptop

# Chisel/FIRRTL is an *integral* part of the “Berkeley stack”

- CPUs
  - *RocketChip*: In-order CPU
    - 2k GitHub stars, 406 citations
  - *BOOM*: Out-of-order CPU
  - *Sodor*: Educational
- Accelerators
  - *Hwacha*: Vectors
  - *Gemmini*: DNNs
  - *Genesis*: Genomics
  - *Protobuf accelerator*
- Hardware design frameworks
  - *Chipyard*
- Simulation frameworks
  - *FireSim*
  - *MIDAS*

# People outside Berkeley use Chisel too!

- Industry
  - Google, IBM, SiFive, Lampro Mellon, Intel
    - Google's Edge TPU was designed with Chisel
- Academia
  - Cornell, Stanford, Technical University of Denmark, LBL, Boston University, UC Davis, Peking

# Some questions to think about

- Is Chisel focusing on the right problems?
- Why do so many people complain about Verilog/VHDL, but they still keep using it?
- Are HDLs even worth creating anymore?
  - Why not focus on HLS instead?
  - What other abstractions for generating hardware are possible?