

Lecture 5:

# High-Level Synthesis

Sitao Huang

[sitaoh@uci.edu](mailto:sitaoh@uci.edu)

February 1, 2022

*Slide courtesy of Prof. Deming Chen, UIUC, ECE 527: SoC Design*



# Logistics

- [Homework 1](#) released, due: **January 31**, 11:59 PM on Canvas
- Homework 2 expected to release later this week, due February 7
- **Midterm: February 10** (Thursday), 8:00-9:20 AM (in class)
- Midterm review session / Q&A: February 8 (Tuesday)
- Project proposal due: February 14
  - Options: (a) literature review paper or (b) compiler + accelerator project
- In-person **(hybrid)** instruction starting next week (week 5)! Classroom: [SSTR 101](#)
- First in-person class: **February 1**

# High-Level Synthesis

- **High-level synthesis (HLS)**, also referred to as C synthesis, electronic system-level (ESL) synthesis, algorithmic synthesis, or behavior synthesis
- HLS takes an abstract behavioral specification of a digital system, and finds a register-transfer level structure that realizes the given behavior

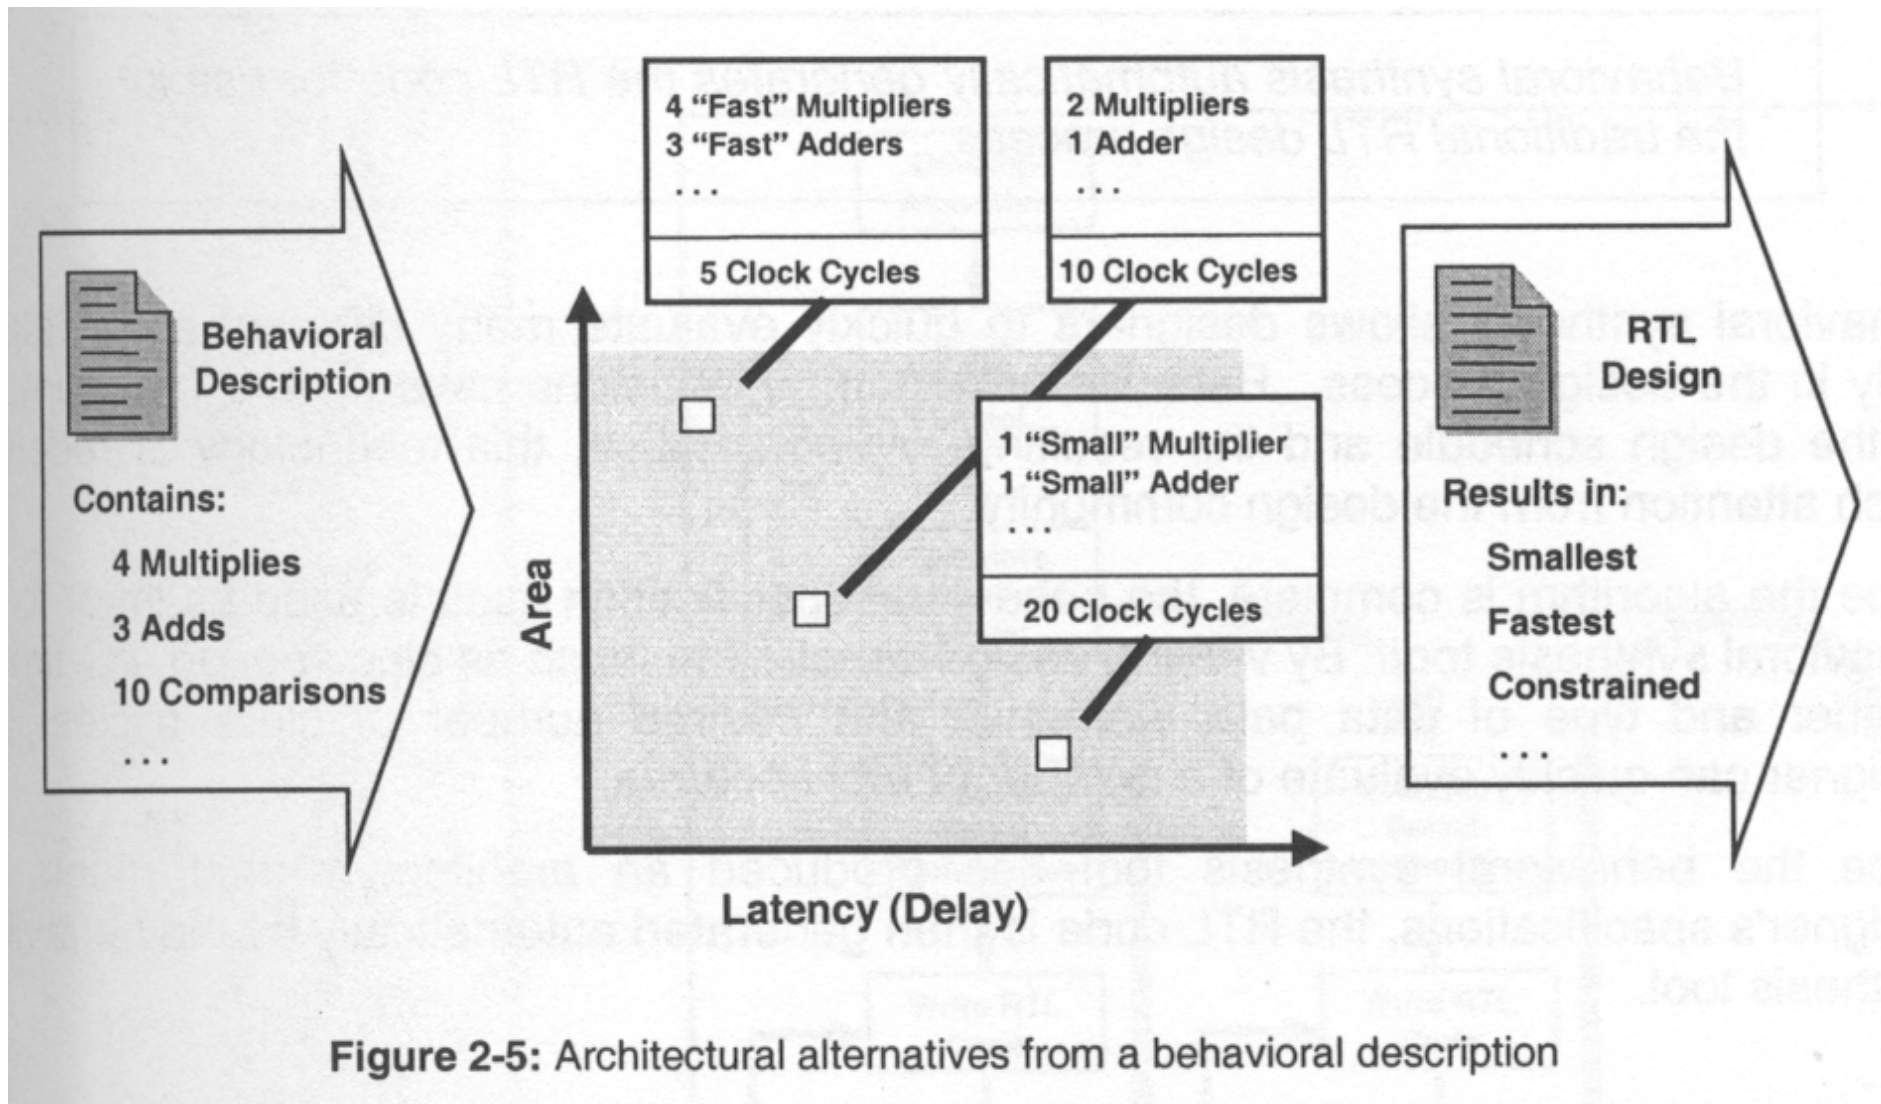
## INPUT:

- A high-level, algorithmic description
  - Control structures (if/else, loop, subroutines)
  - Concurrent and sequential semantics
  - Abstract data types
  - Logical and arithmetic operators
- A set of constraints
  - Speed, power, area, interconnect style
  - A library of pre-specified components


## OUTPUT:

- A register-transfer level description for further synthesis and optimization

# High-Level Synthesis




# Design Space Exploration



Add the Power tradeoff and this becomes a 3-dimensional graph

# Design Space Exploration: RTL vs. Behavioral Synthesis



# High-Level Synthesis Process

## ***Resource Allocation:***

- Allocating resources (library components) to each of the operations, buses, muxes, and registers for storage

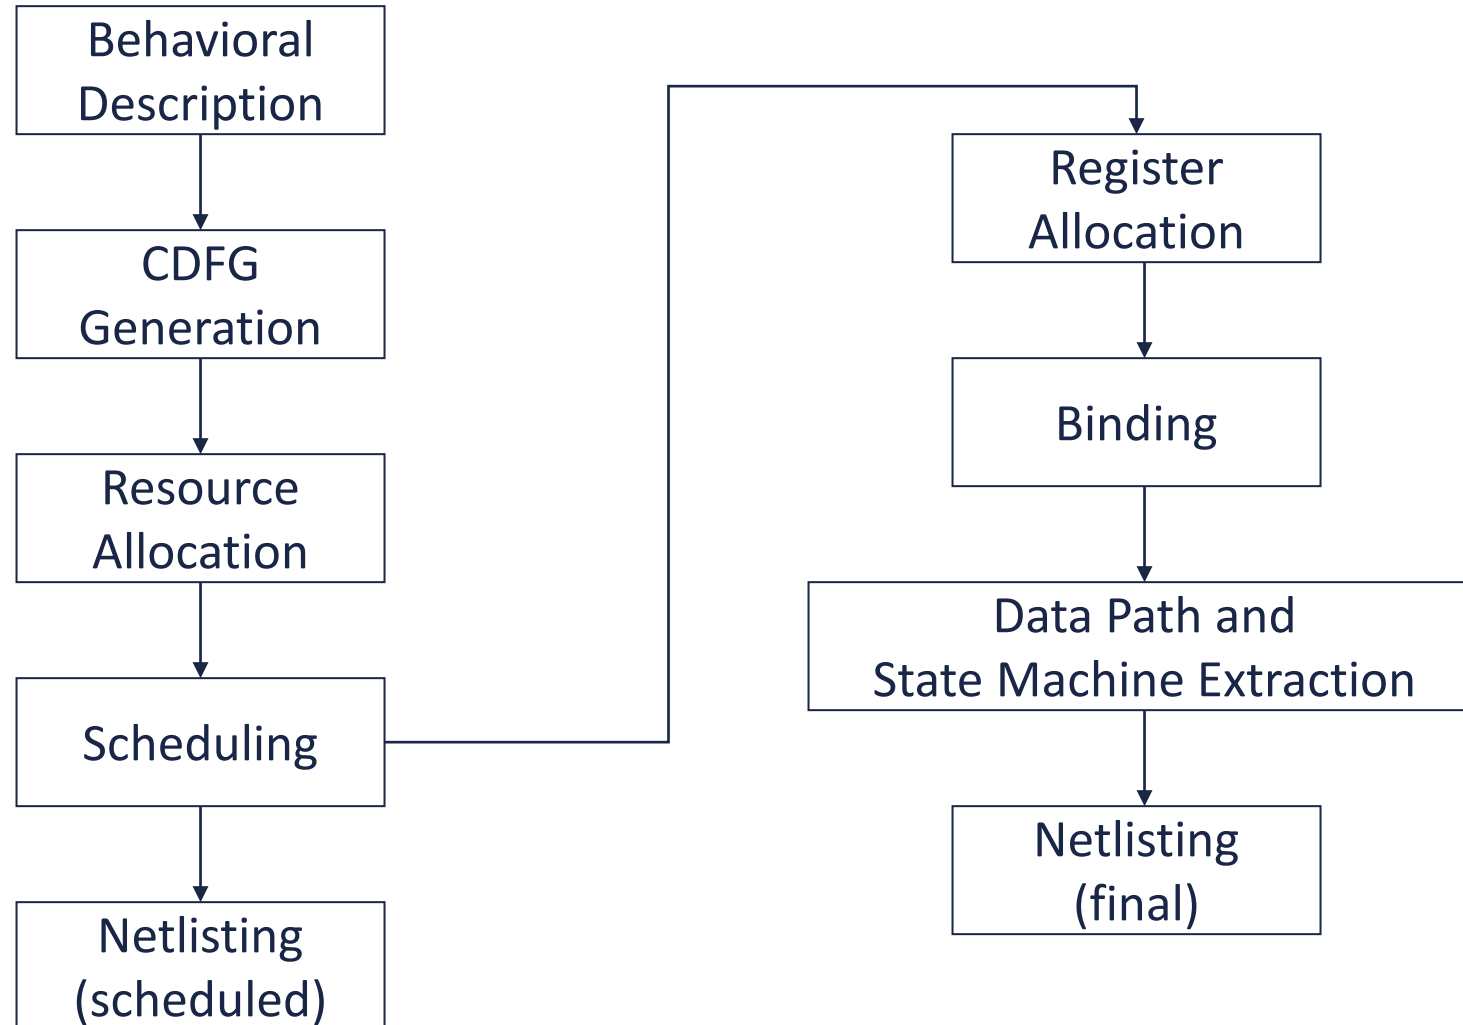
## ***Scheduling:***

- Scheduling the operations in the CDFG to minimize area, time and/or power

## ***Binding:***

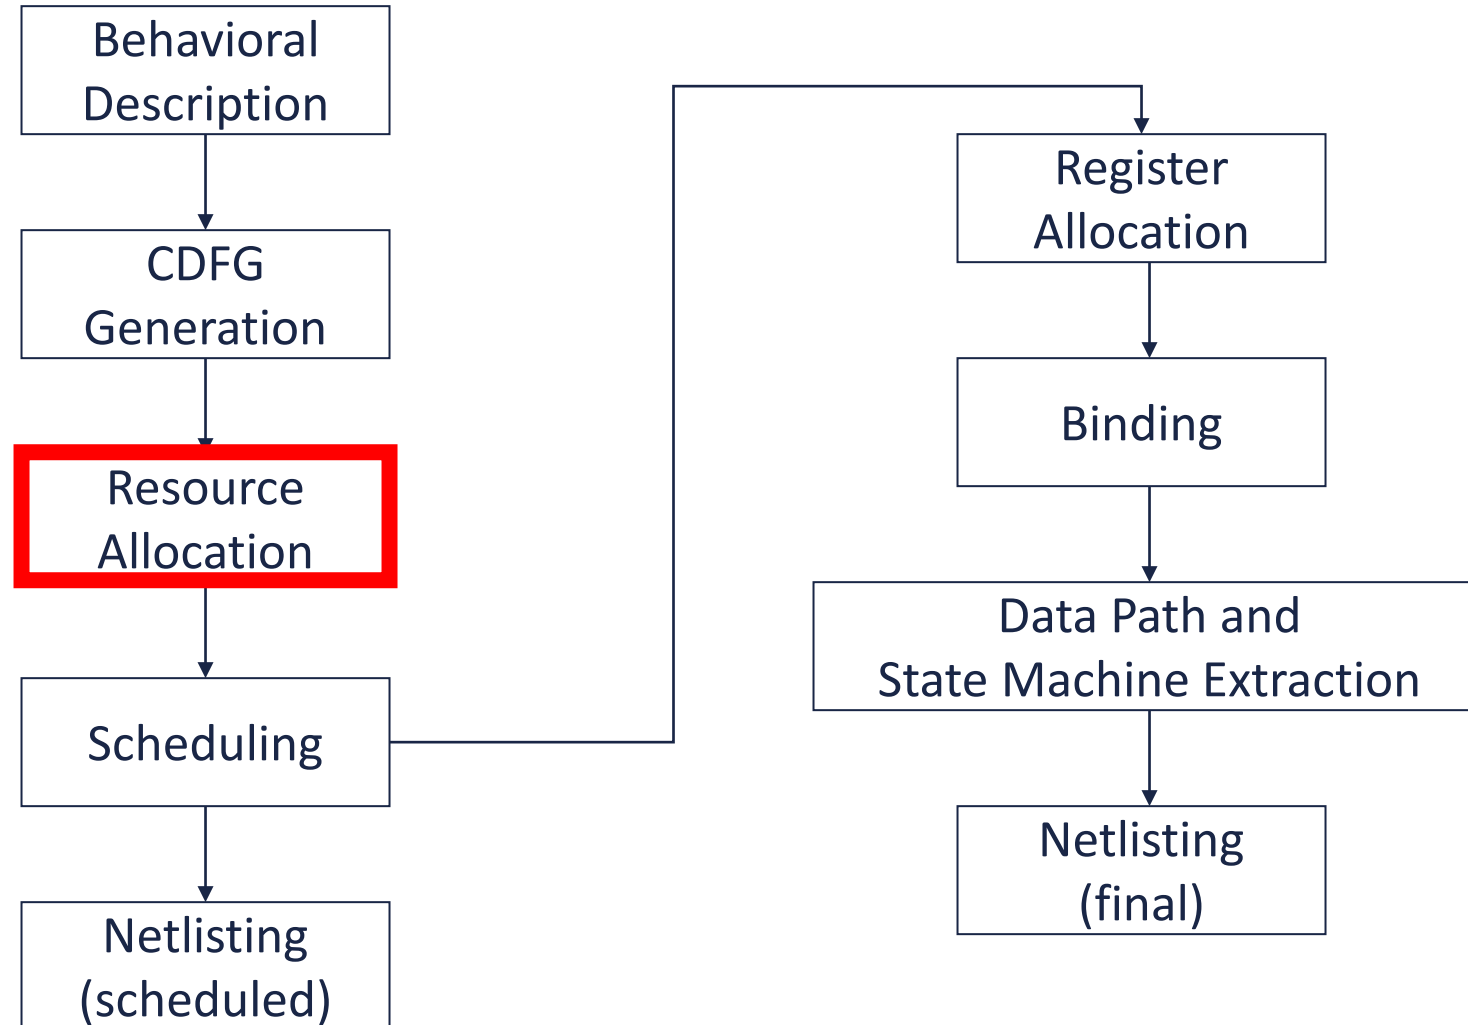
- Determining the time of use of each component
- e.g., which registers to use and when

# High-Level Synthesis Steps





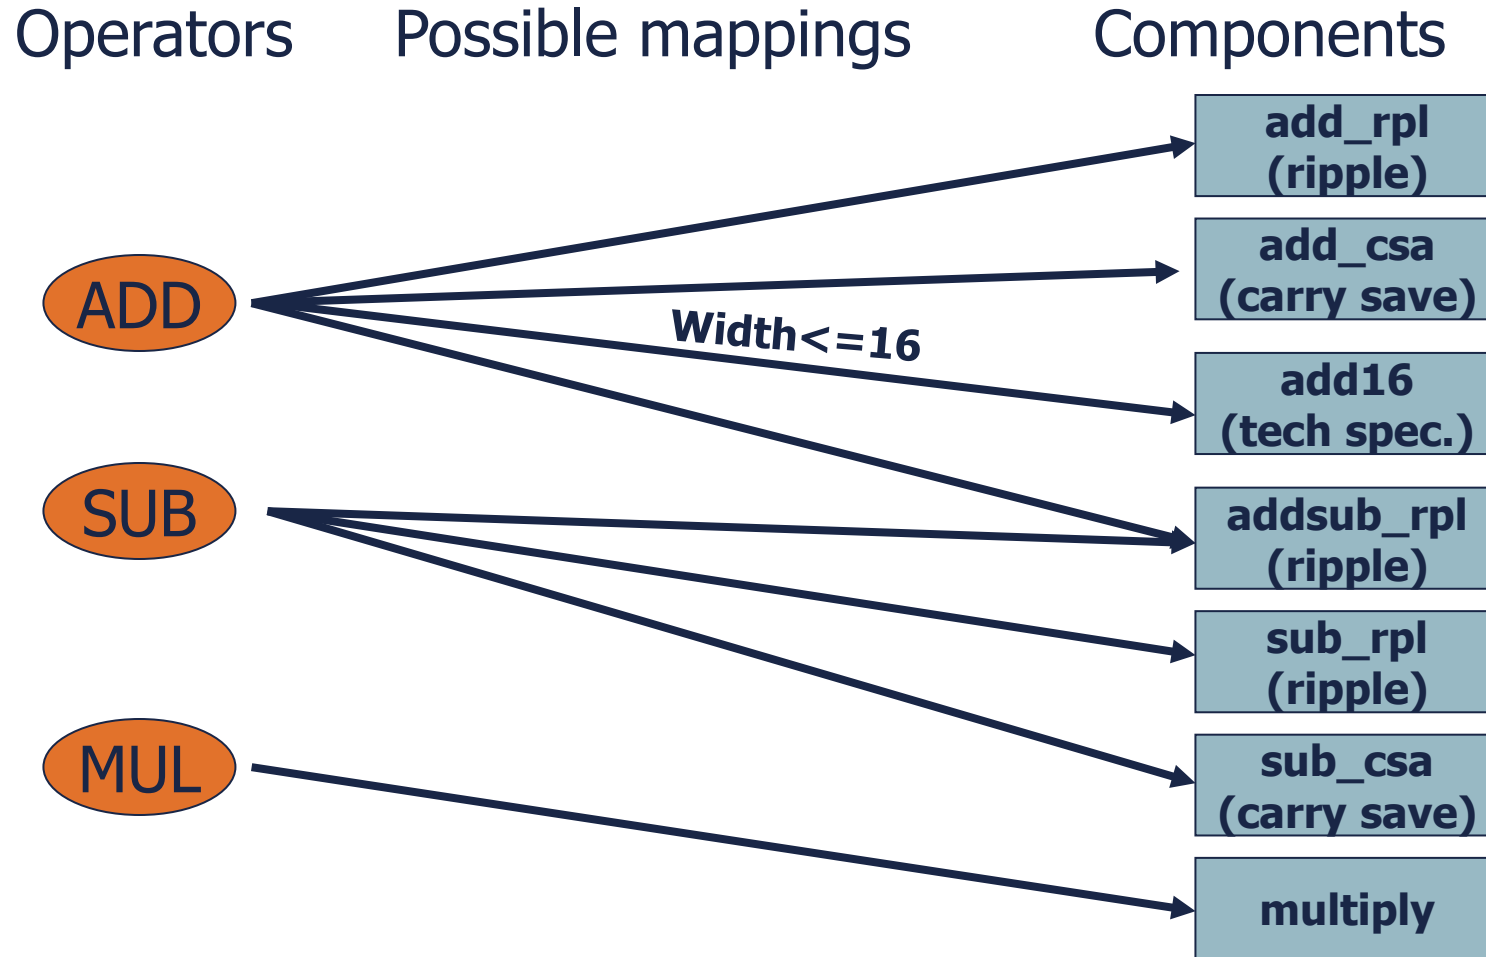
# High-Level Synthesis Steps



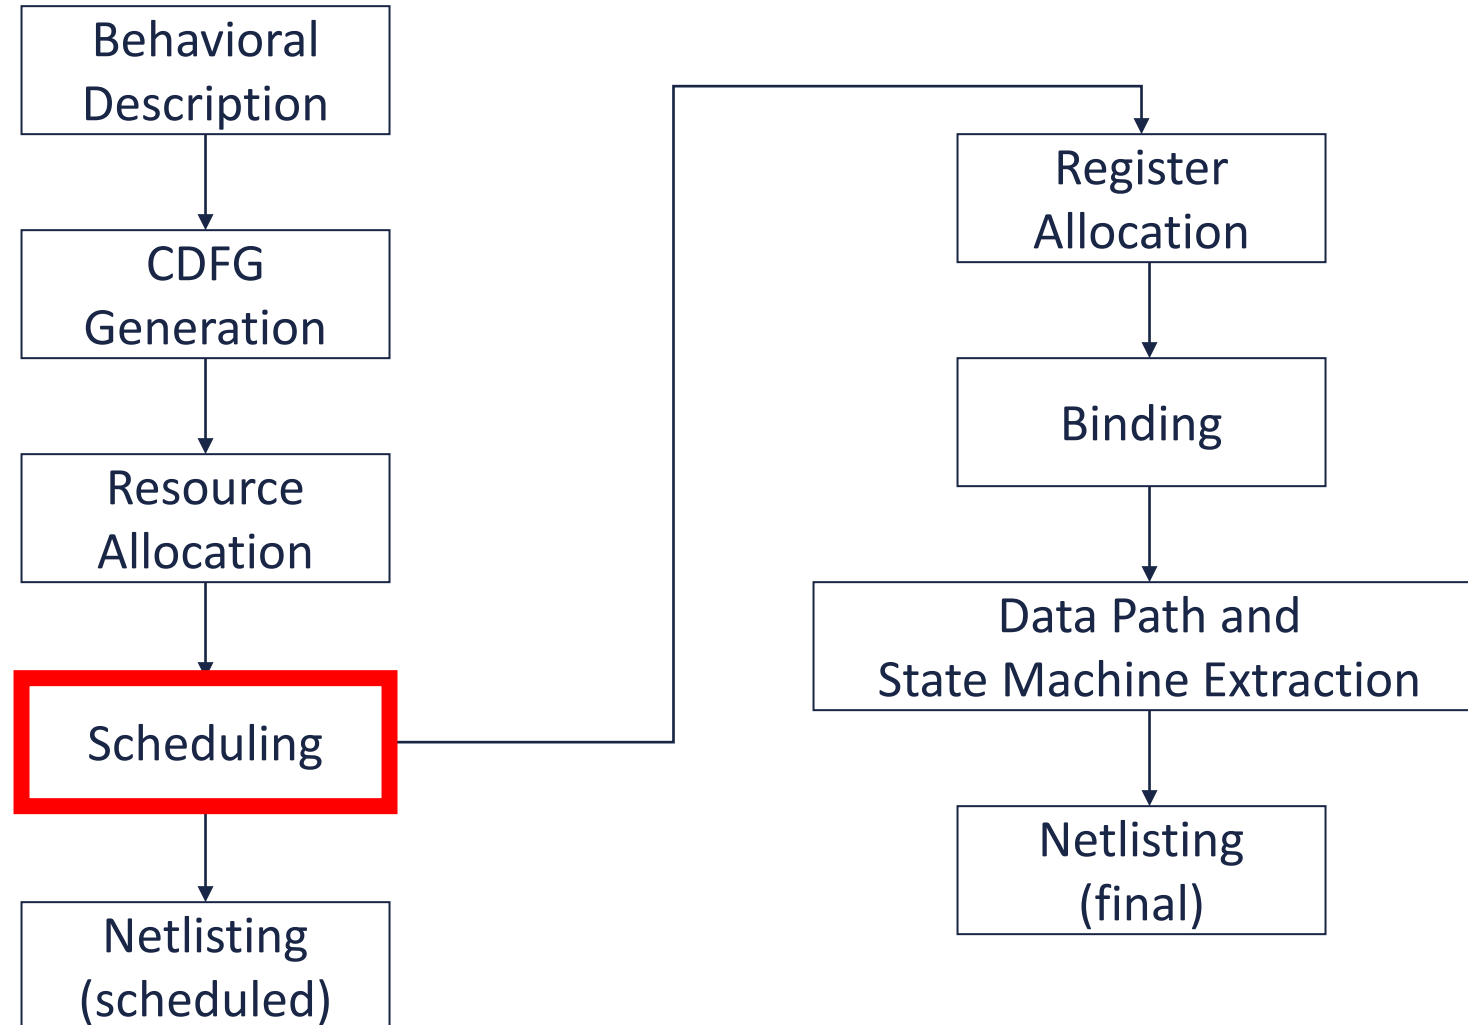
# HLS Step: Resource Allocation

- Deciding how many and which kinds of resources will be used in a given implementation
- This has a major impact on final design
  - Number of operation units (multiple adders?) set the maximum parallelism that the architecture can provide
  - Reuse of overloaded operators (e.g., an adder/subtractor unit) provides smallest designs
  - Choice of buses or muxes provides parallelism vs. size
  - Choice of registers, multi-ported register files or RAM also limits parallelism in data movement

# Example: Mapping of Operators to Components



# High-Level Synthesis Steps



# HLS Step: Scheduling

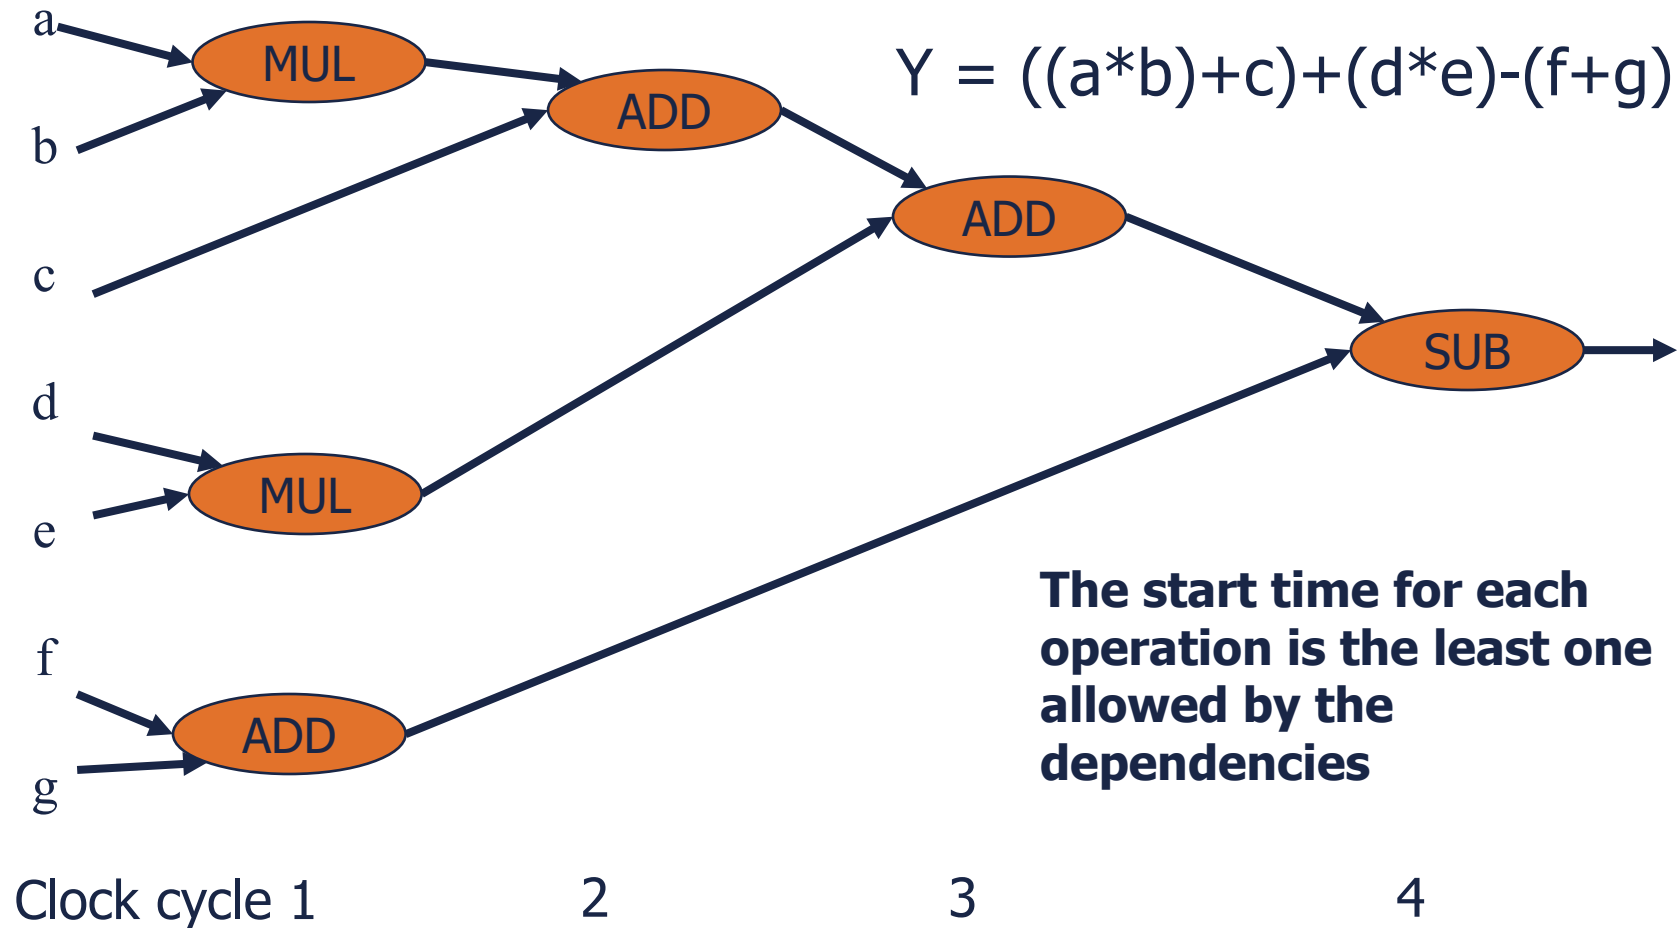
## *Scheduling:*

- Mapping of operations to time slots (cycles)

## *Basic Algorithms:*

- ASAP: As Soon As Possible
- ALAP: As Late As Possible
- List Scheduling

# ASAP Schedule (unconstrained)




*Can we do better than this in terms of the DFG itself? What is the resource usage?*

# ASAP Algorithm

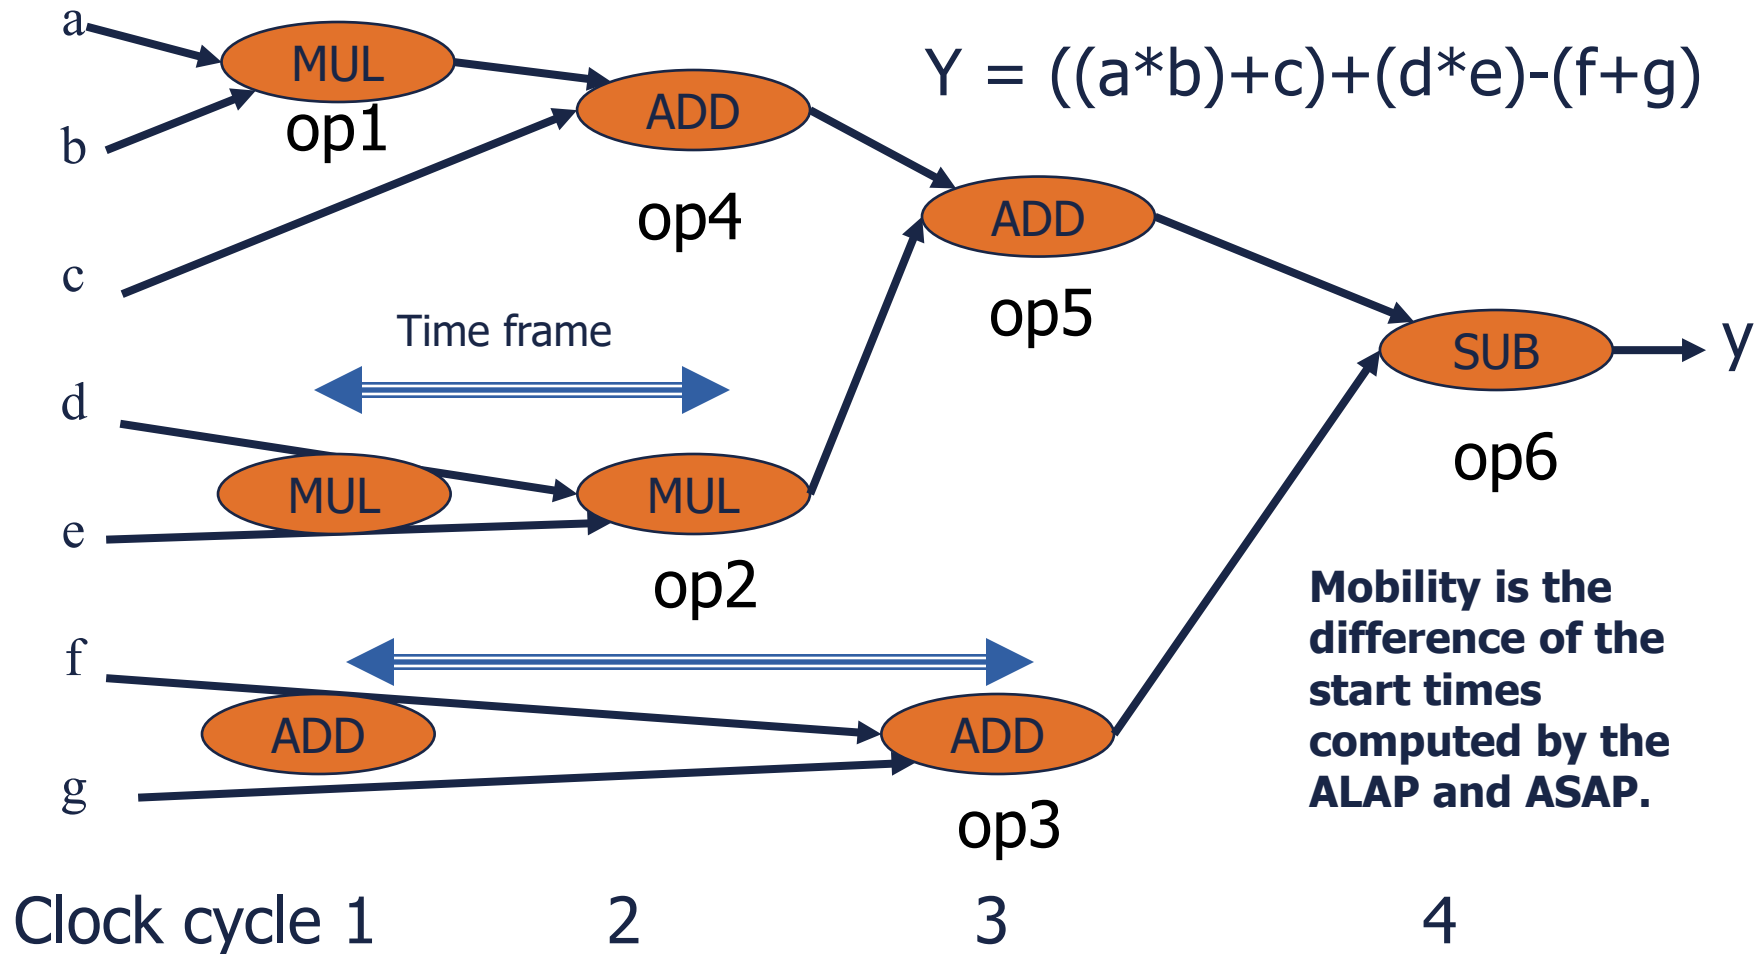
```
ASAP ( $G(V, E)$ ) {  
    Schedule all the nodes driven only by PIs to cycle 1,  
    for all such  $v_i$  nodes,  $t_i$  (starting time) = 1;  
    Repeat {  
        Select a vertex  $v_i$  whose predecessors are all scheduled;  
        Schedule  $v_i$  by setting  $t_i = \text{MAX}(t_j + d_j); (v_j, v_i) \in E$   
    }  
    Until all the nodes are scheduled;  
    Return the schedule in a vector;  
}
```

# ALAP Schedule

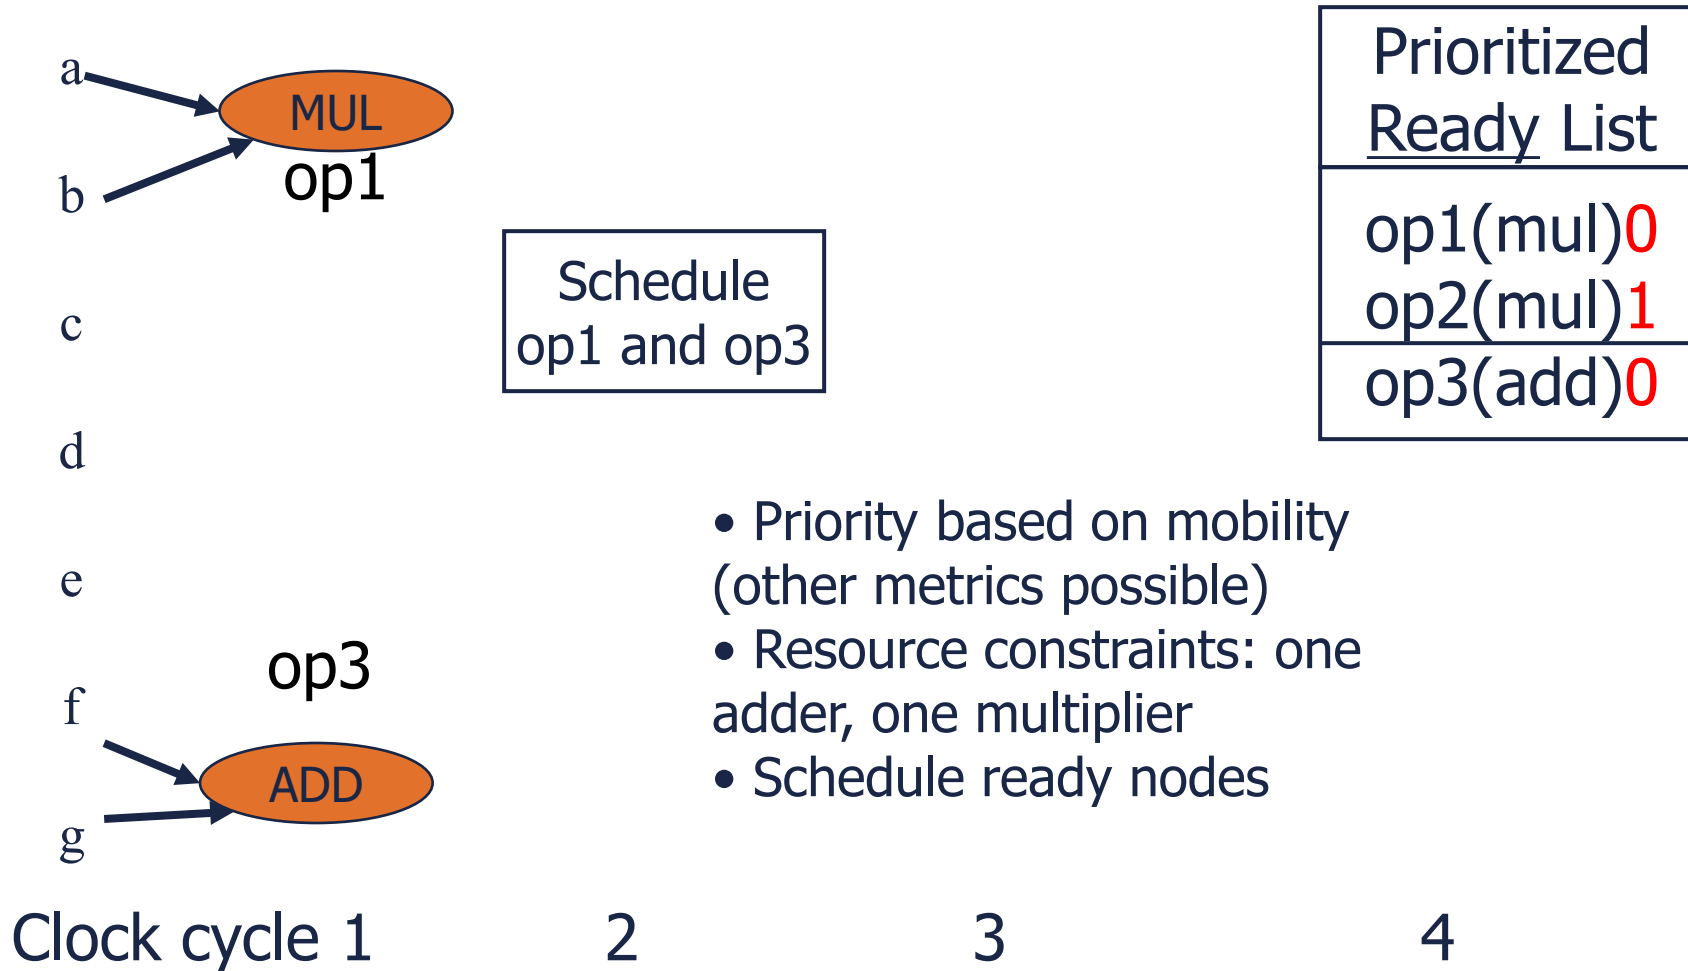




# Mobility (or Slack)

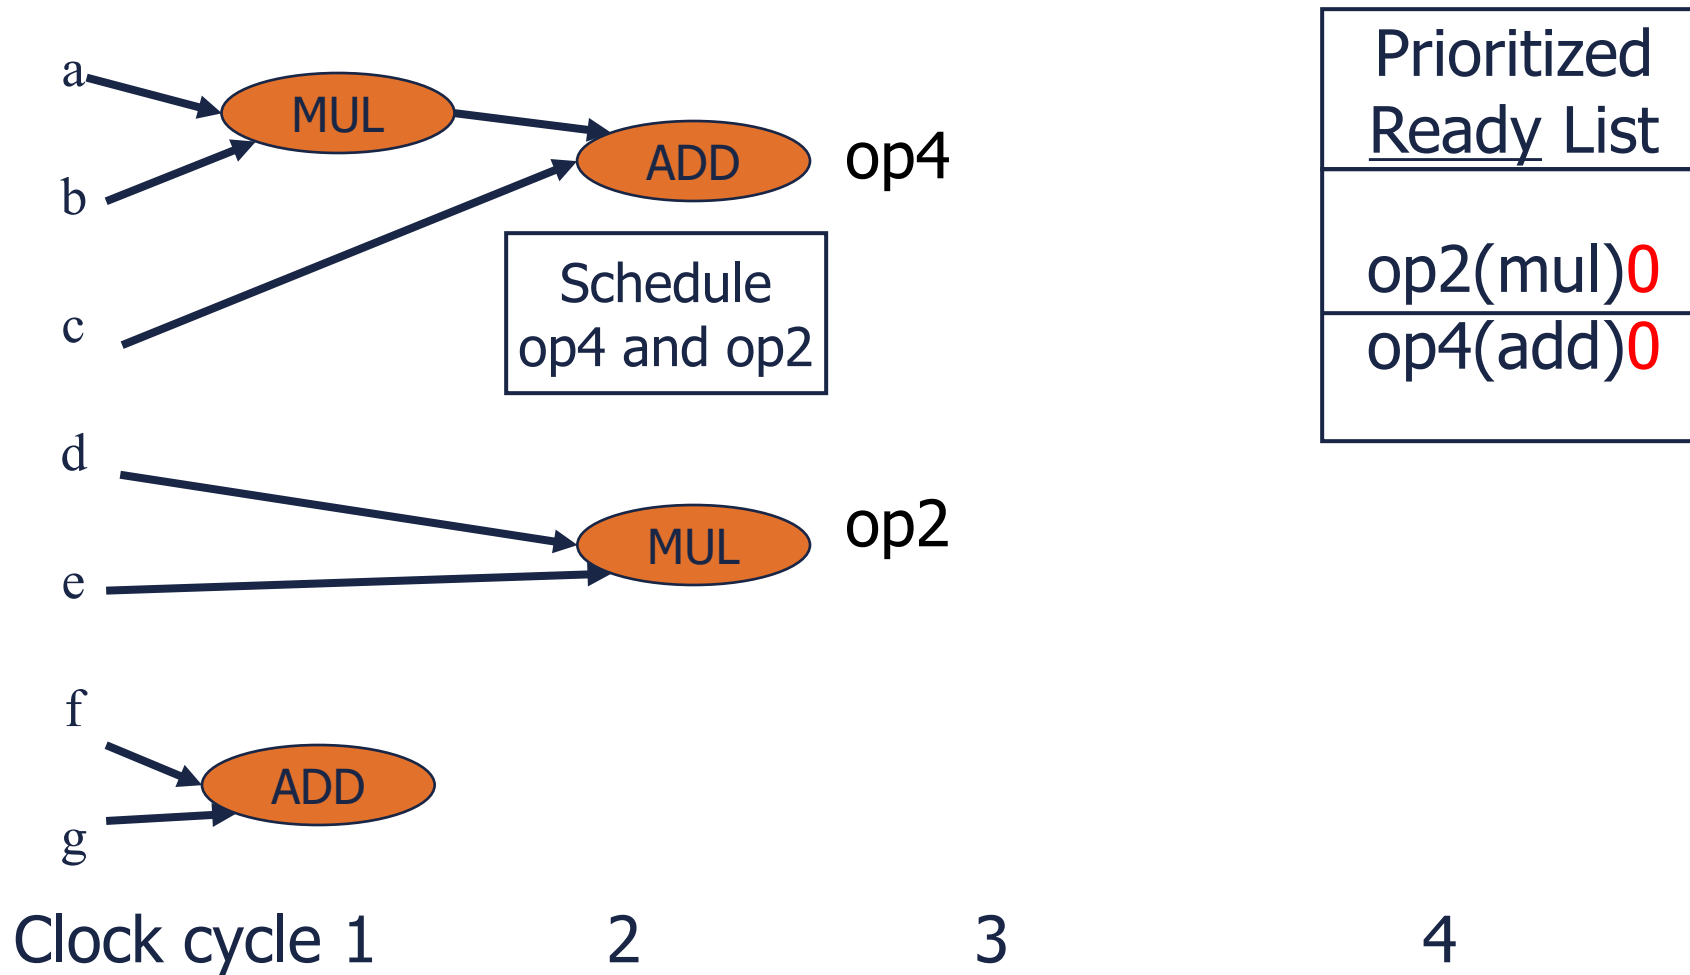


# List Scheduling (1)

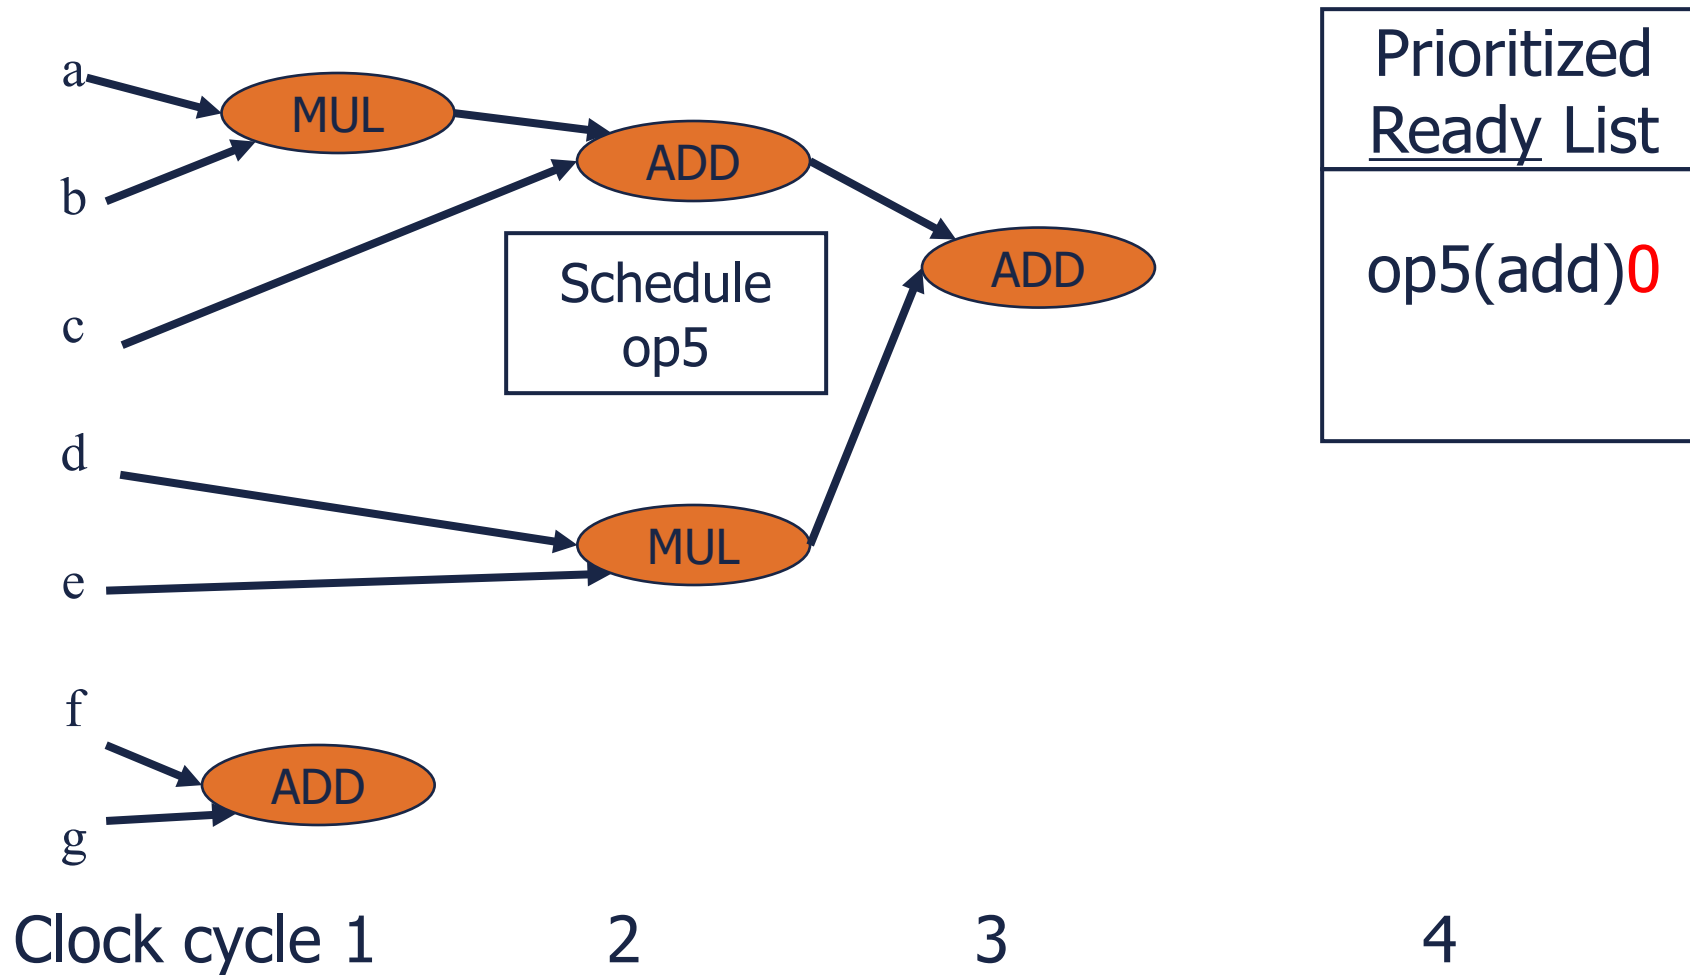


**A popular scheduling algorithm**

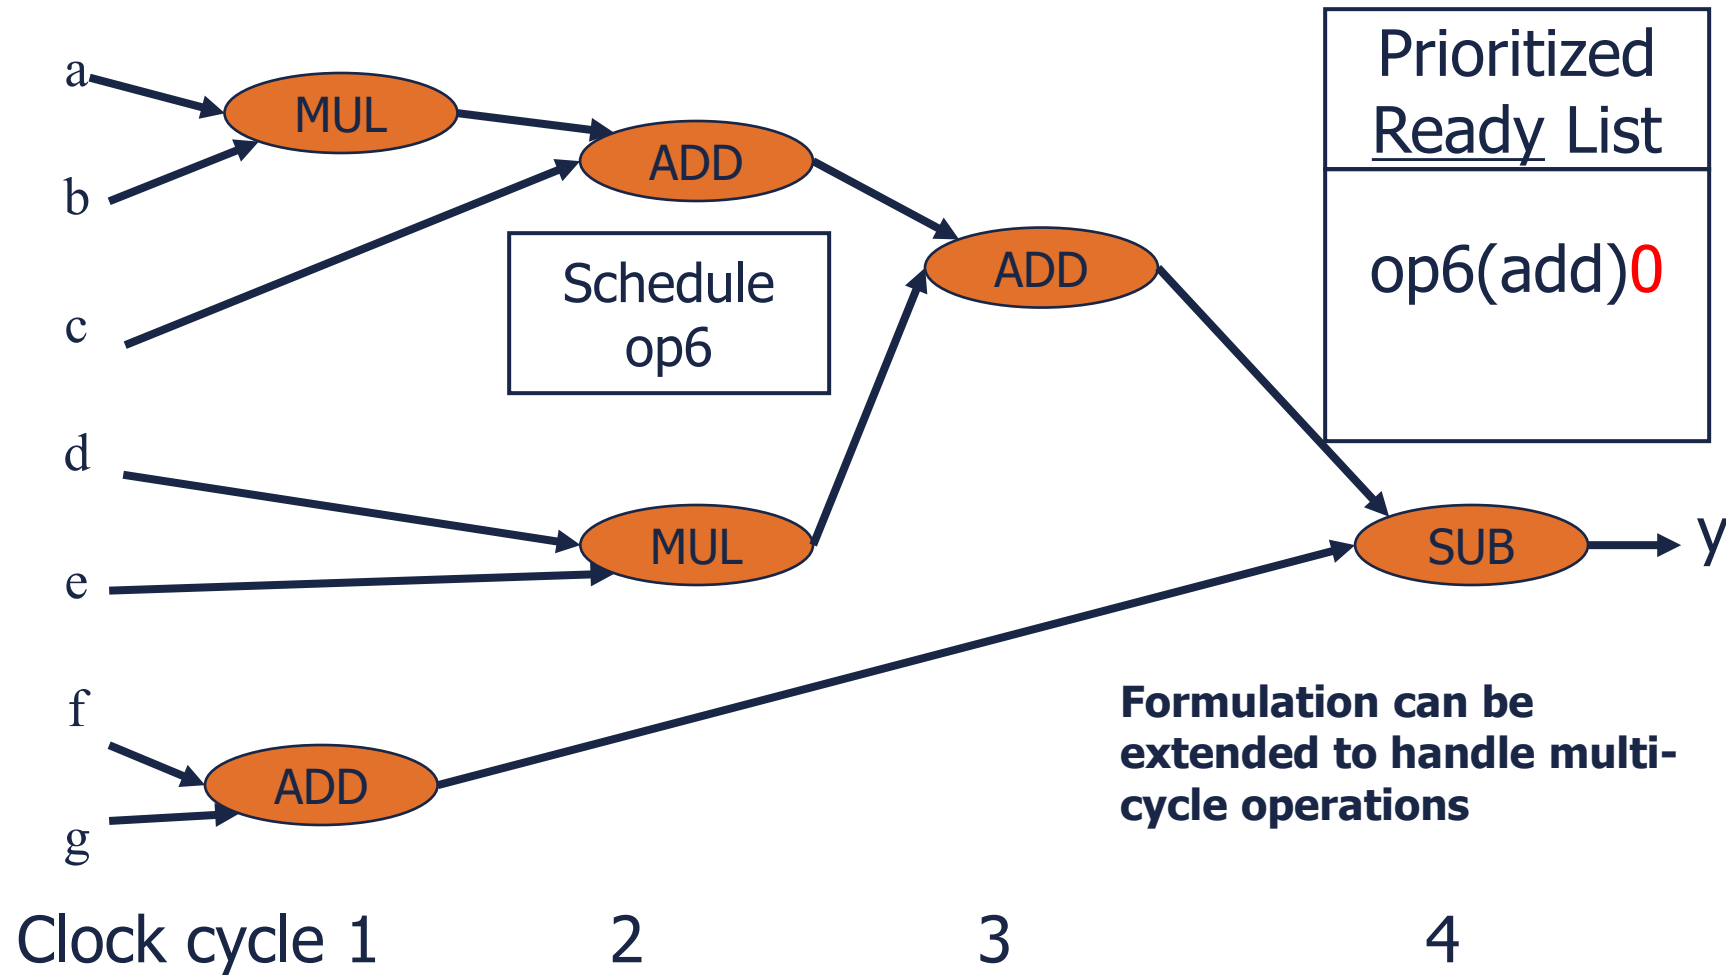
# List Scheduling (2)



# List Scheduling (3)



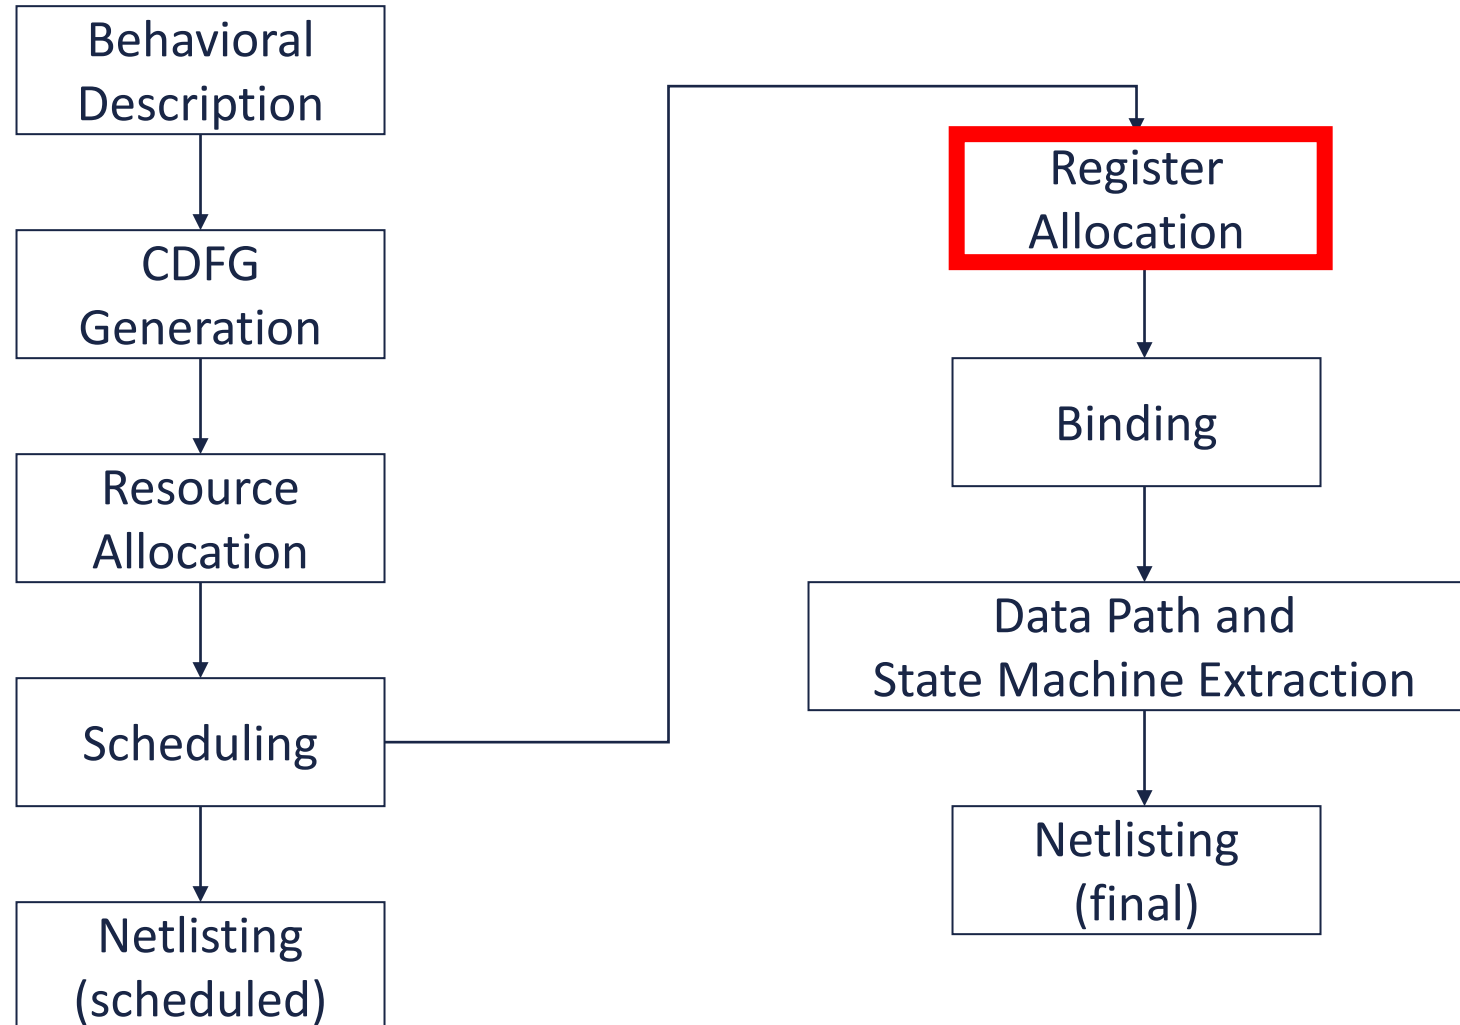
# List Scheduling (4)




# List Scheduling Algorithm

```
LIST_L( $G(V, E)$ ,  $a$ ) {  
     $l = 1$ ;  
    repeat {  
        for each resource type  $k = 1, 2, \dots, n_{res}$  {  
            Determine candidate operations  $U_{l,k}$ ;  
            Determine unfinished operations  $T_{l,k}$ ;  
            Select  $S_k \subseteq U_{l,k}$  vertices, such that  $|S_k| + |T_{l,k}| \leq a_k$ ;  
            Schedule the  $S_k$  operations at step  $l$  by setting  
                 $t_i = l$  for all  $i : v_i \in S_k$ ;  
        }  
         $l = l + 1$ ;  
    }  
    until (all nodes are scheduled);  
    return ( $t$ );  
}
```

# High-Level Synthesis Steps

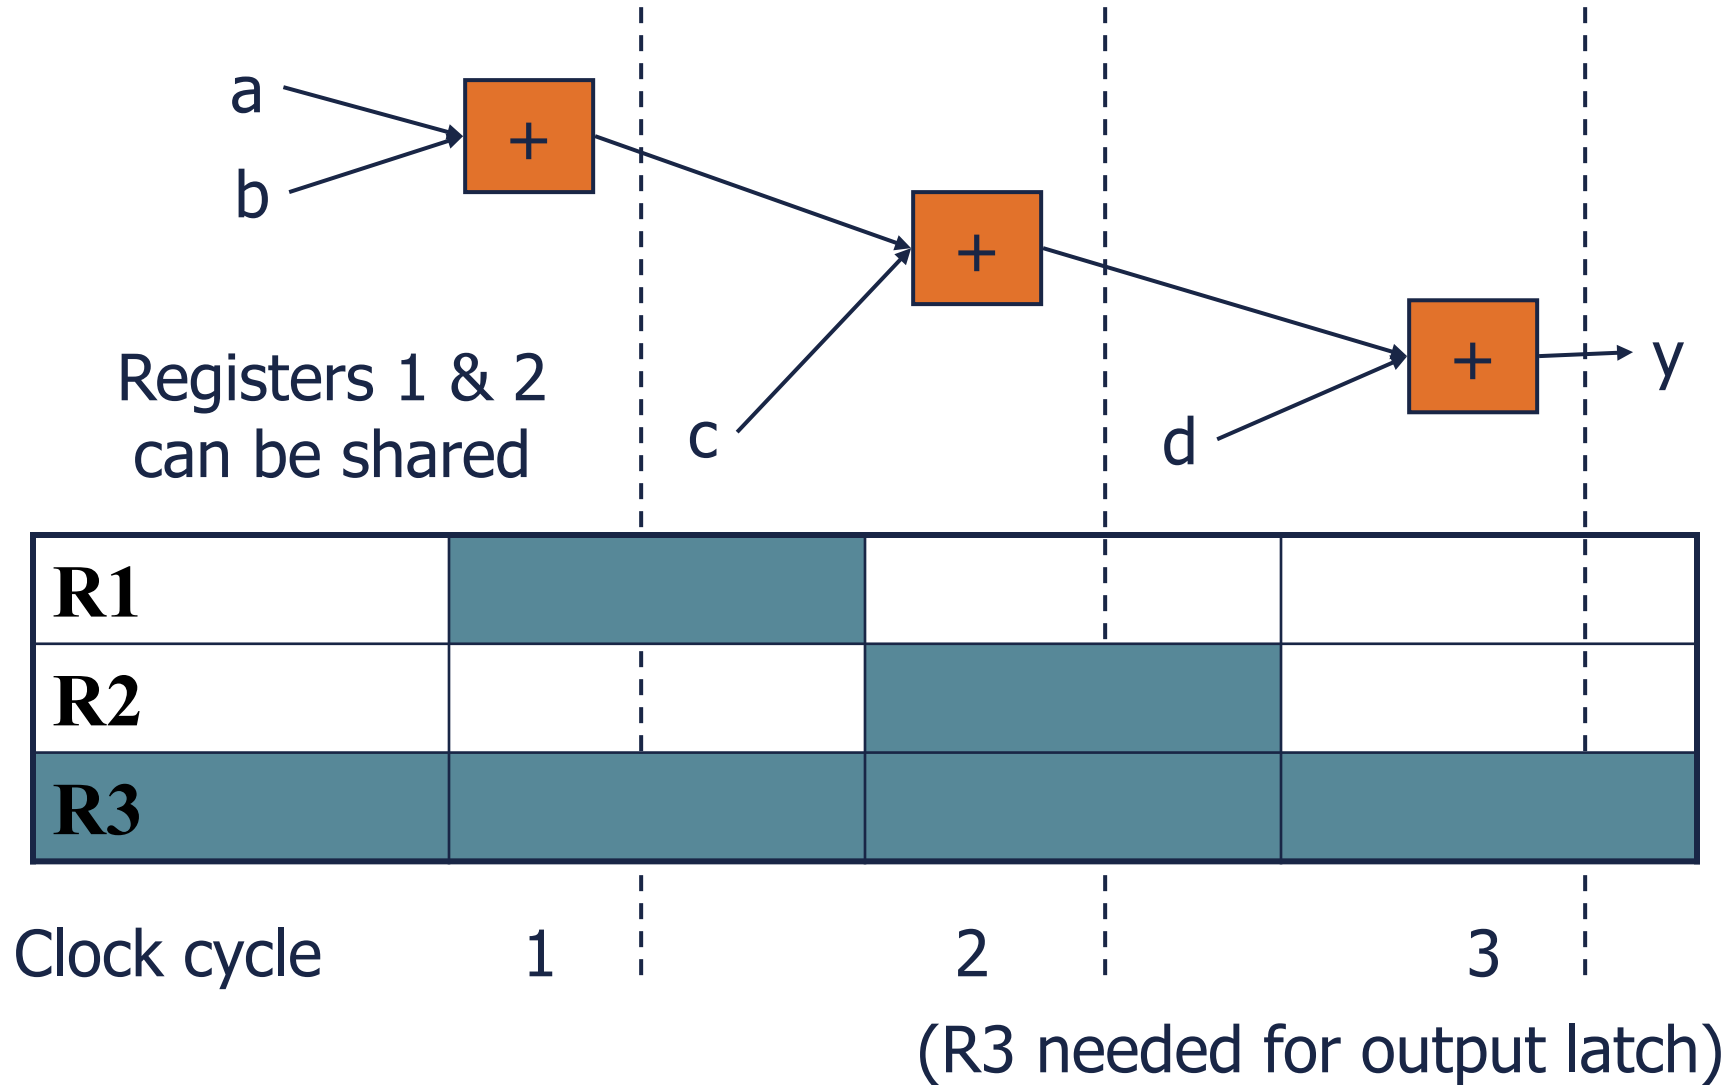


# Register Allocation

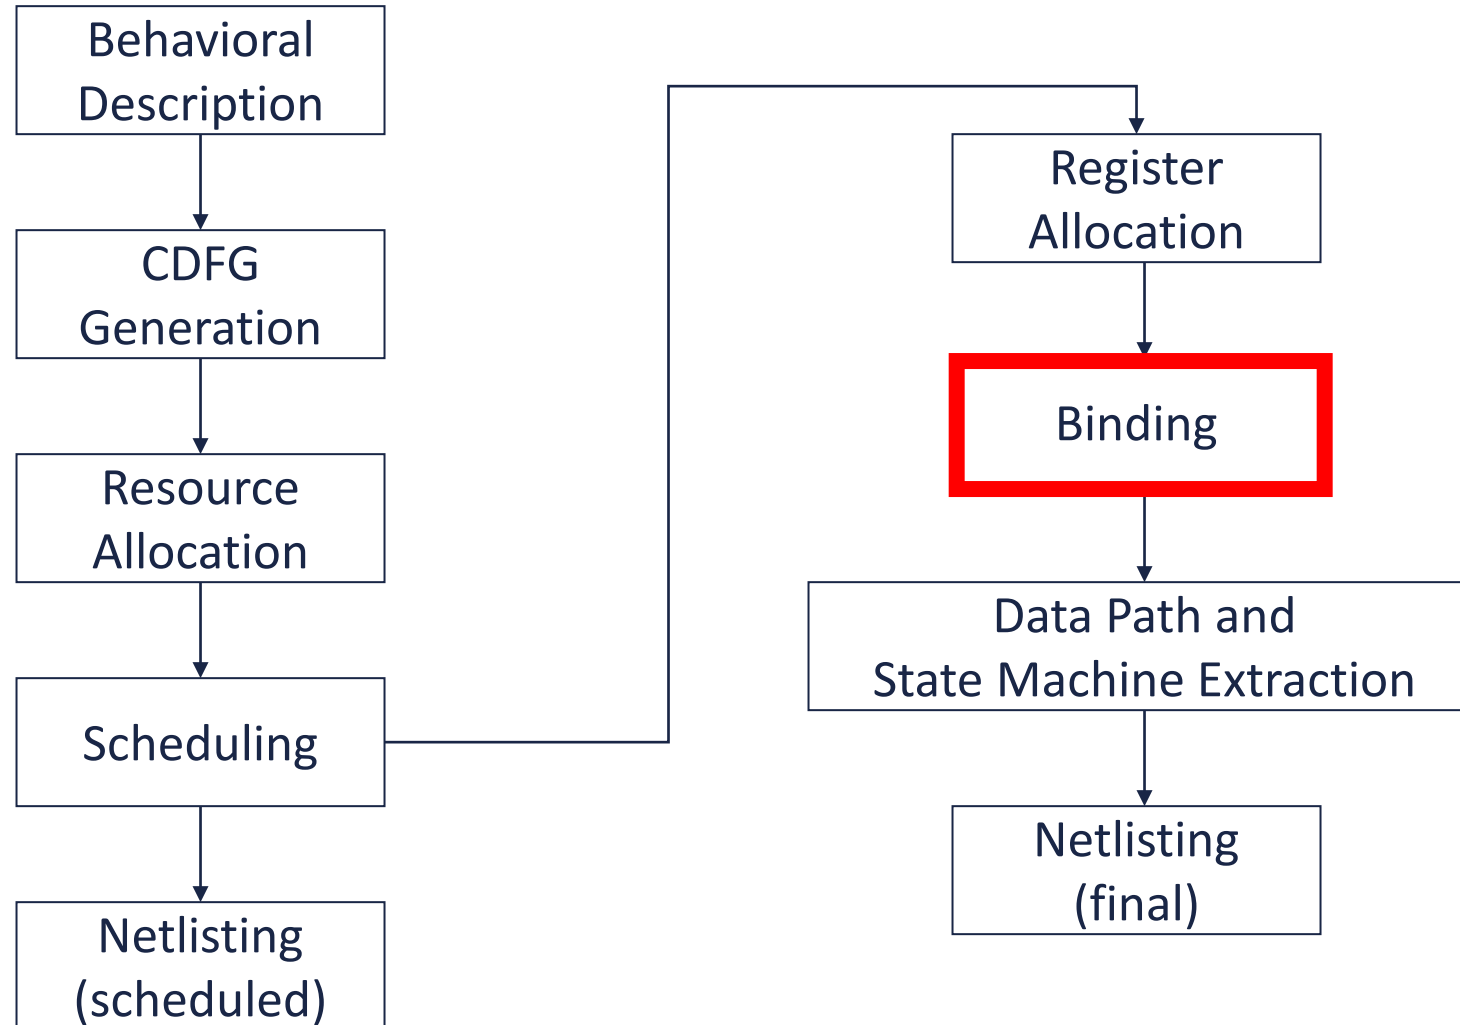





# Lifetime Analysis



# High-Level Synthesis Steps



# Binding



# Binding Choices

Binding 1


Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add1

Binding 2

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add2


# Binding 1 Results

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add1



# Binding 2 Results

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add2



# Linear Programming

- Objective function: a linear function to be maximized or minimized, e.g.,
  - Maximize  $c_1x_1 + c_2x_2$
- Constraints: e.g.,
  - $a_{11}x_1 + a_{12}x_2 \leq b_1$
  - $a_{21}x_1 + a_{22}x_2 \leq b_2$
  - $a_{31}x_1 + a_{32}x_2 \leq b_3$
- Non-negative variables, e.g.,
  - $x_1 \geq 0$
  - $x_2 \geq 0$
- The problem is usually expressed in *matrix form*, and then becomes:
  - Maximize  $\mathbf{c}^T \mathbf{x}$
  - Subject to  $\mathbf{Ax} \leq \mathbf{b}, \mathbf{x} \geq \mathbf{0}$

# Integer Linear Programming (ILP)

- If the variables are all required to be integers, then the problem is called an integer programming (IP) or integer linear programming (ILP) problem.
- In contrast to linear programming, which can be solved efficiently in the worst case, integer programming problems are in the worst case NP-hard.
- 0-1 integer programming is the special case of integer programming where variables are required to be 0 or 1 (rather than arbitrary integers). This method is also classified as NP-hard, and in fact the decision version was one of Karp's 21 NP-complete problems.



# ILP-Based Scheduling

- Use binary decision variables
  - $X = \{x_{il}; i = 1, \dots, n; l = 1, \dots, L\}$
  - $x_{il} \in \{0, 1\}$
- First the start time of each operation is unique

$$\sum_l x_{il} = 1, i = 1, 2, \dots, n$$

The start time of any operation  $v_i$ :

$$t_i = \sum_l l \cdot x_{il}$$

# ILP-Based Scheduling: Constraints

- Given the CDFG represented by  $G(V, E)$   $d_j$  : latency of operation  $v_j$   
 $t_i \geq t_j + d_j \quad \forall i, j : (v_j, v_i) \in E$  implies

$$\sum_l l \cdot x_{il} \geq \sum_l l \cdot x_{jl} + d_j \quad i, j = 1, 2, \dots, n : (v_j, v_i) \in E$$

- Resource bounds must be met at every schedule time step.

$$\sum_{i: T(v_i)=k} \sum_{m=l-d_i+1}^l x_{im} \leq a_k, \quad k = 1, 2, \dots, n_{res}, \quad l = 1, 2, \dots, L$$

# ILP-Based Scheduling: Formulation

- Denote  $t$  the vector whose entries are the start times
  - Minimize  $C^T t$  such that

$$\sum_l x_{il} = 1, i = 1, 2, \dots, n$$

$$\sum_l l \cdot x_{il} - \sum_l l \cdot x_{jl} - d_j \geq 0, i, j = 1, 2, \dots, n : (v_j, v_i) \in E$$

$$\sum_{i: T(v_i)=k} \sum_{m=l-d_i+1}^l x_{im} \leq a_k, k = 1, 2, \dots, n_{res}, l = 1, 2, \dots, L$$

$$x_{il} \in \{0, 1\}, i = 1, 2, \dots, n, l = 1, 2, \dots, L$$

# ILP-Based Scheduling: Vector $c$

- Minimize  $c^T t$

The start time of any operation  $v_i$  :

$$t_i = \sum_l l \cdot x_{il}$$

- $c = [0, \dots, 0, 1]^T$ 
  - Minimize the latency of the schedule
- $c = [1, \dots, 1, 1]^T$ 
  - Finding the earliest start times of all operations

# What designs fit HLS well?

- Data path, uni-directional data streaming pipeline
- Flexible in latency requirement, loose timing relationship between interfaces
- Computational kernels
  - E.g., image/video processing/arithmetic computation
  - Dataflow design is better than control flow design
  - HLS tool may provide IPs for arithmetic and other typical computational operations
- No inter-iteration dependencies
  - Preferably no dependencies between two loop iterations
  - Advanced HLS tool can handle dependencies to certain extent
- Independent kernels and module level designs
  - Not top level/system level designs entirely for HLS

# Challenges for HLS

- May not be a good fit for very high speed design
  - Designers can push until the last FF and create deep pipelines with RTL
- Cycle accurate design, strong timing relationship on interfaces
- Design with feedback loops and strong timing requirement
  - E.g., accumulators whose current result affects next cycle accumulation immediately
- Control-intensive logic -- but is getting better
- HLS may need designer's manual intervention in order to generate hierarchical designs with high quality
- Coarse-grain pipelining
  - Handling multiple kernels at the top level can be a challenge
  - Advanced HLS tools can work adequately in this area (an example to follow)
- Complicated protocol handling
- Analog/mixed signal portions of the design

# Where are these challenges from?

- HLS works with a design entry at a higher abstraction level
  - C/C++ has no idea of “time”
  - It might be do able to capture some timing concept, but the efficiency drops greatly, and we lose the benefit not worth it.
- RTL, as well as SystemC, has “time” built in
  - Fits design with accurate timing requirement best

# Coding with HLS

- Write code with sub functions corresponding to modules
  - With hardware implementation concepts in mind
  - Think about the inter-function communication scheme
- Generate hardware for each function using HLS, leave the arguments as interface/ports
- Generate the glue logic/communication and top-level design
  - Integrate the hardware block generated from HLS into a system (e.g., a virtual platform)
  - If the hardware is hierarchical, to achieve better QoR, manual work may be required



# Software vs. Hardware Compilation

C Construct	Example	Software Compilation	High-level Synthesis
Function	<code>void go() { ... }</code>	groups of instructions	hardware module
Function args and return	<code>int det(int a[16][16]);</code>	pushed onto call stack	input and output ports
Function call	<code>go();</code>	call instructions	submodule
Operations	<code>prod = a.x * b.x + a.y * b.y;</code>	computation instructions	functional units
Local variables	<code>unsigned index; int sum;</code>	architectural registers	physical registers
Arrays	<code>int A[16][16];</code>	stack, heap, static memory	memory blocks
Control flow	<code>for (i = 0; i &lt; 16; i++) { ... }</code>	branch instructions	state machines
source code	all of the above	machine instructions	hardware modules

# C/C++ Code Synthesizability for HLS

- In HLS, your code is the hardware specification

Unsynthesizable C Construct	Example	High-level Synthesis?	Synthesizable alternative
Variable memory allocation	<code>A = malloc(rows * cols * sizeof(A[0][0]));</code>	resizable RAM?	<code>int A[max_rows][max_cols];</code>
Recursion	<code>int bin_search(Array array) { ... bin_search(array); ... }</code>	infinitely nested modules?	use iterative form of binary search
Indirect calls /function pointers	<code>typedef void Callback(); void go(Callback callback) { ... callback(); ... }</code>	dynamically defined submodule?	inline or create “go” function for each callback
Operating system calls	<code>file = fopen(filename, “r”);</code>	I/O? Interrupt?	use input ports for streaming in data

# C/C++ Code Synthesizability for HLS (cont'd)

Construct	Example	Synthesizable?
Debug output	<code>printf("sum: %x\n", sum);</code>	Can be translated to unsynthesizable RTL debug functions.
Time functions	<code>clock_gettime(CLOCK_REALTIME, &amp;cur_time);</code>	Not synthesizable. HLS tools often have profiling support built-in.
C++ memory allocation	<code>Node * node = new Node();</code>	Equivalent to malloc; not synthesizable.
C++ template library	<code>std::vector&lt;int&gt; a; ... std::sort(a.begin(), a.end());</code>	Uses dynamic memory allocation and recursion; not synthesizable.
Floating point operations	<code>float a, b, c; ... float d = a * b + c;</code>	Supported by some HLS engines. Can be expensive, especially on FPGAs.


# Variables

- Customize bit-widths (arbitrary bit-widths) instead of standard C data types
  - E.g., use 24-bit width instead of 32 bits if 24 bits precision is enough
  - Most HLS tools support arbitrary bit-width data types
    - E.g., `ap_int`, `ap_fixed` data types in Vivado HLS
  - Improves efficiency, performance, and saves area

# Functions

- C functions are synthesized to RTL modules
  - By default, sub functions become sub modules
  - Usually, every function call is an instance of the RTL module
  - Inline function is an exception – it flattens the design

```
void top(int a, int b) {  
    foo1(a);  
    foo2(b);  
}  
void foo1(int a) {  
    foo1_1(a);  
}
```



# Functions: Splitting

- Use subfunctions to improve parallelism
  - Use function to maintain hierarchy
  - Most tools support dataflow/block-level pipelining and parallelism
  - May not be feasible due to dependencies

```
void top (int a[100], int b[100]) {  
    for(i=0; i<100; i++) {  
        a[i] = i;  
    }  
    for(j = 0; j < 100; j++) {  
        z += b[j];  
    }  
}
```



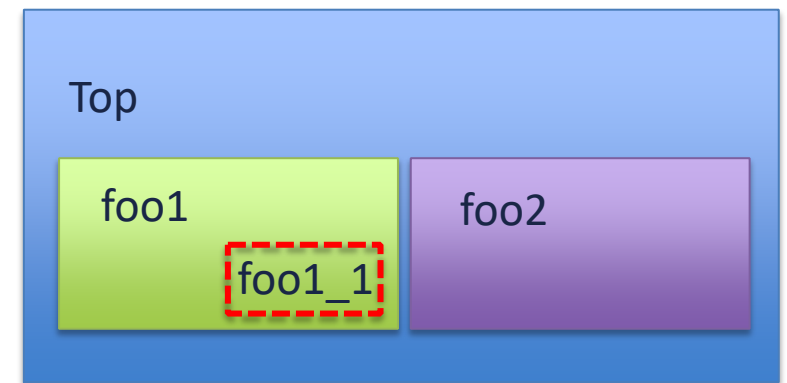
```
void foo1 (int a[100]) {  
    for(i=0; i<100; i++) {  
        a[i] = i;  
    }  
}  
void foo2 (int b[100]) {  
    for(j = 0; j < 100; j++) {  
        z += b[j];  
    }  
}  
void top (int a[100], int b[100]) {  
    foo1(a);  
    foo2(b);  
}
```



# Functions: Inlining

- Inline functions to flatten the hierarchy
- Pros:
  - Enable more optimization opportunities
    - Creates more objects to schedule and optimize
  - Enable resource sharing
  - Eliminates function call overhead
- Cons:
  - Module instance is effectively copied for each inlined call
  - Inlining of multiple function calls increases area cost

```
void top(){int a, int b){  
    foo1(a);  
    foo2(b);  
}  
void foo1(int a){  
    foo1_1(a);  
}  
inline void foo1_1(int a){}
```



# Inlining: Example


- Merge foo1 and foo2, and apply loop pipelining

```
int foo1(int a, int b, int c) {  
    z0 = a * b;          //s1_1  
    z1 = c * b;          //s1_2  
    return (z0+z1);  
}  
int foo2(int c, int d, int a){  
    z2 = c * d;          //s2_1  
    z3 = d * a;          //s2_2  
    return (z2+z3);  
}  
int top(int a, int b, int c, int d){  
    y1 = foo1(a, b, c);  
    y3 = foo2(c, d, a);  
    return (y3+y1);  
}
```



```
int top(int a, int b, int c, int d){  
    z0 = a * b;          //s1_1  
    z1 = c * b;          //s1_2  
    z2 = c * d;          //s2_1  
    z3 = d * a;          //s2_2  
    return (z0+z1+z2+z3);  
}
```

Assume each multiplication takes multiple cycles, and it can be pipelined



Also, enable resource sharing, e.g. for the multiplier



# Loops: Bounds

- Not recommend using variable loop bounds, because:
  - HLS tools may not be able to estimate the performance
  - Limits the optimization opportunities
    - E.g., cannot deliver complete loop unrolling, thus prevent the pipelining
    - Solutions: Set the bounds to the maximum value, and conditionally execute the loop body

```
for (i=0; i < len; i++) {  
    A[i] += b;  
}
```




```
for (i=0; i < N-1; i++) {  
    if(i < len)  
        A[i] += b;  
}
```

# Loop Unrolling

- Most HLS tools support automatic loop unrolling (by using HLS unroll pragmas)
  - But have strict limits, e.g., constant loop bounds
- May need manual loop unrolling to
  - Improve pipeline opportunity
  - Improve parallelism
- Things to notice
  - Need to consider memory port limit when accessing array
  - Expensive hardware replication
  - Need to balance the unroll level

```
void top(){  
    for(i=0; i<=N; i++){  
        a[i] = i;  
        ...  
    }  
}
```



# Loop Fusion/Merging

- Merge several loops into one loop
- Some tools support loop fusion
  - Strict bound check of loops, e.g., identical loop bounds
- Manually loop fusion can be used to
  - Improve data locality
  - Shorten latency
  - Share resources

```
void top(){  
    for(i=0; i<N; i++){  
        a[i] = i;  
    }  
    for(i=1; i<N; i++){  
        x+=a[i-1];  
    }  
}
```



```
void top(){  
    for(i=0; i<N; i++){  
        a[i] = i;  
        if(i>=1)  
            x+=a[i-1];  
    }  
}
```

After fusion, loop bodies are optimized together

# Loop Perfectization

- Convert loops to perfect loop
  - All statements are in the innermost loop
  - An example: Loop Sinking

```
for (i=0, i<20; i++){  
    a = 0;  
    for(j=0; j<20; j++){  
        a += A[j] + j;  
    }  
    B[i] = a * 20;  
}
```



```
for (i=0, i<20; i++){  
    for(j=0; j<20; j++){  
        if(j==0) a =0;  
        a += A[j] + j;  
        if(j==19) B[i] = a*20; }  
}
```

HLS result (Use Vivado\_HLS):

```
Loop_I:  
    *Trip count: 20  
    *Latency: 480  
Loop_J:  
    *Trip count: 20  
    *Latency: 21  
    *Pipeline II: 1  
    *Pipeline depth: 2
```



```
Loop_I_Loop_J:  
    *Trip count: 400  
    *Latency: 401  
    *Pipeline II: 1  
    *Pipeline depth: 2
```

**The original code can only pipeline Loop\_J. After perfection, two loops can be flattened and pipelined.**

# Why Loop Perfectization?

- Form very regular loop structure
  - Some tools can only flatten the perfect loops
  - Easy for HLS scheduling and discovering optimization opportunities
  - Improves performance by eliminating the loop entering/exiting/cost
- Different ways of loop perfection
  - Loop sinking
  - Loop unrolling
  - Loop splitting

# Loop Perfectization: Code Sinking

```
for (i=0, i<20; i++){  
    a = 0;  
    for(j=0; j<100; j++){  
        a += A[j] + j;  
    }  
    B[i] = a * 20;  
}
```



```
for (i=0, i<20; i++){  
    for(j=0; j<20; j++){  
        if(j==0) a = 0;  
        a += A[j] + j;  
        if(j==19) B[i] = a*20;  
    }  
}
```

- Pros:
  - Easy to implement
  - No performance penalty for branch operations in pipeline
- Cons:
  - Expensive for hardware
    - sometimes the branches can be messy when code is complicated
  - Could sequentialize loop nest and affect performance
    - Change the data dependency

# Loop Perfectization: Loop Unrolling

```
for (i=0; i<100; i++) {  
    z = a*b[i];  
    z+= c*d[i];  
    for(j=0; j<4; j++) {  
        z+=x[j] + j;  
    }  
}
```



```
for (i=0, i<100; i++) {  
    z = a*b[i];  
    z += c*d[i];  
    z += x[0] + 0;  
    z += x[1] + 1;  
    z += x[2] + 2;  
    z += x[3] + 3;  
}
```

- Pros:
  - Enable pipelining at outer loop
  - Enable resource sharing
    - More statements to schedule in the same loop level
  - Higher performance, good for expensive operations
- Cons:
  - Not always feasible
  - Expensive hardware, best if the innermost loop is small

# Loop Perfectization: Loop Splitting

```
for (i=0; i<100; i++){  
    z = a*b[i];  
    z += c*d[i];  
    for(j=0; j<4; j++){  
        z+=x[i] + j;  
    }  
}
```



```
for(i=0; i<100; i++) {  
    z = a*b[i];  
    z += c*d[i];  
}  
for(i=0, i<100; i++) {  
    for(j=0; j<4; j++) {  
        z += x[0] + 0;  
        z+= x[1] + 1;  
        z+= x[2] + 2;  
        z+= x[3] + 3;  
    }  
}
```

- Pros:
  - Enable optimization per loop
  - Opportunity for dataflow streaming between two loops
- Cons:
  - Not always possible because of data dependency



# Arrays

- By default, it is implemented as RAM/ROM
- Expensive and have limited ports
- Optimizations:
  - Array streaming
    - Replace the RAM with FIFO, if possible
    - Requires array access to be sequential
    - Requires the array access order to be the same
    - Transform the code to make array access sequential
  - Array partitioning
    - Eliminate the ports limit


# Array Streaming: Example

```
for (i=0; i<100; i++)
    for(j=0; j<100; j++){
        A[i][j] = a * b; //row
        ...
    }
}
for (i=0; i<100; i++)
    for(j=0; j<100; j++) {
        C[i][j] = A[j][i]; //col
    }
}
```




```
hls::stream<int> A;
for (i=0, i<100; i++){
    for(j = 0; j<100; j++)
        A.write(a*b);
}
for(i=0, i<100; i++){
    for(j = 0; j<100; j++){
        A.read(C[j][i]);
    }
}
```

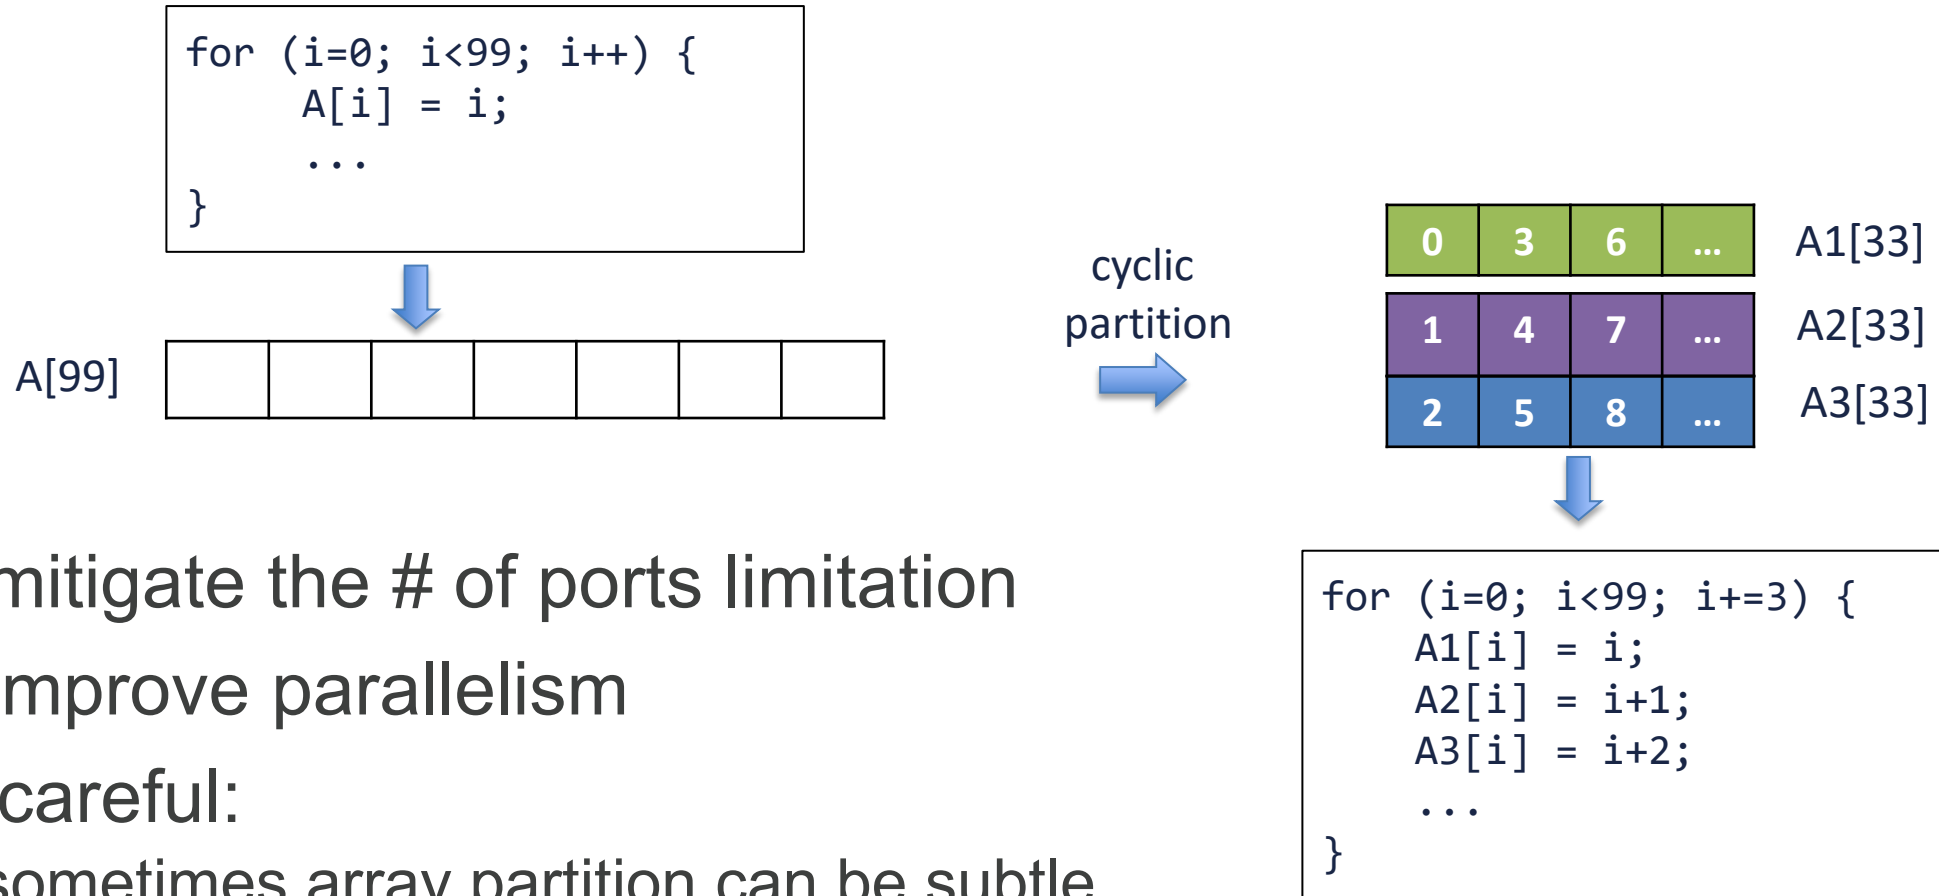
- Original array A access order:
  - Loop1: row; Loop2: column
  - No dataflow pipeline
  - No streaming



- New array A access order:
  - Loop1: row; Loop2: row
  - Dataflow pipeline
  - Streaming



# Array Partitioning



- To mitigate the # of ports limitation
- To improve parallelism
- Be careful:
  - sometimes array partition can be subtle
  - can easily introduce bugs

# Some Tips for C-based HLS Design

- Understand the HLS tool behavior and the optimization goal
  - No unique coding style is “useful” or “good” for all optimizations
- Reuse data and minimize array access
  - Once the data has been read into a block or buffer and reused, it can improve the parallelism.
  - Array ports are limited. Array partitioning can be an alternative, but it's not free.
- Loop transformation to enable parallelism
  - Pipelining/Unrolling/Dataflow
- Bit-width selection for efficient hardware
- The optimization goals for CPU and Hardware are different
  - E.g., branch mis-prediction
  - But many software code optimization techniques are useful for hardware design too

EECS 221:

# Languages and Compilers for Hardware Accelerators

*(Winter 2022)*

Sitao Huang

[sitaoh@uci.edu](mailto:sitaoh@uci.edu)

